
Time-MMD: Multi-Domain Multimodal Dataset for Time Series Analysis

Haoxin Liu^{†,*}, Shangqing Xu[†], Zhiyuan Zhao[†], Ling kai Kong[†],
Harshavardhan Kamarthi[†], Aditya B. Sasanur[†], Megha Sharma[†],
Jiaming Cui[†], Qingsong Wen[§], Chao Zhang[†], B. Aditya Prakash^{†,*}

[†]Georgia Institute of Technology [§]Squirrel AI

Abstract

Time series data are ubiquitous across a wide range of real-world domains. While real-world time series analysis (TSA) requires human experts to integrate numerical series data with multimodal domain-specific knowledge, most existing TSA models rely solely on numerical data, overlooking the significance of information beyond numerical series. This oversight is due to the untapped potential of textual series data and the absence of a comprehensive, high-quality multimodal dataset. To overcome this obstacle, we introduce Time-MMD, the first multi-domain, multimodal time series dataset covering 9 primary data domains. Time-MMD ensures fine-grained modality alignment, eliminates data contamination, and provides high usability. Additionally, we develop MM-TSFlib, the first multimodal time-series forecasting (TSF) library, seamlessly pipelining multimodal TSF evaluations based on Time-MMD for in-depth analyses. Extensive experiments conducted on Time-MMD through MM-TSFlib demonstrate significant performance enhancements by extending unimodal TSF to multimodality, evidenced by over 15% mean squared error reduction in general, and up to 40% in domains with rich textual data. More importantly, our datasets and library revolutionize broader applications, impacts, research topics to advance TSA. The dataset and library are available at <https://github.com/AdityaLab/Time-MMD> and <https://github.com/AdityaLab/MM-TSFlib>.

1 Introduction

Time series (TS) data are ubiquitous across a wide range of domains, including economics, urban computing, and epidemiology (51; 55; 28). Analytical tasks on such datasets hence find broad applications in various real-world scenarios such as energy forecasting, traffic planning, and epidemic policy formulation. Human experts typically complete such Time-Series Analysis (TSA) tasks by integrating multiple modalities of time-series data. For instance, epidemiologists combine numerical data on influenza infections with textual domain knowledge, policies, and reports to predict future epidemiological trends. However, most existing TSA models (57; 41; 76; 43; 79; 31; 78; 47) are unimodal, solely using numerical series.

Recently, with the development of Large Language Models (LLMs), the field of TSA is also undergoing an exciting transformative moment with the integration of natural language (71; 33). Existing LLM-based TSA methods incorporate endogenous text derived from numerical series, such as linguistic descriptions of statistical information, which has demonstrated promising benefits (21; 24; 6; 39). However, the potential of exogenous or auxiliary textual signals—such as information on concurrent

*Correspondence to: Haoxin Liu <hliu763@gatech.edu>, B. Aditya Prakash <badityap@cc.gatech.edu>

events and policies that provide additional context to time series—remains untapped. This observation prompts a crucial question for multimodal TSA: **Can multimodal TSA models utilize these exogenous textual signals effectively, thereby enhancing current TSA tasks and enabling new applications?**

The primary obstacle in addressing this question lies in the absence of a comprehensive, high-quality multimodal TS dataset, as evidenced by three significant gaps: (1) **Narrow data domains.** Data characteristics and patterns vary between different domains, such as the periodicity of numerical data and the sparsity of textual data. However, current multimodal TS datasets (15; 68; 11; 48; 6) focus solely on stock prediction tasks in the financial domain, which are unable to represent the diverse data domains. (2) **Coarse-grained modality alignment.** Existing multimodal TS datasets only ensure that the text and numerical data come from the same domain, such as general stock news and the prices of one specific stock. Clearly, an abundance of irrelevant text diminishes the effectiveness of multimodal TSA. (3) **Inherent data contamination.** Existing multimodal TS datasets overlook two main reasons of data contamination: (1) Textual data often contains predictions. For example, influenza outlook is a regular section in influenza reports. (2) Outdated test set, particularly the textual data, may have been exposed to LLMs, which are pretrained on vast corpuses. For example, the knowledge cutoff for Llama3-70B is December 2023, which is later than the cutoff dates for most existing multimodal TS datasets. These reasons lead to biased evaluations of general or LLM-based TSA models.

To address the identified gaps, this work aims to introduce a comprehensive, high-quality multimodal TS dataset that spans diverse domains and can be validated through its effectiveness and benefits for TSA. The main contributions of our work are:

- **Pioneering Multi-Domain Multimodal Time-Series Dataset: Time-MMD.** We introduce Time-MMD, the first multi-domain multimodal time-series dataset that addresses the aforementioned gaps: (1) encompasses 9 primary data domains. (2) ensures fine-grained modality alignment through meticulously selected data sources and rigorous filtering steps. (3) disentangles facts and predictions from text; ensures all cutoff dates are up to May 2024. To the best of our knowledge, Time-MMD stands as the inaugural high-quality and comprehensive multimodal time-series dataset. We envision Time-MMD offering exciting opportunities to significantly advance time series analysis through multimodal extensions.
- **Pilot Multimodal Time-Series Forecasting Library: MM-TSFlib.** We develop the first multimodal time-series forecasting (TSF) library, MM-TSFlib, piloting multimodal TSA research based on Time-MMD. Our library MM-TSFlib features an end-to-end pipeline with a seamless interface that allows the integration of any open-source language models with arbitrary TSF models, thereby enabling multimodal TSF tasks. MM-TSFlib facilitates easy exploration of Time-MMD and supports future advancements in multimodal TSA.
- **Extensive Evaluations with Significant Improvement.** We conducted over 1,000 experiments of multimodal TSF on Time-MMD using MM-TSFlib. The multimodal versions outperformed corresponding unimodal versions in 95% of cases, reducing the mean squared error by an average of over 20% and up to 40% in some domains with rich textual data. This significant and consistent improvement demonstrates the high quality of Time-MMD, the effectiveness of MM-TSFlib, and the superiority of multimodal extensions for TSF.

We include additional related works in Appendix A and limitations in Appendix B.

2 Multi-Domain Multimodal Time-Series Dataset: Time-MMD

We first introduce the key challenges in constructing Time-MMD, followed by the construction pipeline. We then detail each component of the pipeline with corresponding data quality verification. Finally, we discuss considerations for fairness and data release.

Challenges. Creating a high-quality, multi-domain numerical-text series dataset presents significant challenges, encompassing the effective gathering, filtering, and alignment of useful textual data. First, textual sources are sparse. Unlike numerical data, typically provided by a "packaged" source, textual data are collected from a variety of dispersed sources, such as reports and news articles, necessitating extensive individual collection efforts. Second, textual information is noisy. Raw textual data often contains large portions of irrelevant information and potential data contamination, such as expert predictions in reports, requiring rigorous filtering processes to ensure data quality.




Figure 1: Overview of the Time-MMD construction. We first construct numerical data, then construct textual data from search and report sources with LLM preprocessing targeted at the numerical data, and finally annotate the data with binary timestamps to support various downstream tasks.

Third, textual data requires precise alignment. It is essential to achieve temporal alignment between textual and numerical data by synchronizing reported times with numerical time steps (e.g., the time step where text is posted) and ensuring that the effective duration of textual information matches the relevant time frames at various granularities (e.g., a seasonal report should correspond to 12 time steps in a weekly time series). Additionally, the dataset faces challenges regarding ease of use, maintenance, and regular updates to remain relevant and useful for ongoing research and applications.

Pipeline Overview. We propose a comprehensive pipeline for constructing a text-numeric series dataset utilizing modern LLMs. As illustrated in Figure 1, the construction process is divided into three key steps: (1) Numerical Series Data Construction. We gather numerical data from reputable sources to ensure reliability and accuracy. (2) Textual Series Data Construction. Textual data is collected for fine-grained matching with the numerical data. The quality of this matching is ensured through human selection of data sources and raw text filtering by LLMs. Additionally, LLMs are employed to disentangle facts and predictions and generate summaries. (3) Numerical-Textual Alignment. We use binary timestamps to mark the start and end dates as a universal temporal alignment method between numerical and textual series, supporting the requirements of various downstream TSA tasks.

2.1 Numerical Series Data Construction

Data Source Selection. We select data sources that are (1) reliable, containing verified knowledge; (2) actively released, allowing for updates with new data; and (3) multi-domain, covering various TSA patterns. Appendix C provides considerations for domain selection. Based on these principles, we choose nine data sources from different domains, as shown in Appendix D. Most sources are from government agencies, with the lowest update frequency being six months.

Target Variable Selection. For each domain, we select target variables with significant real-world implications, indicating easier text matching, as shown in Table 1. These variables span three distinct frequencies: daily, weekly, and monthly.

Collection & Preprocessing. We collect raw data for all available times, either from batch-released files or through individual scraping. We preprocess the data by discarding early years with a high proportion of missing values. We maintain the original frequency for most domains, adjusting it for security and climate domains due to irregular releases and difficult text matching, respectively. Figure 2 illustrates the diverse patterns present in each domain, such as periodicity and trends.

Data Quality & Property. As shown in Table 1 and Figure 2, the constructed numerical data provides comprehensive temporal coverage, ranging from the earliest in 1950 to the present, and exhibits distinct patterns, such as periodicity and trends.

2.2 Textual Series Data Construction

Data Source Selection: Selected Reports and Web Search Results. The choice of data sources should take into account both extensive coverage and initial strong relevance to the numerical data. Consequently, we combine two appropriate data source types as follows: (1) Selected Reports: For

Table 1: Overview of numerical data in Time-MMD, covering key variables across nine domains with daily, weekly, or monthly frequencies, sourced from reputable government departments. Eight domains are updated to May 2024; the environment domain update is scheduled for June 2024. We focus on univariate time series forecasting for the target variable, as Time-MMD requires constructing aligned textual data. Covariates are provided for some datasets to inspire future research which are revealed in the dimension column.

Domain	Target	Dimension	Frequency	Number of Samples	Timespan
Agriculture	Retail Broiler Composite	1	Monthly	496	1983 - Present
Climate	Drought Level	5	Monthly	496	1983 - Present
Economy	International Trade Balance	3	Monthly	423	1989 - Present
Energy	Gasoline Prices	9	Weekly	1479	1996 - Present
Environment	Air Quality Index	4	Daily	11102	1982 - 2023
Health	Influenza Patients Proportion	11	Weekly	1389	1997 - Present
Security	Disaster and Emergency Grants	1	Monthly	297	1999 - Present
Social Good	Unemployment Rate	1	Monthly	900	1950 - Present
Traffic	Travel Volume	1	Monthly	531	1980 - Present




Figure 2: Visualization of Time-MMD, highlighting distinct characteristics across different domains.

each target variable, we manually select 1-2 highly relevant report series with guaranteed updates. For instance, the weekly influenza report² published by the Centers for Disease Control and Prevention of the United States is chosen as one of the report sources for the weekly influenza patients proportion of the United States. (2) Web Search Results: For each target variable, we design 2-3 highly relevant keywords used for web searching.

These two data sources complement each other: report data ensures higher relevance but cannot guarantee all-time coverage, while search results cover all times but are highly redundant; search results aggregate multiple data sources, while report data, usually in PDF or TXT format, cannot be extracted by searching. For frequency alignment, the textual data covers multiple frequencies, with search texts enabling daily precision and reports ranging from weekly to monthly and yearly intervals. Appendix D provides a complete list of keywords and reports source.

Data Collection: Searching and Crawling. For keyword web searching, we use the official Google API³ as the entry point. For each keyword, we collect the timestamp, source, title, and content from the top 10 results located each week from 1980 to present. For report data, we parse all available reports from each data source and preserve only plain-text paragraphs.

Data Preprocessing: Filtering, Disentangling, and Summarizing. To curate the collected raw text data, we introduce three key preprocessing steps: (1) Filtering to improve relevance; (2) Disentangling facts with predictions to mitigate data contamination; (3) Summarizing for better usability. Given the impracticality of performing these steps manually, we leverage the state-of-the-art LLM, Llama3-70B, to accomplish these tasks.

The prompt used for LLMs is detailed in Appendix F. We incorporate three specific strategies to alleviate the hallucination issue in LLMs and enhance preprocessing quality: (1) A concise introduction of the text. (2) Mandating the LLM to reference the data source, aiding constraint and

²<https://www.cdc.gov/flu/weekly/weeklyarchives2023-2024/week04.htm>

³<https://developers.google.com/custom-search/v1/overview>

Table 2: Statistics for text data. Relevance indicates the percentage of text data with relevant content. Coverage describes the proportion of numerical series data being covered by at least one fact. Details are provided in Appendix E. The statistics highlight the need for both reports and search data.

Source	Raw Tokens	Preprocessed Extracted Facts		Tokens	Preprocessed Extracted Prediction		
		Relevance(%)	Coverage (%)		Relevance(%)	Coverage (%)	Tokens
Report	17.4k	84.3 \pm 27.2	34.1 \pm 26.8	37.6 \pm 11.9	82.3 \pm 28.5	33.8 \pm 27.3	74.6 \pm 16.1
Search	54.4	16.8 \pm 3.3	90.7 \pm 13.5	38.4 \pm 4.0	16.0 \pm 4.3	81.9 \pm 17.2	62.8 \pm 3.3




Figure 3: Visualization of relevant report (a, left) and search (b, right) counts in Time-MMD over time. Text counts from both reports and searches increase over time. Domains receiving more attention, such as the economy, contain more available relevant text data..

verification. (3) Permitting the LLM to indicate ‘not available’ when relevance is uncertain, to avoid fabrication. Appendix G provides a showcase of the text before and after processing.

Data Quality & Property. Overall, Figure 3 visualizes the extracted fact count per month over time by domain. Note that the Agriculture report data is of high volume around 2020 and therefore produces a peak. We make the following observations: (a) The search data count exhibits a gradual increasing trend, benefiting from the development of the Internet; the report data count has stabilized in recent years, indicating that release schedule has become stable. (b) The sparsity of textual data varies across different domains, with high-profile fields often accompanied by richer textual data. These validate the extensibility and updatability of Time-MMD and highlight the importance of its coverage across 9 diverse domains.

We further validate the effectiveness of key steps in textual data construction:

(a) Data sources selection. We use *relevance* and *coverage* ratio to describe the percentage of relevant texts and the proportion of numerical series data being covered by at least one fact, respectively. As demonstrated by Table 2, report data exhibits higher relevance but lower coverage; search data display the opposite pattern. Thus, our combined usage serves as a comprehensive solution.

(b) Data preprocessing. Figure 4 provides word cloud visualizations of constructed text data in the health domain, respectively for extracted facts, extracted predictions, and discarded text. Recall that the target variable here is the influenza patients proportion. Highly relevant words such as "pandemic", "vaccine", and "flu" appear more frequently in the extracted facts; research paper-related words such as "edu", "mdpi", and "university" are more common in the discarded text. Besides, the prediction text primarily contains words describing future, such as "will" and "next". These validate the effectiveness of LLM filtering and disentangling. Furthermore, Table 2 presents a comparison of the token count before and after preprocessing. The substantial decrease validates that LLM summarizing improves usability. Appendix H provides the manual verification results on a subset of the data to further validate the effectiveness of preprocessing using LLMs.

2.3 Binary Time Stamps for Diverse TSA Tasks

To enable the Time-MMD for versatile and flexible use, we maintain binary timestamps for all numerical and text data, storing the manually verified start dates and end dates. Such binary stamps can be easily referred to while serving different tasks. For report text data, we manually verify the




Figure 4: Word cloud visualization for influenza patient proportion from the health domain. The discarded texts are identified by the LLM as irrelevant to the target variable. The results validate the effectiveness of LLM preprocessing.

timestamps based on the release notes or report contents. For search data, we integrate adjacent search results within each week and mark timestamps correspondingly.

2.4 Considerations for Fairness and Data Release

To consider fairness, we gather data from both the United States and African regions in the Health domain. As depicted in Figure 2, the numerical data of African region exhibits weaker periodicity. Figure 3 shows that the African region has considerably fewer reports compared to the United States. We encourage researchers to consider underrepresented groups when conducting multimodal TSA tasks.

To support various existing and potential novel TSA tasks, we include the following metadata when releasing Time-MMD: (1) Numerical Data: start & end time, target variable, other variables; (2) Text Data: start & end time, fact text (content & data source), prediction text (content & data source).

3 Multimodal Time-Series Forecasting Library: MM-TSFlib

In this section, we aim to illustrate the potential benefits of our Time-MMD for multimodal TSA by focusing on time-series forecasting (TSF), a fundamental TSA task. TSF involves predicting future events or trends based on historical time-series data. While most existing TSF methods primarily depend on numerical series, we aim to extend these unimodal TSF methods to multimodality. To achieve this, we contribute both formulating the multimodal TSF problem as well as introducing MM-TSFlib, the first comprehensive multimodal TSF library.

3.1 Problem Formulation

Conventional unimodal TSF models take a numerical series as input and output future values of some or all of its features. Let the input variable of the numerical series be denoted as $\mathbf{X} \in \mathbb{R}^{l \times d_{\text{in}}}$, where l is the length of the *lookback window* decided by domain experts and d_{in} is the feature dimension at each time step. The output variable of the forecasts generated of *horizon window* length h is denoted as $\mathbf{Y} \in \mathbb{R}^{h \times d_{\text{out}}}$, where d_{out} is the dimension of targets at each time step. For the sample at time step t , denoted as $(\mathbf{X}_t, \mathbf{Y}_t)$, $\mathbf{X}_t \in \mathbf{X} = [\mathbf{x}_{t-l+1}, \mathbf{x}_{t-l+2}, \dots, \mathbf{x}_t]$ and $\mathbf{Y}_t \in \mathbf{Y} = [\mathbf{y}_{t+1}, \mathbf{y}_{t+2}, \dots, \mathbf{y}_{t+h}]$. Thus, the unimodal TSF model parameterized by θ is denoted as $f_\theta : \mathcal{X} \rightarrow \mathcal{Y}$.

For multimodal TSF, the input variable of the textual series is also considered, which can be denoted as $\mathbf{S} \in \mathbb{R}^{k \times d_{\text{txt}}}$, where k is the lookback window length of the text series, independent of l , and d_{txt} is the feature dimension of the text. Although the text variable may have inconsistent feature dimensions, we slightly abuse the notation d_{txt} here for brevity. Thus, the multimodal TSF model parameterized by θ is denoted as $g_\theta : \mathcal{X} \times \mathcal{S} \rightarrow \mathcal{Y}$.

3.2 Pioneering Solution for Multimodal TSF

Multimodal Integration Framework. We propose a pioneering multimodal integration framework to extend existing unimodal TSF models to their multimodal versions. As illustrated in Figure 5, our framework features an end-to-end pipeline that integrates open-source language models with

various TSF models. Numerical and textual series are independently modeled using unimodal TSF models and LLMs with projection layers. These outputs are then combined using a learnable linear weighting mechanism to produce the final prediction. To reduce computational costs, we freeze the LLM parameters and train only the additional projection layers. We employ pooling layers to address the inconsistent dimensions of textual variables. This framework features an end-to-end training manner, with minimal overhead in trainable parameters.

Multimodal TSF Library. Building upon the multimodal dataset Time-MMD and integration framework, we present the first multimodal TSF library, named **MM-TSFLib**. MM-TSFLib supports multimodal extensions of over 20 unimodal TSF algorithms through 7 open-source (large) language models, including BERT (14), GPT-2 (49) (Small, Medium, Large, Extra-Large), Llama-2-7B (56), and Llama-3-8B⁴. We detail the implementations and language models in Appendix J.

MM-TSFLib is designed for ease of use with Time-MMD in multimodal TSA. Additionally, MM-TSFLib serves as a pilot toolkit for evaluating the multimodal extensibility of existing TSF models.




Figure 5: Overall structure of the MM-TSFLib. MM-TSFLib uses a model-agnostic multimodal integration framework that independently models numerical and textual series within an end-to-end training manner. MM-TSFLib slightly increases the number of trainable parameters, balancing effectiveness and efficiency.




Figure 6: Average normalized MSE results for each TSF backbone. Blue areas represent the performance gap between unimodal and multimodal results. The multimodal experiments significantly and consistently outperform corresponding unimodal ones. Detailed results are provided in Appendix O

4 Experiments for Multimodal TSF

Based on the constructed MM-TSFLib, we further conduct comprehensive experiments to demonstrate the superiority of multimodal TSF and the high quality of Time-MMD. We further investigate the impact of data domains, horizon window size and the text modeling method.

4.1 Experimental Setup

We adhere to the general setups following existing TSF literatures (65; 64; 47). Regarding the horizon window length, we consider a wider range **from short- to long-term TSF** tasks, with four different lengths for each dataset according to frequency. We conduct TSF tasks on **all 9 domains of Time-MMD**. We employ the widely-adopted mean squared error (MSE) as the evaluation metric. A higher MSE indicates a better performance.

We comprehensively consider **12 advanced unimodal TSF methods across 4 types** including: (1) Transformer-based: Transformer (57), Reformer (31), Informer (78), Autoformer (65), Crossformer (76), Non-stationary Transformer (43), FEDformer (80), iTransformer (41). (2) MLP-based: DLinear (72). (3) Agnostic: FiLM (79). (4) LLM-based: Time-LLM (24). Unless otherwise specified, we use GPT-2-Small as the LLM backbone in MM-TSFLib. We deployed two sets of experiments upon each TSF model, i.e., both unimodal and multimodal versions. More details about experimental setup are provided in Appendix K.

⁴<https://llama.meta.com/llama3>




Figure 7: Results of exploratory experiments. (a) Influence of data domains: inherent characteristics and text richness influence performance. (b) Influence of LLM backbones: unclear correlation between multimodal TSF performance and LLM natural language capabilities. (c) Influence of horizon window size: improvements via multimodality are robust to horizon window size.

4.2 Experimental Results

Our experiments aim to investigate the following five aspects. Appendix O provides detailed results.

Effectiveness of multimodal TSF. Figure 6 shows average MSE results for corresponding unimodal and multimodal versions of each TSF backbone. The multimodal versions consistently outperform corresponding unimodal versions. As detailed in Appendix O, the multimodal versions outperformed their unimodal versions in **95%** of over 1,000 experiments, reducing the mean squared error by over **15%** in average and up to **40%** in domains with rich textual data. Such consistent improvements fully validate the superiority of multimodal TSF and the effectiveness of our proposed multimodal framework in Section 3.

Furthermore, we observe that different TSF backbones benefit from multimodal extension to varying degrees. For example, the originally inferior Informer exhibits strong multimodal performance, which we attribute to its intrinsic design for modeling long-range dependencies that may benefit more from textual cues.

Additional experiments on text integration approach approaches are provided in Appendix L. We hope these results inspire more advanced multimodal TSF solutions.

Quality of Time-MMD dataset. Figure 6 shows that SOTA unimodal TSF models, such as iTransformer and PatchTST, maintain leading unimodal performance, validating the quality of Time-MMD’s numerical data. Moreover, multimodal extension consistently and significantly improves performance by incorporating textual data, confirming the quality of Time-MMD’s textual data.

Influence of data domains. Figure 7a shows the relationship between the relevant fact count and the reduced MSE via multimodal extension for each domain. The scatter plot generally illustrates a positive linear correlation, aligning with the innovation of integrating textual information. Besides, domain characteristics also influence multimodal performance, even with a similar fact count. For example, the security domain, focusing on disasters and emergency grants, exhibits higher unpredictability in the future thus benefits less from the historical textual information. This observation highlights the importance of Time-MMD’s coverage of 9 diverse domains.

Influence of the horizon window size. Figure 7c shows the relationship between horizon window size and the average MSE reduction for each domain. Overall, the MSE reduction is stable and promising across different horizon window size, from short term to long term. These results demonstrate that the effectiveness of multimodal TSF is robust to different forecast horizon requirements.

Influence of the text modeling method. Firstly, we varied the LLM backbone in MM-TSFlib and evaluated corresponding multimodal performance on health domain. As shown in Figure 7b, the choice of LLM backbone does not exhibit a significant correlation with multimodal TSF performance. For the GPT2 series, the scaling law is unclear for multimodal TSF, indicating no clear positive correlation between the parameter scale and TSF performance. Across different LLMs, multimodal TSF performance is relatively similar, even between the advanced Llama-3-8B and the earlier BERT. There might be three possible reasons: (1) Our proposed multimodal framework, although effective, still does not fully utilize the power of LLMs, particularly by only fine-tuning through projection layers. (2) Existing LLMs, pre-trained for natural language tasks, may not be directly suitable for multimodal TSF. (3) The embedding dimension of BERT is 768, much lower than the 4096 dimension

size of Llama-3-8B, thereby makes it easier to fine-tune an effective projection layer with limited training data.

Although LLMs are currently the mainstream approach for text modeling, they may not be suitable for certain scenarios, such as unfamiliar domains or unsupported languages. Therefore, we introduce Doc2Vec (32) a text embedding model trained from scratch, as an alternative for text modeling. The experimental results on three datasets across four prediction horizons are provided in Table 6. The results show that Doc2Vec is also effective but generally performs worse than BERT. Doc2Vec is further included in our MM-TSFlib, which will greatly enhance applicability to underrepresented languages and domains.

We provide additional experiments and discussions on multimodal modeling approaches in Appendix N, including the introduction of attention mechanisms (57) and the use of closed-source LLMs, such as GPT-3.5⁵.

In summary, all these observations suggest that there is significant room for improvement in the methodological research for multimodal TSF.

5 Potential Future Works

Beyond its efficacy in enhancing time-series forecasting accuracy through multimodality extension (Section 4), Time-MMD holds significant potential in advancing time series analysis across a wide spectrum. In the following section, we discuss how Time-MMD might transform conventional approaches, facilitate novel methodologies, and may broadly impact the time series analysis domain in future works.

Multimodal Time-Series Imputation. Missing values in time series data, caused by sensor failures, system instability, or privacy concerns, pose a significant challenge in analysis. Conventional time-series imputation (TSI) methods (8; 19; 3; 4) often overlook valuable information captured in textual formats alongside the numerical data. For instance, incident reports, weather conditions, and special events can provide crucial context for imputing missing data points in traffic time series, but current methods fail to effectively incorporate this information. By constructing missing values using existing toolkits (16), Time-MMD can directly serve as multimodal imputation datasets. Time-MMD enables the integration of textual contextual information with numerical time series data, opening new avenues for multimodal time series analysis and enhancing imputation accuracy.

Multimodal Time-Series Anomaly Detection. Detecting anomalies in time series data is crucial for identifying unusual patterns that may indicate faults, fraud, or other significant events (10; 5; 73). However, conventional anomaly detection methods (46; 7; 67; 52) are limited to expected pattern deviations in numerical data, overlooking valuable information in textual formats. For instance, news articles, social media posts, and market reports can provide critical context that influences financial market behavior and helps identify anomalies not evident from numerical data alone. Time-MMD enables a feasible way to support multimodal time series anomaly detection. For example, in influenza (health) dataset, if "flu outbreak" is mentioned in the text data, we can initially label the aligned numerical data as anomalous.

Multimodal Foundation Time-Series Models. The introduction of Time-MMD, a comprehensive text-numeric TS dataset, is expected to significantly advance multimodal-based TS methods, including the development of multimodal foundation TS models beyond existing unimodal models (29). Time-MMD will facilitate further exploration with more complex and informative prompts, enhancing the performance and capabilities of fine-tuning methods. Additionally, it might spur research and development of multimodal models specifically tailored for TSA, an area with limited exploration compared to other domains like vision and video generation (53; 60; 66; 17).

6 Ethics Statement

While collecting data from government and news websites, we rigorously adhered to ethical standards to ensure compliance with website policies and avoid potential conflicts of interest. Mindful of copyright and regional policies, we restricted our collection to content freely available without

⁵<https://platform.openai.com/docs/models/gpt-3-5-turbo>

premium access or subscription requirements. In collecting data from web searches, we used Google’s official API to ensure that the data strictly complied with ethical standards.

7 Conclusion

In this work, we propose Time-MMD, the first multi-domain multimodal time series dataset, and develop MM-TSFlib, the first multimodal time series forecasting library, which facilitates a pilot study for multimodal time series analysis on Time-MMD. We conduct extensive experiments to demonstrate the high quality of Time-MMD, the effectiveness of MM-TSFlib, and the superiority of integrating textual information for time series analysis. We envision that this work catalyzes the transformation of time series analysis from unimodal to multimodal by integrating natural language.

Acknowledgements: This paper was supported in part by the NSF (Expeditions CCF-1918770, CAREER IIS-2028586, CAREER IIS-2144338, Medium IIS-1955883, Medium IIS-2106961, Medium IIS-2403240, PIPP CCF-2200269), CDC MInD program, Meta faculty gifts, and funds/computing resources from Georgia Tech and GTRI. This work used the Delta GPU Supercomputer at NCSA of UIUC through allocation CIS240288 from the Advanced Cyberinfrastructure Coordination Ecosystem: Services & Support (ACCESS) program.

References

- [1] Abdul Fatir Ansari, Lorenzo Stella, Caner Turkmen, Xiyuan Zhang, Pedro Mercado, Huibin Shen, Oleksandr Shchur, Syama Sundar Rangapuram, Sebastian Pineda Arango, Shubham Kapoor, et al. Chronos: Learning the language of time series. *arXiv preprint arXiv:2403.07815*, 2024.
- [2] Mattia Atzeni, Amna Dridi, and Diego Reforgiato Recupero. Fine-grained sentiment analysis on financial microblogs and news headlines. In *Semantic Web Challenges: 4th SemWebEval Challenge at ESWC 2017, Portoroz, Slovenia, May 28-June 1, 2017, Revised Selected Papers*, pages 124–128. Springer, 2017.
- [3] Parikshit Bansal, Prathamesh Deshpande, and Sunita Sarawagi. Missing value imputation on multidimensional time series. *Proceedings of the VLDB Endowment*, 14(11):2533–2545, 2021.
- [4] Marin Biloš, Kashif Rasul, Anderson Schneider, Yuriy Nevmyvaka, and Stephan Günnemann. Modeling temporal data as continuous functions with stochastic process diffusion. In *International Conference on Machine Learning*, pages 2452–2470. PMLR, 2023.
- [5] Ane Blázquez-García, Angel Conde, Usue Mori, and Jose A Lozano. A review on outlier/anomaly detection in time series data. *ACM Computing Surveys (CSUR)*, 54(3):1–33, 2021.
- [6] Defu Cao, Furong Jia, Sercan O Arik, Tomas Pfister, Yixiang Zheng, Wen Ye, and Yan Liu. Tempo: Prompt-based generative pre-trained transformer for time series forecasting. *arXiv preprint arXiv:2310.04948*, 2023.
- [7] Ander Carreño, Iñaki Inza, and Jose A Lozano. Analyzing rare event, anomaly, novelty and outlier detection terms under the supervised classification framework. *Artificial Intelligence Review*, 53:3575–3594, 2020.
- [8] Zhengping Che, Sanjay Purushotham, Kyunghyun Cho, David Sontag, and Yan Liu. Recurrent neural networks for multivariate time series with missing values. *Scientific reports*, 8(1):6085, 2018.
- [9] Si-An Chen, Chun-Liang Li, Nate Yoder, Sercan O Arik, and Tomas Pfister. Tsmixer: An all-mlp architecture for time series forecasting. *arXiv preprint arXiv:2303.06053*, 2023.
- [10] Andrew A Cook, Göksel Mısırlı, and Zhong Fan. Anomaly detection for iot time-series data: A survey. *IEEE Internet of Things Journal*, 7(7):6481–6494, 2019.

- [11] Keith Cortis, André Freitas, Tobias Daudert, Manuela Huerlimann, Manel Zarrouk, Siegfried Handschuh, and Brian Davis. Semeval-2017 task 5: Fine-grained sentiment analysis on financial microblogs and news. In *Proceedings of the 11th international workshop on semantic evaluation (SemEval-2017)*, pages 519–535, 2017.
- [12] Abhimanyu Das, Weihao Kong, Andrew Leach, Shaan Mathur, Rajat Sen, and Rose Yu. Long-term forecasting with tide: Time-series dense encoder. *arXiv preprint arXiv:2304.08424*, 2023.
- [13] Abhimanyu Das, Weihao Kong, Rajat Sen, and Yichen Zhou. A decoder-only foundation model for time-series forecasting. In *Forty-first International Conference on Machine Learning*, 2024.
- [14] Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. Bert: Pre-training of deep bidirectional transformers for language understanding. *arXiv preprint arXiv:1810.04805*, 2018.
- [15] Zihan Dong, Xinyu Fan, and Zhiyuan Peng. Fnspid: A comprehensive financial news dataset in time series. *arXiv preprint arXiv:2402.06698*, 2024.
- [16] Wenjie Du, Jun Wang, Linglong Qian, Yiyuan Yang, Fanxing Liu, Zepu Wang, Zina Ibrahim, Haoxin Liu, Zhiyuan Zhao, Yingjie Zhou, et al. Tsi-bench: Benchmarking time series imputation. *arXiv preprint arXiv:2406.12747*, 2024.
- [17] Yilun Du, Sherry Yang, Bo Dai, Hanjun Dai, Ofir Nachum, Josh Tenenbaum, Dale Schuurmans, and Pieter Abbeel. Learning universal policies via text-guided video generation. *Advances in Neural Information Processing Systems*, 36, 2024.
- [18] Saeede Anbaee Farimani, Majid Vafaei Jahan, Amin Milani Fard, and Gholamreza Haffari. Leveraging latent economic concepts and sentiments in the news for market prediction. In *2021 IEEE 8th International Conference on Data Science and Advanced Analytics (DSAA)*, pages 1–10. IEEE, 2021.
- [19] Vincent Fortuin, Dmitry Baranchuk, Gunnar Rätsch, and Stephan Mandt. Gp-vae: Deep probabilistic time series imputation. In *International conference on artificial intelligence and statistics*, pages 1651–1661. PMLR, 2020.
- [20] Sebastian Gerard, Yu Zhao, and Josephine Sullivan. Wildfirespreadts: A dataset of multi-modal time series for wildfire spread prediction. *Advances in Neural Information Processing Systems*, 36:74515–74529, 2023.
- [21] Nate Gruver, Marc Finzi, Shikai Qiu, and Andrew G Wilson. Large language models are zero-shot time series forecasters. *Advances in Neural Information Processing Systems*, 36, 2024.
- [22] Albert Gu and Tri Dao. Mamba: Linear-time sequence modeling with selective state spaces. *arXiv preprint arXiv:2312.00752*, 2023.
- [23] Furong Jia, Kevin Wang, Yixiang Zheng, Defu Cao, and Yan Liu. Gpt4mts: Prompt-based large language model for multimodal time-series forecasting. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 38, pages 23343–23351, 2024.
- [24] Ming Jin, Shiyu Wang, Lintao Ma, Zhixuan Chu, James Y Zhang, Xiaoming Shi, Pin-Yu Chen, Yuxuan Liang, Yuan-Fang Li, Shirui Pan, et al. Time-llm: Time series forecasting by reprogramming large language models. In *The Twelfth International Conference on Learning Representations*, 2023.
- [25] Yiqiao Jin, Mohit Chandra, Gaurav Verma, Yibo Hu, Munmun De Choudhury, and Srijan Kumar. Better to ask in english: Cross-lingual evaluation of large language models for healthcare queries. In *Web Conference*, pages 2627–2638, 2024.
- [26] Yiqiao Jin, Minje Choi, Gaurav Verma, Jindong Wang, and Srijan Kumar. Mm-soc: Benchmarking multimodal large language models in social media platforms. In *ACL*, 2024.

- [27] Yiqiao Jin, Qinlin Zhao, Yiyang Wang, Hao Chen, Kaijie Zhu, Yijia Xiao, and Jindong Wang. Agentreview: Exploring peer review dynamics with llm agents. In *EMNLP*, 2024.
- [28] Harshavardhan Kamarthi, Lingkai Kong, Alexander Rodriguez, Chao Zhang, and B Aditya Prakash. When in doubt: Neural non-parametric uncertainty quantification for epidemic forecasting. *Advances in Neural Information Processing Systems*, 34:19796–19807, 2021.
- [29] Harshavardhan Kamarthi and B Aditya Prakash. Large pre-trained time series models for cross-domain time series analysis tasks. *Advances in Neural Information Processing Systems*, 2024.
- [30] Kunio Kashino, Gavin Smith, and Hiroshi Murase. Time-series active search for quick retrieval of audio and video. In *1999 IEEE International Conference on Acoustics, Speech, and Signal Processing. Proceedings. ICASSP99 (Cat. No. 99CH36258)*, volume 6, pages 2993–2996. IEEE, 1999.
- [31] Nikita Kitaev, Łukasz Kaiser, and Anselm Levskaya. Reformer: The efficient transformer. *arXiv preprint arXiv:2001.04451*, 2020.
- [32] Quoc Le and Tomas Mikolov. Distributed representations of sentences and documents. In *International conference on machine learning*, pages 1188–1196. PMLR, 2014.
- [33] Jun Li, Che Liu, Sibao Cheng, Rossella Arcucci, and Shenda Hong. Frozen language model helps ecg zero-shot learning. In *Medical Imaging with Deep Learning*, pages 402–415. PMLR, 2024.
- [34] Bryan Lim, Serkan Ö Arık, Nicolas Loeff, and Tomas Pfister. Temporal fusion transformers for interpretable multi-horizon time series forecasting. *International Journal of Forecasting*, 37(4):1748–1764, 2021.
- [35] Shengsheng Lin, Weiwei Lin, Wentai Wu, Feiyu Zhao, Ruichao Mo, and Haotong Zhang. Segrnn: Segment recurrent neural network for long-term time series forecasting. *arXiv preprint arXiv:2308.11200*, 2023.
- [36] Desheng Liu and Shanshan Cai. A spatial-temporal modeling approach to reconstructing land-cover change trajectories from multi-temporal satellite imagery. *Annals of the Association of American Geographers*, 102(6):1329–1347, 2012.
- [37] Haoxin Liu, Harshavardhan Kamarthi, Lingkai Kong, Zhiyuan Zhao, Chao Zhang, B Aditya Prakash, et al. Time-series forecasting for out-of-distribution generalization using invariant learning. In *Forty-first International Conference on Machine Learning*, 2024.
- [38] Haoxin Liu, Ziwei Zhang, Peng Cui, Yafeng Zhang, Qiang Cui, Jiashuo Liu, and Wenwu Zhu. Signed graph neural network with latent groups. In *Proceedings of the 27th ACM SIGKDD conference on knowledge discovery & data mining*, pages 1066–1075, 2021.
- [39] Haoxin Liu, Zhiyuan Zhao, Jindong Wang, Harshavardhan Kamarthi, and B Aditya Prakash. LSTPrompt: Large language models as zero-shot time series forecasters by long-short-term prompting. In *The 62nd Annual Meeting of the Association for Computational Linguistics*, 2024.
- [40] Shizhan Liu, Hang Yu, Cong Liao, Jianguo Li, Weiyao Lin, Alex X Liu, and Schahram Dustdar. Pyraformer: Low-complexity pyramidal attention for long-range time series modeling and forecasting. In *International conference on learning representations*, 2021.
- [41] Yong Liu, Tengge Hu, Haoran Zhang, Haixu Wu, Shiyu Wang, Lintao Ma, and Mingsheng Long. itransformer: Inverted transformers are effective for time series forecasting. In *The Twelfth International Conference on Learning Representations*, 2023.
- [42] Yong Liu, Chenyu Li, Jianmin Wang, and Mingsheng Long. Koopa: Learning non-stationary time series dynamics with koopman predictors. *Advances in Neural Information Processing Systems*, 36, 2024.

- [43] Yong Liu, Haixu Wu, Jianmin Wang, and Mingsheng Long. Non-stationary transformers: Exploring the stationarity in time series forecasting. *Advances in Neural Information Processing Systems*, 35:9881–9893, 2022.
- [44] Yong Liu, Haoran Zhang, Chenyu Li, Xiangdong Huang, Jianmin Wang, and Mingsheng Long. Timer: Generative pre-trained transformers are large time series models. In *Forty-first International Conference on Machine Learning*, 2024.
- [45] Yongtuo Liu, Sara Magliacane, Miltiadis Kofinas, and Stratis Gavves. Amortized equation discovery in hybrid dynamical systems. In *Forty-first International Conference on Machine Learning*, 2024.
- [46] Mohsin Munir, Shoaib Ahmed Siddiqui, Andreas Dengel, and Sheraz Ahmed. Deepant: A deep learning approach for unsupervised anomaly detection in time series. *Ieee Access*, 7:1991–2005, 2018.
- [47] Yuqi Nie, Nam H Nguyen, Phanwadee Sinthong, and Jayant Kalagnanam. A time series is worth 64 words: Long-term forecasting with transformers. *arXiv preprint arXiv:2211.14730*, 2022.
- [48] Xiao Ding Philippe Remy. Financial news dataset from bloomberg and reuters. <https://github.com/philipperemy/financial-news-dataset>, 2015.
- [49] Alec Radford, Jeffrey Wu, Rewon Child, David Luan, Dario Amodei, Ilya Sutskever, et al. Language models are unsupervised multitask learners. *OpenAI blog*, 1(8):9, 2019.
- [50] Filipe Rodrigues, Ioulia Markou, and Francisco C Pereira. Combining time-series and textual data for taxi demand prediction in event areas: A deep learning approach. *Information Fusion*, 49:120–129, 2019.
- [51] Omer Berat Sezer, Mehmet Ugur Gudelek, and Ahmet Murat Ozbayoglu. Financial time series forecasting with deep learning: A systematic literature review: 2005–2019. *Applied soft computing*, 90:106181, 2020.
- [52] Lifeng Shen, Zhuocong Li, and James Kwok. Timeseries anomaly detection using temporal hierarchical one-class network. *Advances in Neural Information Processing Systems*, 33:13016–13026, 2020.
- [53] Amanpreet Singh, Ronghang Hu, Vedanuj Goswami, Guillaume Couairon, Wojciech Galuba, Marcus Rohrbach, and Douwe Kiela. Flava: A foundational language and vision alignment model. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 15638–15650, 2022.
- [54] Ankur Sinha, Satishwar Kedas, Rishu Kumar, and Pekka Malo. Sentfin 1.0: Entity-aware sentiment analysis for financial news. *Journal of the Association for Information Science and Technology*, 73(9):1314–1335, 2022.
- [55] Anika Tabassum, Supriya Chinthavali, Varisara Tansakul, and B. Aditya Prakash. Actionable insights in urban multivariate time-series. *CIKM '21*, page 1774–1783, New York, NY, USA, 2021. Association for Computing Machinery.
- [56] Hugo Touvron, Thibaut Lavril, Gautier Izacard, Xavier Martinet, Marie-Anne Lachaux, Timothée Lacroix, Baptiste Rozière, Naman Goyal, Eric Hambro, Faisal Azhar, et al. Llama: Open and efficient foundation language models. *arXiv preprint arXiv:2302.13971*, 2023.
- [57] Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N Gomez, Łukasz Kaiser, and Illia Polosukhin. Attention is all you need. *Advances in neural information processing systems*, 30, 2017.
- [58] Huiqiang Wang, Jian Peng, Feihu Huang, Jince Wang, Junhui Chen, and Yifei Xiao. Micn: Multi-scale local and global context modeling for long-term series forecasting. In *The Eleventh International Conference on Learning Representations*, 2022.

- [59] Shiyu Wang, Haixu Wu, Xiaoming Shi, Tengge Hu, Huakun Luo, Lintao Ma, James Y Zhang, and Jun Zhou. Timemixer: Decomposable multiscale mixing for time series forecasting. *arXiv preprint arXiv:2405.14616*, 2024.
- [60] Wenhui Wang, Hangbo Bao, Li Dong, Johan Bjorck, Zhiliang Peng, Qiang Liu, Kriti Aggarwal, Owais Khan Mohammed, Saksham Singhal, Subhojit Som, et al. Image as a foreign language: Beit pretraining for vision and vision-language tasks. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 19175–19186, 2023.
- [61] Thomas Wolf, Lysandre Debut, Victor Sanh, Julien Chaumond, Clement Delangue, Anthony Moi, Pierric Cistac, Tim Rault, Rémi Louf, Morgan Funtowicz, et al. Huggingface’s transformers: State-of-the-art natural language processing. *arXiv preprint arXiv:1910.03771*, 2019.
- [62] G Woo, C Liu, D Sahoo, A Kumar, and S Hoi. Etsformer: Exponential smoothing transformers for time-series forecasting. *arxiv 2022. arXiv preprint arXiv:2202.01381*, 2022.
- [63] Gerald Woo, Chenghao Liu, Akshat Kumar, Caiming Xiong, Silvio Savarese, and Doyen Sahoo. Unified training of universal time series forecasting transformers. In *Forty-first International Conference on Machine Learning*, 2024.
- [64] Haixu Wu, Tengge Hu, Yong Liu, Hang Zhou, Jianmin Wang, and Mingsheng Long. Timesnet: Temporal 2d-variation modeling for general time series analysis. In *International Conference on Learning Representations*, 2023.
- [65] Haixu Wu, Jiehui Xu, Jianmin Wang, and Mingsheng Long. Autoformer: Decomposition transformers with auto-correlation for long-term series forecasting. *Advances in neural information processing systems*, 34:22419–22430, 2021.
- [66] Jinbo Xing, Menghan Xia, Yuxin Liu, Yuechen Zhang, Y He, H Liu, H Chen, X Cun, X Wang, Y Shan, et al. Make-your-video: Customized video generation using textual and structural guidance. *IEEE Transactions on Visualization and Computer Graphics*, 2024.
- [67] Jiehui Xu, Haixu Wu, Jianmin Wang, and Mingsheng Long. Anomaly transformer: Time series anomaly detection with association discrepancy. In *International Conference on Learning Representations*, 2021.
- [68] Yumo Xu and Shay B Cohen. Stock movement prediction from tweets and historical prices. In *Proceedings of the 56th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 1970–1979, 2018.
- [69] Yumo Xu and Shay B Cohen. Stock movement prediction from tweets and historical prices. In *Proceedings of the 56th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 1970–1979, 2018.
- [70] Kun Yi, Qi Zhang, Wei Fan, Shoujin Wang, Pengyang Wang, Hui He, Ning An, Defu Lian, Longbing Cao, and Zhendong Niu. Frequency-domain mlps are more effective learners in time series forecasting. *Advances in Neural Information Processing Systems*, 36, 2024.
- [71] Xinli Yu, Zheng Chen, Yuan Ling, Shujing Dong, Zongyi Liu, and Yanbin Lu. Temporal data meets llm—explainable financial time series forecasting. *arXiv preprint arXiv:2306.11025*, 2023.
- [72] Ailing Zeng, Muxi Chen, Lei Zhang, and Qiang Xu. Are transformers effective for time series forecasting? In *Proceedings of the AAAI conference on artificial intelligence*, volume 37, pages 11121–11128, 2023.
- [73] Kexin Zhang, Qingsong Wen, Chaoli Zhang, Rongyao Cai, Ming Jin, Yong Liu, James Y Zhang, Yuxuan Liang, Guansong Pang, Dongjin Song, et al. Self-supervised learning for time series analysis: Taxonomy, progress, and prospects. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 2024.
- [74] Tianping Zhang, Yizhuo Zhang, Wei Cao, Jiang Bian, Xiaohan Yi, Shun Zheng, and Jian Li. Less is more: Fast multivariate time series forecasting with light sampling-oriented mlp structures. *arXiv preprint arXiv:2207.01186*, 2022.

- [75] Xingxuan Zhang, Linjun Zhou, Renzhe Xu, Peng Cui, Zheyang Shen, and Haoxin Liu. Towards unsupervised domain generalization. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 4910–4920, 2022.
- [76] Yunhao Zhang and Junchi Yan. Crossformer: Transformer utilizing cross-dimension dependency for multivariate time series forecasting. In *The eleventh international conference on learning representations*, 2022.
- [77] Zhiyuan Zhao, Alexander Rodriguez, and B Aditya Prakash. Performative time-series forecasting. *arXiv preprint arXiv:2310.06077*, 2023.
- [78] Haoyi Zhou, Shanghang Zhang, Jieqi Peng, Shuai Zhang, Jianxin Li, Hui Xiong, and Wancai Zhang. Informer: Beyond efficient transformer for long sequence time-series forecasting. In *Proceedings of the AAAI conference on artificial intelligence*, volume 35, pages 11106–11115, 2021.
- [79] Tian Zhou, Ziqing Ma, Qingsong Wen, Liang Sun, Tao Yao, Wotao Yin, Rong Jin, et al. Film: Frequency improved legendre memory model for long-term time series forecasting. *Advances in Neural Information Processing Systems*, 35:12677–12690, 2022.
- [80] Tian Zhou, Ziqing Ma, Qingsong Wen, Xue Wang, Liang Sun, and Rong Jin. Fedformer: Frequency enhanced decomposed transformer for long-term series forecasting. In *International conference on machine learning*, pages 27268–27286. PMLR, 2022.

Appendix

A Additional Related Dataset Work

Existing datasets (2; 69; 50; 18; 54; 15; 6) primarily focus on stock analysis tasks in the finance domain. Other multi-modal datasets for traffic demand prediction (50), news impact prediction (23) and electrocardiogram classification (33) are proposed recently. However, these datasets still do not address the aforementioned gaps. Especially constructing datasets for diverse domains, such as agriculture and security, is more challenging but holds substantial real-world impact.

Besides, some datasets focus on combining images with numerical sequences, where the images are typically sourced from spatial domains, such as satellite imagery. These works (20; 36) are centered on spatiotemporal data analysis, whereas our work focuses on time series analysis.

Additional multi-modal datasets (23) for news impact prediction and electrocardiogram classification (33) are proposed recently. However, these datasets still do not address the aforementioned gaps. Especially constructing multimodal datasets for other domains, such as agriculture and security, is more challenging but holds substantial real-world impact.

B Limitations

Our work provides a comprehensive, high-quality multimodal time series dataset, but it still has limitations in terms of dataset diversity, as all the text data comes from English. We plan to extend Time-MMD to multilingual versions to better address diversity and leverage data from multiple languages (25). In addition, utilizing Time-MMD to support other multimodal time series analysis tasks, such as multivariate time series forecasting and time series classification, requires additional dataset curation work. Moreover, considering multimodal time series datasets with images and audio is a worthwhile direction for future work. In this work, we chose to include text for the following reasons: (1) Text and numerical series commonly coexist across multiple domains. (2) With the rise of LLMs, an increasing number of time-series methods are incorporating text based on large language models. Collecting data that further align images (20), audio (30) or graph (38) with text and numeric is meaningful but beyond the scope of this work. Lastly, this work focuses on time series forecasting and does not cover interacting dynamical systems or equation discovery (45), where language could provide insights into the qualitative behavior of the system.

Our multimodal time series forecasting library is built upon a simple integration framework that only uses a projection layer for fine-tuning LLMs. How to fine-tune LLMs more efficiently and effectively for time series analysis remains an interesting topic. Moreover, addressing the issue of distributional shift (75; 77; 37) in the text should also be considered in future work to build robust models.

Despite these limitations, we hope that the datasets and library we have constructed will facilitate broader research and applications in multimodal time series analysis, such as benchmark (26) and agent (27).

C Considerations for Domain Selection

Our domain selection is based on the following three considerations:

- Most of our selected domains are also widely used in existing time-series works (44; 1; 13; 63) but lack textual and numerical alignment. We chose these domains to facilitate comparison by researchers, and all these domains have significant real-world importance.
- The remaining three domains, including social good, agriculture, and security, are often overlooked by researchers. We included these domains to contribute to a more comprehensive dataset.
- Additionally, we also considered underrepresented groups, as reflected in the Health (Africa) dataset, compared with the Health (U.S.).

These domains exhibit differences in numerical data, textual data, and performances. To the best of our knowledge, our Time-MMD is the most domain-inclusive multimodal time-series dataset.

D Details of Data Sources and Variables by Domain

D.1 Agricultural Domain Dataset Sources

D.1.1 Numeric Data

The sequence raw data is sourced from the *retail broiler composite*, provided by the United States Department of Agriculture (USDA) Economic Research Service (ERS). The raw data can be accessed at the USDA ERS website⁶.

D.1.2 Text Data

Report: Source 1

- **Frequency:** Weekly
- **Name:** USDA Broiler Market News Report.
- **URL:** <https://usda.library.cornell.edu/concern/publications/5999n3392?locale=en&page=30#release-items>.
- **Description:** This report provides a summary of the national broiler market, with data consolidated from a number of linked reports. The release contains price information for broiler parts by U.S. region as well as in Mexico City, slaughter estimates, availability, and national fowl market and miscellaneous poultry information.

Report: Source 2

- **Frequency:** Daily
- **Name:** Daily National Broiler Market at a Glance.
- **URL:** <https://usda.library.cornell.edu/concern/publications/vx021f146?locale=en&page=142#release-items>.
- **Description:** This report summarizes the daily trends in the national broiler market, with commentary on price, demand and supply, market activity, and trade.

Search

- **Frequency:** 10 results per week.
- **Keywords:** United States Broiler Market; Chicken prices

D.2 Climate Domain Dataset Sources

The data is collected by the NOAA National Centers for Environmental Information and spans from 1895 to the present. It offers a comprehensive historical dataset for analyzing long-term climate trends and their implications for energy and environmental policies. The raw data can be accessed at the NOAA.

D.2.1 Numeric Data

Source :<https://www.drought.gov/historical-information?dataset=0&selectedDateUSDM=20240514>

D.2.2 Text Data

Report: Source 1

- **Frequency:** Month
- **Name:** Drought Report of NOAA National Centers for Environmental Information⁷.
- **Description:** Drought Highlights;National Overview;Regional Overview

⁶<https://www.ers.usda.gov/data-products/meat-price-spreads/documentation/>

⁷<https://www.ncei.noaa.gov/access/monitoring/monthly-report/drought/200602>

Report: Source 2

- **Frequency:** Month
- **Name:** National Climate Report of NOAA National Centers for Environmental Information⁸.
- **Description:** National Overview; Monthly and Seasonal Highlights;

Search

- **Frequency:** 10 results per week.
- **Keywords:** Drought and extreme weather; Precipitation

D.3 Economy Domain Dataset Sources

D.3.1 Numeric Data

Source : <https://www.census.gov/foreign-trade/balance/c0015.html>.

D.3.2 Text Data

Report: Source 1

- **Frequency:** Monthly
- **Name:** U.S. International Trade in Goods and Services ⁹.

Report: Source 2

- **Frequency:** Monthly
- **Name:** Advance Economic Indicators Report¹⁰.

Search

- **Frequency:** 10 results per week.
- **Keywords:** International Trade Balance

D.4 Energy Domain Dataset Sources

This subsection outlines the sources and characteristics of the sequence and raw textual data that form the energy component of the series-text dataset. This domain specifically focuses on gasoline prices, which are crucial for analyzing economic stability, consumer spending patterns, and the overall energy market dynamics.

D.4.1 Numeric Data

The sequence raw data comprises weekly statistics of U.S. gasoline prices, measured in dollars per gallon. Monitoring these prices provides insights into energy market fluctuations and helps predict changes in economic conditions and policy adjustments.

Source and Accessibility This data is collected by the U.S. Energy Information Administration (EIA) and spans from April 5, 1993, to the present, offering a comprehensive historical view of fuel economics. The raw data can be accessed at the EIA website¹¹.

D.4.2 Text Data

We collect raw text from various EIA publications and reports, as they provide contextual information and expert analysis related to the Numeric Data.

⁸<https://www.ncei.noaa.gov/access/monitoring/monthly-report/national/200602>

⁹https://www.census.gov/foreign-trade/Press-Release/ft900_index.html

¹⁰<https://www.census.gov/econ/indicators/current/index.html>

¹¹<https://www.eia.gov/petroleum/gasdiesel/>

Report Source 1: Annual Energy Outlook

- **Frequency:** Annually
- **Source:** U.S. Energy Information Administration¹².
- **Type:** Annual Energy Outlook reports
- **Coverage:** From 1979 to present

Report Source 2: Weekly Petroleum Status Report

- **Frequency:** Weekly
- **Source:** U.S. Energy Information Administration¹³.
- **Type:** Weekly Petroleum Status Reports
- **Coverage:** From 2011 to present

Search

- **Frequency:** 10 results per week.
- **Keywords:** Gasoline prices

D.5 Social Good Domain Dataset Sources

This subsection describes the sources and characteristics of the sequence and raw textual data that comprise the social good component of the series-text dataset. This domain focuses on unemployment statistics in the United States, segmented by various racial groups, reflecting the societal impact of economic disparities.

D.5.1 Numeric Data

The sequence raw data consists of monthly unemployment statistics for the United States, disaggregated by race. These statistics include data for the following racial groups: White, Black or African American, Asian, American Indian or Alaska Native, and Native Hawaiian or Other Pacific Islander.

This data is sourced from the United States Bureau of Labor Statistics (BLS) and spans from 1954 to the present, providing a long-term view of employment trends across different racial demographics. The raw data can be accessed at the BLS website¹⁴.

D.5.2 Text Data

The raw textual data is sourced from official reports which provide a comprehensive analysis of employment conditions, enriched with context that is crucial for understanding the nuances of unemployment across different racial groups.

Report Source 1

- **Frequency:** Monthly
- **Source:** U.S. Department of Labor, BLS¹⁵.
- **Type:** Official report on monthly employment situations
- **Coverage:** From 1994 to present
- **Quantity:** TBD
- **Relevance:** Directly relevant (Siamese)

¹²<https://www.eia.gov/outlooks/aeo/archive.php>

¹³<https://www.eia.gov/petroleum/supply/weekly/archive/>

¹⁴<https://data.bls.gov/home.htm>

¹⁵<https://www.bls.gov/bls/news-release/empsit.htm>

Report Source 2

- **Frequency:** Annually
- **Source:** U.S. Department of Labor, BLS¹⁶.
- **Type:** Official report on annual labor force characteristics by race and ethnicity
- **Coverage:** From 2015 to present
- **Quantity:** TBD
- **Relevance:** Directly relevant

Search

- **Frequency:** 10 results per week.
- **Keywords:** Unemployment rate; Employment situation

D.6 Public Health (United States) Domain Dataset Sources

This subsection outlines the sources and characteristics of the sequence and raw textual data that form the public health component of the series-text dataset. This domain specifically focuses on Influenza-Like Illness (ILI) statistics, which are crucial for monitoring seasonal and pandemic influenza outbreaks, guiding public health interventions, and planning healthcare resources effectively.

D.6.1 Numeric Data

The sequence raw data includes weekly statistics of ILI cases in the United States, sourced from the Centers for Disease Control and Prevention (CDC). ILI, representing a category of health conditions characterized by symptoms similar to those of influenza, serves as an important public health indicator for tracking influenza activity.

Source and Accessibility The data is collected and made available by the CDC's ILINet system, which has been tracking ILI cases since 1954. This extensive dataset allows for robust historical trend analysis and modeling of influenza patterns. The raw influenza patients data can be accessed at the CDC website¹⁷.

D.6.2 Text Data

Raw textual data in the public health domain includes official reports and related news, providing insights into the context and implications of the ILI statistics.

Report Source 1: Weekly U.S. Influenza Surveillance Report

- **Frequency:** Weekly
- **Source:** CDC¹⁸.
- **Type:** Official Weekly U.S. Influenza Surveillance Report
- **Coverage:** From 1999 to present

Report Source 2: Annual Flu Season Key Studies and News Reports

- **Frequency:** Annually
- **Source:** CDC¹⁹.
- **Type:** Key studies and news reports from the flu season
- **Coverage:** From 2019 to present

¹⁶<https://www.bls.gov/opub/reports/race-and-ethnicity/2015/home.htm>

¹⁷<https://gis.cdc.gov/grasp/fluview/fluportaldashboard.html>

¹⁸<https://www.cdc.gov/flu/weekly/pastreports.htm>

¹⁹<https://www.cdc.gov/flu/spotlights/index.htm>

Search

- **Frequency:** 10 results per week.
- **Keywords:** influenza; epidemic

D.7 Public Health (Africa) Domain Dataset Sources

D.7.1 Numeric Data

Source: World Health Organization ²⁰.

D.7.2 Text Data

Report: Source 1

- **Frequency:** Weekly
- **Name:** Influenza Virological Surveillance in the WHO African Region²¹.

Report: Source 2

- **Frequency:** Daily
- **Name:** Press of Africa Centres for Diseases Control²².
- **Description:**

Search

- **Frequency:** 10 results per week.
- **Keywords:** influenza; epidemic

D.8 Environment Domain Dataset Sources

D.8.1 Numeric Data

Source :<https://www.epa.gov/outdoor-air-quality-data>.

D.8.2 Text Data

Report: Source 1

- **Frequency:** Daily
- **Name:** Press about Air Quality of Department of Environmental Conservation of New York²³.

Report: Source 2

- **Frequency:** Daily
- **Name:** Article about air quality from New York National Broadcasting Company²⁴.

Search

- **Frequency:** 10 results per week.
- **Keywords:** New York air quality; New York air pollution

²⁰<https://app.powerbi.com/view?r=eyJrIjoibWVjM2Y4NjktMjJmMC00Y2NjLWFmOWQtODQONjZkNWMyZnMiwiwidCI6ImY2MTBjM>

²¹<https://www.afro.who.int/publications/influenza-virological-surveillance-who-african-region-epidemiological-week-39-2017>

²²<https://africacdc.org/news-item/africa-centres-for-diseases-control-and-prevention-launches-new-network>

²³<https://dec.ny.gov/news/press-releases/2021/5/dec-directs-pvs-chemical-solutions-inc-to-temporarily-cease-operations>

²⁴<https://www.nbcnewyork.com/tag/air-quality/>

D.9 Traffic Domain Dataset Sources

D.9.1 Numeric Data

Source :https://www.fhwa.dot.gov/policyinformation/travel_monitoring/tvtfaq.cfm.

D.9.2 Text Data

Report: Source 1

- **Frequency:** Weekly
- **Name:** Weekly Traffic Volume Report²⁵.
- The Traffic Volume report estimates the vehicle miles traveled (VMT) for interstate highways and how the total travel measured by VMT compares with travel that occurred in the same week of the previous year. The VMT is further split into passenger vehicle and truck components. The information gives new insights into the effect on traffic by storm activity, economic fluctuations, and other variables that could not be evaluated using the monthly report.

Search

- **Frequency:** 10 results per week.
- **Keywords:** Travel; Mobility

D.10 Security Domain Dataset Sources

D.10.1 Numeric Data

Source :<https://www.fema.gov/about/openfema/data-sets>.

D.10.2 Text Data

Report: Source 1

- **Frequency:** Uncertain
- **Name:** Billion-Dollar Weather and Climate Disasters²⁶.

Report: Source 2

- **Frequency:** Uncertain
- **Name:** Disaster and emergency declarations²⁷.

²⁵<https://datahub.transportation.gov/stories/s/Weekly-Traffic-Volume-Report/3g63-ik4u/>

²⁶<https://www.ncei.noaa.gov/access/billions/>

²⁷<https://www.fema.gov/about/openfema/data-sets>

E Statistics of Textual Data

We show statistics of collected report data and search data in Table 3 and Table 4. As the security reports themselves contain manually-written summaries, we didn't perform a LLM preprocessing on them.

Table 3: Detailed relevance ratio, coverage ratio and average token counts of LLM-processed report data on each domain. Token counts are collected by GPT2Tokenizer from Huggingface (61). As the collected security reports already contain hand-written summaries, we didn't perform a LLM preprocessing on them.

Domain	Raw Tokens	Fact			Prediction		
		Relevance(%)	Coverage(%)	Tokens	Relevance(%)	Coverage(%)	Tokens
Agriculture	3850.62	99.81	10.08	31.63	99.67	10.08	65.71
Climate	6501.57	98.64	61.09	36.71	98.64	61.09	69.99
Economy	62315.99	100.00	81.56	31.55	100.00	81.56	66.11
Energy	2335.59	100.00	23.94	45.97	100.00	23.94	92.53
Environment	495.24	90.38	1.80	44.75	75.00	1.53	82.33
Health (US)	3705.83	56.25	34.63	33.68	56.37	34.70	62.10
Health (AFR)	934.50	21.99	6.25	35.53	17.84	3.12	91.19
Security	10.78	-	-	-	-	-	-
Social Good	75878.06	93.53	37.56	18.17	93.53	37.56	47.58
Traffic	1002.40	97.82	50.28	60.69	100.00	50.28	93.68

Table 4: Detailed relevance ratio, coverage ratio and average token counts of LLM-processed search data on each domain

Domain	Raw Tokens	Fact			Prediction		
		Relevance(%)	Coverage(%)	Tokens	Relevance(%)	Coverage(%)	Tokens
Agriculture	54.26	10.58	81.85	35.39	10.96	61.09	64.76
Climate	54.33	12.49	98.59	35.55	8.17	57.06	67.91
Economy	53.16	17.25	94.09	35.97	16.86	91.25	63.05
Energy	55.45	19.12	78.43	36.64	16.85	76.81	57.62
Environment	55.61	16.20	95.97	37.48	15.92	83.22	66.60
Health(US)	56.46	19.49	99.71	48.40	21.12	97.77	60.05
Health(AFR)	56.66	22.20	99.39	41.99	21.48	95.20	63.95
Security	51.49	16.61	99.66	37.97	12.39	97.64	64.48
SocialGood	52.67	17.21	59.11	37.42	19.22	59.11	59.35
Traffic	53.42	17.12	100.00	36.89	17.46	99.62	60.53

F Prompt designed for LLM preprocessing

```
Instructions:
You are expert of {domain}.
Instructions:
1. Carefully read through the following {reportname} from {start_date} to {
  end_date} published by {Author}. Description: {Description}. Due to the
  length of the report, only an excerpt is provided below.
2. Filter the results to find information useful to making predictions about
  {keyword}. Discard any irrelevant information.
3. Summarize the useful filtered information into the following 5 parts:
  - Objective facts about the {keyword} situation
  - Analysis of the current situation
  - Predictions for the short-term future (next {short_term})
  - Predictions for the long-term future (next {long_term})
4. Format your output as follows:
  - Start each part with a tag indicating the type of content, using these
    tags:
    - #F# for objective facts
    - #In# for insights
    - #A# for analysis of current situation
    - #SP# for short-term predictions
    - #LP# for long-term predictions
  - Write each part concisely, using no more than 3 sentences.
  - For objective facts, cite the source at the end using a [Source] tag.
  - For your own opinions, use a [LLM] tag.
  - If no useful information can be found, simply write "NA" for that
    part.
Remember, focus only on information relevant to predicting {keyword}. You are
allowed to use "NA".
```

Figure 8: LLM prompt template used for preprocessing.

We show the prompt we use while doing LLM preprocessing in Figure 8. Corresponding keywords come from manually written domain-specific instructions.

G Showcase of Raw and Preprocessed Text

We show the pre and post-processing text content of a report in Climate domain in Figure 9. We observe that LLMs are able to accurately summarize the factual content from original report.

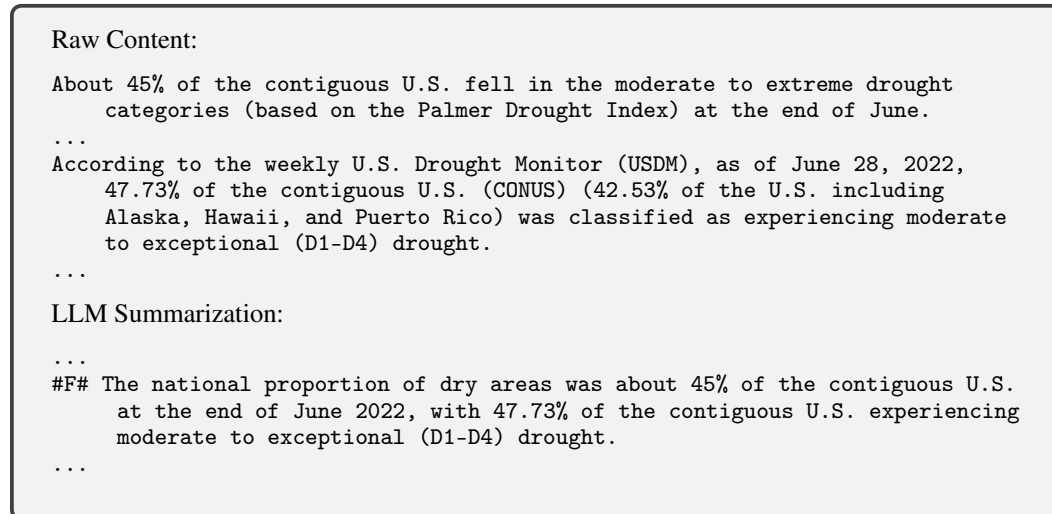


Figure 9: Show case of a text report before and after LLM preprocessing, sampled from Climate domain.

H Manual Verification of LLM Preprocessing

To validate the quality of LLM preprocessing, we manually inspected 100 text samples. Our observations are as follows:

- Among the 127 extracted facts, 8 were fabricated by the LLM.
- In the 8 fabricated instances, 5 were labeled with the data source as LLM, indicating that they can be filtered out. This results in a true hallucination rate of 3/127.
- Among the 52 text samples discarded by the LLM, 4 were manually identified as containing relevant information, yielding an error discard rate of 4/52.

I Case Study of Semantic Alignment

A case example of textual data explaining numerical data changes is shown in Figure 10 from the Health dataset. As the reports and search data are collected while ensuring strong time alignments, their textual content can provide clear illustrations on numerical series changes. For example, during Feb 2008, the health care numerical data experienced a significant upward trend, while the summarized report data states "..... The proportion of deaths attributed to pneumonia and influenza was above the epidemic threshold for the fourth consecutive week". Explicit alignments like this can be easily found in our data.




Figure 10: Case Study of Semantic Alignment from the Health Dataset.

J Details of MM-TSFlib

In terms of implementation, MM-TSFlib chooses to extend the widely used Time-Series Library (TSlib)²⁸, thus ensuring ease of use. For LLM invocation, MM-TSFlib utilizes the popular and active Hugging Face²⁹. For the projection layer, we use a multilayer perceptron (MLP) to keep it simple, apply instance normalization followed by adjusting the mean to the corresponding historical average to keep it rational. We tune the initialization values of the linear weighting mechanism to prevent overfitting or reduce learning difficulty. To use Time-MMD dataset for TSF task, MM-TSFlib constrains the latest end date of the input text sequence to be earlier than the latest end date of the input sequence, in order to avoid information leakage. Overall, MM-TSFlib supports multimodal extensions of over 20 unimodal TSF algorithms via 7 open-source LLM models.

J.1 List of Supported TSF

So far, our models supports following TSF models: TimeMixer (59), TSMixer (9), iTransformer (41), PatchTST (47), TimesNet (64), DLinear (72), LightTS (74), ETSformer (62), Non-stationary Transformer (43), FEDformer (80), Pyraformer (40), Autoformer (65), Informer (78), Reformer (31), Transformer (57), Mamba (22), SegRNN (35), Koopa (42), FreTS (70), TiDE (12), FiLM (79), MICN (58), Crossformer (76), TFT (34).

²⁸<https://github.com/thuml/Time-Series-Library>

²⁹<https://huggingface.co/models>

J.2 List of Supported LLM

For LLM, MM-TSFlib supports BERT (14), GPT-2 (49) (Small, Medium, Large, Extra-Large), Llama-2-7B (56), and Llama-3-8B³⁰.

K More Details of Experimental Setup

K.1 Time-Series Forecasting Backbones

We deployed two sets of experiments upon TSF models: (1) Transformer-based, including Transformer (57), Reformer (31), Informer (78), Autoformer (65), Crossformer (76), Non-stationary Transformer (43), FEDformer (80), iTransformer (41). (2) MLP-based: DLinear (72). (3) Agnostic: FiLM (79). (4) LLM-based: Time-LLM (24).

- **Transformer**, a classic sequence-to-sequence model basing on multi-head attention mechanism.
- **Reformer**, an computational-efficient Transformer with advancements in attention hashing and reversible residual layers
- **Informer**, an advanced Transformer designed to tackle long-term forecasting problem with sparse attention layers and self-attention distilling.
- **Autoformer**, a Transformer-based model that keeps encoder-decoder structure but alters attention computations by auto-correlation mechanism in order to benefit long-term forecasting.
- **Crossformer**, a multi-variate Transformer-based model that explicitly explores and utilizes cross-dimension dependencies.
- **Non-stationary Transformer**, a Transformer that is designed to capture non-stationarity patterns instead of temporal correlation.
- **FEDformer**, a Transformer that explicitly use Fourier decomposition results to enhance long-term forecasting ability.
- **iTransformer**, a inverted Transformer that tokenizes multivariate time-series upon each timestamps/
- **DLinear**, a linear model that performs forecasting by a direct regression upon historical time series with a one-layer linear model.
- **FiLM**, a model-agnostic method that introduces Legendre and Fourier projections to denoise series and approximate historical information.
- **Time-LLM**, a framework that integrates LLM for time-series forecasting by reprogramming input series and then aligning with text prototypes.

K.2 LLM Backbones

We use GPT-2 Small (49) for the majority of experiments, while other GPT-2 models and BERT, Llama-2, and Llama-3 are used in the ablation study.

- **BERT** a bidirectional transformer pre-trained on large text corpora with tasks like masked language modeling and next sentence prediction
- **GPT-2** an advanced language model that generates coherent and contextually relevant text by predicting subsequent words in a sentence, pre-trained on diverse internet text and capable of performing a variety of language tasks without task-specific fine-tuning
- **Llama-2** an accessible, open-source LLM designed to generate coherent and contextually relevant text by leveraging advanced transformer-based architecture
- **Llama-3** the latest iteration of the Llama model, offering enhanced text generation and comprehension abilities, further advancing the performance and versatility of its predecessors

³⁰<https://llama.meta.com/llama3>

K.3 Evaluation Metrics

We use mean squared error (MSE) as the evaluation metric for all experiments. The MSE is defined as the average of the squares of the errors. The error here is the difference between the forecasted value y_i and the actual value \hat{y}_i . That is:

$$\text{MSE} = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2 \quad (1)$$

K.4 Details of Implementation

All experiments are conducted using Pytorch on a 80GB NVIDIA A100 GPU. We follow the commonly adopted setup for defining the forecasting horizon window length, as outlined in prior works (65; 64; 47). Specifically, for daily reported datasets, the forecasting horizon windows are chosen from the set [48, 96, 192, 336], with a fixed lookback window size of 96 and a consistent label window size of 48 for the decoder. Similarly, for the weekly reported dataset, we employ forecasting horizon windows from [12, 24, 36, 48], with a fixed lookback window size of 36 and a constant label window size of 18 for the decoder. Besides, for the monthly reported dataset, we employ forecasting horizon windows from [6, 8, 10, 12], with a fixed lookback window size of 8 and a constant label window size of 4 for the decoder.

L Additional Experiments on Text Integration Approaches

The previously proposed multimodal integration approaches (24; 6) are designed specifically for prompt-based text rather than exogenous textual series. Specifically, the text used in these models serves as prompts, such as task descriptions, rather than our exogenous textual series. Prompts are static, fixed in length, highly templated, and do not provide additional information, which makes them significantly different from the textual series used in our Time-MMD. We choose Time-LLM (24), an advanced model in this category, as a representative. The experimental results are shown in Table 5.

Model	Health	Energy	Traffic
TimeLLM (Unimodal)	1.92	0.30	0.23
TimeLLM (Multimodal_TimeLLM)	2.51	0.39	0.27
TimeLLM (Multimodal_Our)	1.42	0.25	0.22

Table 5: Results of Additional Experiments on Multimodal Integration Approaches

The results show that the 'prompt as prefix' approach in TimeLLM, referred to as Multimodal_TimeLLM, is not suitable for handling textual series. Possible reasons include the longer length of the text series potentially overwhelming the numerical series information, among others. In contrast, our framework, although simple, proves to be effective."

M Additional Experiments on Text Modeling Approaches

Results are provided in Table 6.

N Additional Experiments on Multimodal Modeling Approaches

N.1 Using attention mechanisms

We use the output of the time series model to calculate attention scores for each token in the LLM output and perform weighted aggregation. For detailed implementation, please refer to our library. Results are provided in Table 7. This basic attention method has not yet demonstrated a clear advantage over the current simple combining approach. We look forward to future method designs.

Dataset	Health	Energy	Traffic
Reformer (Uni)	1.85	0.43	0.31
Reformer (Multi_Bert)	1.27	0.40	0.17
Reformer (Multi_Doc2Vec)	1.35	0.43	0.22
Informer (Uni)	1.53	0.33	0.28
Informer (Multi_Bert)	1.22	0.29	0.17
Informer (Multi_Doc2Vec)	1.32	0.31	0.21

Table 6: Results of Additional Text Modeling Solutions

Dataset	Health	Energy	Traffic
Reformer (Multi_Our)	1.27	0.41	0.17
Reformer (Multi_Att)	1.26	0.39	0.22
Informer (Multi_Our)	1.22	0.29	0.17
Informer (Multi_Att)	1.23	0.29	0.16

Table 7: Results of Using Attention Mechanisms for Multimodal TSF

N.2 Using closed-source LLMs

We first used closed-source LLMs, such as GPT-3.5, to generate text-based predictions, and then encoded them using BERT. For detailed implementation, please check our library. Results are provided in Table 8. We initially observed that fully frozen closed-source LLMs are less effective than fine-tuned open-source LLM.

Dataset	Health	Energy	Traffic
Reformer (Uni)	1.86	0.43	0.31
Reformer (Multi_Bert)	1.27	0.40	0.17
Reformer (Multi_GPT3.5)	1.38	0.46	0.19
Informer (Uni)	1.53	0.33	0.28
Informer (Multi_Bert)	1.22	0.28	0.17
Informer (Multi_GPT3.5)	1.33	0.32	0.23

Table 8: Results of Using Closed-Source LLMs for Multimodal TSF

O Detailed Results

Detailed results on all domains are shown. All results in tables are averaged results on three seeded runs. We report MSE as described in Appendix K.3. Quick reference: Agriculture: Table 9; Climate: Table 10; Economy: Table 11; Energy: Table 12; Environment: Table 13; Health(US): Table 14; Security: Table 15; Social Good: Table 16; Traffic: Table 17.

Horizon Window Length		6	8	10	12
Model	Modal				
FiLM	Uni	0.07	0.09	0.12	0.15
	Multi	0.06	0.09	0.10	0.14
DLinear	Uni	0.11	0.24	0.18	0.23
	Multi	0.09	0.17	0.13	0.16
Transformer	Uni	0.24	0.35	0.40	0.50
	Multi	0.14	0.16	0.24	0.24
Reformer	Uni	0.38	0.26	0.51	0.70
	Multi	0.19	0.18	0.27	0.32
Informer	Uni	0.51	0.61	0.66	0.80
	Multi	0.16	0.21	0.34	0.34
Autoformer	Uni	0.08	0.09	0.11	0.14
	Multi	0.08	0.09	0.10	0.13
FEDformer	Uni	0.06	0.08	0.10	0.13
	Multi	0.06	0.07	0.10	0.13
Nonstationary Transformer	Uni	0.06	0.07	0.10	0.12
	Multi	0.05	0.07	0.09	0.12
Crossformer	Uni	0.31	0.35	0.38	0.46
	Multi	0.11	0.16	0.20	0.26
PatchTST	Uni	0.06	0.08	0.10	0.13
	Multi	0.06	0.08	0.10	0.13
iTransformer	Uni	0.06	0.08	0.10	0.13
	Multi	0.06	0.08	0.10	0.13
Time-LLM	Uni	0.08	0.09	0.10	0.14
	Multi	0.06	0.08	0.10	0.13

Table 9: Agriculture

Horizon Window Length		6	8	10	12
Model	Modal				
FiLM	Uni	1.42	1.39	1.40	1.40
	Multi	1.15	1.15	1.14	1.17
DLinear	Uni	1.35	1.41	1.36	1.36
	Multi	1.06	1.05	1.07	1.08
Transformer	Uni	1.04	1.14	1.12	1.11
	Multi	0.97	1.01	1.00	1.00
Reformer	Uni	1.24	1.06	1.13	1.16
	Multi	0.97	0.95	0.94	0.98
Informer	Uni	1.08	1.11	1.08	1.07
	Multi	1.04	1.03	1.04	1.02
Autoformer	Uni	1.30	1.24	1.28	1.25
	Multi	1.08	1.02	1.05	1.05
FEDformer	Uni	1.32	1.36	1.28	1.27
	Multi	0.98	1.00	1.03	1.02
Nonstationary Transformer	Uni	1.30	1.32	1.36	1.32
	Multi	1.00	1.02	1.02	1.01
Crossformer	Uni	1.12	1.10	1.12	1.10
	Multi	1.00	0.99	1.00	1.01
PatchTST	Uni	1.36	1.33	1.27	1.28
	Multi	0.99	1.01	1.04	1.06
iTransformer	Uni	1.16	1.23	1.24	1.22
	Multi	0.99	1.01	1.04	1.06
Time-LLM	Uni	1.36	1.26	1.27	1.27
	Multi	0.99	1.01	1.04	1.07

Table 10: Climate

Horizon Window Length		6	8	10	12
Model	Modal				
FiLM	Uni	0.04	0.04	0.03	0.04
	Multi	0.03	0.03	0.03	0.03
DLinear	Uni	0.08	0.18	0.10	0.11
	Multi	0.04	0.10	0.03	0.03
Transformer	Uni	0.60	0.69	1.19	1.12
	Multi	0.14	0.31	0.23	0.36
Reformer	Uni	0.65	0.76	1.02	0.55
	Multi	0.22	0.21	0.32	0.28
Informer	Uni	1.59	1.82	1.74	2.08
	Multi	0.28	0.36	0.49	0.50
Autoformer	Uni	0.08	0.10	0.07	0.06
	Multi	0.07	0.08	0.07	0.07
FEDformer	Uni	0.05	0.06	0.05	0.06
	Multi	0.04	0.03	0.04	0.03
Nonstationary Transformer	Uni	0.02	0.02	0.02	0.03
	Multi	0.02	0.02	0.02	0.02
Crossformer	Uni	0.95	1.06	1.22	1.24
	Multi	0.17	0.16	0.27	0.29
PatchTST	Uni	0.02	0.02	0.02	0.02
	Multi	0.02	0.02	0.02	0.02
iTransformer	Uni	0.02	0.02	0.02	0.02
	Multi	0.01	0.02	0.02	0.02
Time-LLM	Uni	0.06	0.08	0.02	0.05
	Multi	0.02	0.02	0.02	0.02

Table 11: Economy

Horizon Window Length		12	24	36	48
Model	Modal				
FiLM	Uni	0.21	0.30	0.37	0.49
	Multi	0.17	0.28	0.36	0.48
DLinear	Uni	0.26	0.32	0.39	0.50
	Multi	0.22	0.29	0.36	0.47
Transformer	Uni	0.18	0.26	0.36	0.44
	Multi	0.13	0.22	0.32	0.42
Reformer	Uni	0.28	0.38	0.49	0.57
	Multi	0.25	0.38	0.43	0.54
Informer	Uni	0.18	0.29	0.35	0.48
	Multi	0.15	0.24	0.32	0.44
Autoformer	Uni	0.18	0.31	0.34	0.47
	Multi	0.16	0.27	0.32	0.45
FEDformer	Uni	0.11	0.24	0.34	0.45
	Multi	0.09	0.21	0.32	0.44
Nonstationary Transformer	Uni	0.11	0.21	0.34	0.48
	Multi	0.10	0.20	0.28	0.46
Crossformer	Uni	0.14	0.29	0.36	0.41
	Multi	0.13	0.26	0.36	0.41
PatchTST	Uni	0.10	0.21	0.30	0.42
	Multi	0.10	0.21	0.29	0.41
iTransformer	Uni	0.10	0.21	0.30	0.42
	Multi	0.09	0.19	0.29	0.41
Time-LLM	Uni	0.16	0.27	0.31	0.45
	Multi	0.10	0.20	0.29	0.41

Table 12: Energy

Horizon Window Length		48	96	192	336
Model	Modal				
FiLM	Uni	0.32	0.35	0.35	0.32
	Multi	0.30	0.32	0.32	0.30
DLinear	Uni	0.41	0.57	0.73	0.59
	Multi	0.32	0.40	0.46	0.42
Transformer	Uni	0.32	0.32	0.48	0.44
	Multi	0.59	0.61	0.70	0.32
Reformer	Uni	0.39	0.45	0.51	0.48
	Multi	0.29	0.35	0.36	0.32
Informer	Uni	0.39	0.42	0.46	0.48
	Multi	0.31	0.33	0.39	0.34
Autoformer	Uni	0.43	0.36	0.52	0.37
	Multi	0.35	0.35	0.35	0.34
FEDformer	Uni	0.36	0.43	0.42	0.35
	Multi	0.30	0.34	0.34	0.33
Nonstationary Transformer	Uni	0.31	0.39	0.43	0.38
	Multi	0.29	0.31	0.32	0.30
Crossformer	Uni	0.34	0.33	0.73	0.53
	Multi	0.29	0.30	0.36	0.36
PatchTST	Uni	0.35	0.38	0.36	0.32
	Multi	0.31	0.32	0.32	0.30
iTransformer	Uni	0.28	0.29	0.30	0.28
	Multi	0.28	0.29	0.29	0.27
Time-LLM	Uni	0.38	0.37	0.45	0.33
	Multi	0.29	0.30	0.31	0.28

Table 13: Environment

Horizon Window Length		12	24	36	48
Model	Modal				
FiLM	Uni	2.53	2.59	2.46	2.38
	Multi	1.67	1.83	1.80	1.81
DLinear	Uni	2.37	2.61	2.50	2.48
	Multi	1.62	1.67	1.68	1.78
Transformer	Uni	1.22	1.56	1.43	1.55
	Multi	0.93	1.34	1.26	1.29
Reformer	Uni	1.63	1.99	1.91	1.90
	Multi	1.06	1.30	1.33	1.39
Informer	Uni	1.24	1.61	1.61	1.67
	Multi	0.98	1.23	1.28	1.40
Autoformer	Uni	1.99	2.25	2.26	2.39
	Multi	1.43	1.74	1.76	1.69
FEDformer	Uni	1.08	1.58	1.69	1.76
	Multi	0.92	1.25	1.36	1.42
Nonstationary Transformer	Uni	1.19	1.68	1.91	2.02
	Multi	0.94	1.14	1.17	1.30
Crossformer	Uni	1.45	1.57	1.62	1.65
	Multi	1.01	1.29	1.28	1.37
PatchTST	Uni	1.23	1.63	1.78	1.86
	Multi	0.98	1.27	1.49	1.60
iTransformer	Uni	1.14	1.62	1.84	1.89
	Multi	0.97	1.38	1.71	1.72
Time-LLM	Uni	1.60	1.94	1.95	2.17
	Multi	0.98	1.36	1.65	1.69

Table 14: Health(US)

Horizon Window Length		6	8	10	12
Model	Modal				
FiLM	Uni	127.92	122.53	125.10	127.06
	Multi	105.92	107.25	109.80	110.94
DLinear	Uni	107.12	111.72	114.42	116.65
	Multi	104.28	106.76	108.53	109.90
Transformer	Uni	128.42	130.78	131.65	133.04
	Multi	122.54	126.66	125.95	127.46
Reformer	Uni	124.48	125.68	126.51	127.26
	Multi	122.11	123.85	124.54	125.01
Informer	Uni	129.43	131.43	132.84	133.54
	Multi	121.81	124.24	126.98	126.73
Autoformer	Uni	125.67	119.82	119.72	126.67
	Multi	107.95	114.98	113.34	116.33
FEDformer	Uni	114.98	112.14	117.92	121.90
	Multi	109.28	109.55	115.99	117.22
Nonstationary Transformer	Uni	105.77	106.90	109.72	111.16
	Multi	104.15	105.99	110.49	110.69
Crossformer	Uni	124.70	124.97	127.49	128.47
	Multi	120.99	122.40	124.34	125.50
PatchTST	Uni	107.31	111.42	114.90	113.11
	Multi	108.75	111.26	114.61	114.67
iTransformer	Uni	113.20	115.26	116.32	117.59
	Multi	114.58	116.49	116.25	116.31
Time-LLM	Uni	112.96	116.53	116.46	119.60
	Multi	110.86	114.54	115.56	115.63

Table 15: Security

Horizon Window Length		6	8	10	12
Model	Modal				
FiLM	Uni	1.03	1.07	1.16	1.24
	Multi	0.87	0.97	1.05	1.06
DLinear	Uni	1.01	1.13	1.20	1.24
	Multi	0.91	1.02	1.14	1.15
Transformer	Uni	0.78	0.93	0.98	1.02
	Multi	0.73	0.84	0.91	0.90
Reformer	Uni	0.81	0.90	0.96	1.05
	Multi	0.80	0.85	0.95	1.03
Informer	Uni	0.77	0.86	0.93	0.94
	Multi	0.74	0.79	0.86	0.85
Autoformer	Uni	0.89	1.05	1.05	1.17
	Multi	0.88	0.96	1.04	1.12
FEDformer	Uni	0.81	0.92	1.02	1.11
	Multi	0.80	0.91	0.98	1.04
Nonstationary Transformer	Uni	0.85	1.17	1.29	1.31
	Multi	0.78	1.08	1.16	1.23
Crossformer	Uni	0.79	0.83	0.97	0.94
	Multi	0.75	0.83	0.89	0.93
PatchTST	Uni	0.88	0.99	1.07	1.18
	Multi	0.82	0.92	1.02	1.10
iTransformer	Uni	1.03	1.11	1.21	1.22
	Multi	0.90	1.08	1.12	1.15
Time-LLM	Uni	0.90	1.00	1.07	1.19
	Multi	0.90	0.99	1.08	1.10

Table 16: SocialGood

Horizon Window Length		6	8	10	12
Model	Modal				
FiLM	Uni	0.28	0.28	0.27	0.31
	Multi	0.25	0.24	0.24	0.29
DLinear	Uni	0.35	0.42	0.34	0.38
	Multi	0.28	0.30	0.25	0.31
Transformer	Uni	0.29	0.29	0.29	0.29
	Multi	0.15	0.15	0.16	0.18
Reformer	Uni	0.31	0.28	0.31	0.33
	Multi	0.16	0.16	0.17	0.19
Informer	Uni	0.26	0.27	0.27	0.30
	Multi	0.15	0.16	0.17	0.18
Autoformer	Uni	0.20	0.20	0.20	0.29
	Multi	0.19	0.19	0.20	0.27
FEDformer	Uni	0.21	0.21	0.20	0.27
	Multi	0.16	0.16	0.17	0.23
Nonstationary Transformer	Uni	0.19	0.21	0.20	0.27
	Multi	0.19	0.20	0.19	0.25
Crossformer	Uni	0.25	0.24	0.25	0.26
	Multi	0.17	0.16	0.17	0.19
PatchTST	Uni	0.21	0.21	0.22	0.28
	Multi	0.18	0.19	0.20	0.27
iTransformer	Uni	0.20	0.21	0.22	0.28
	Multi	0.19	0.20	0.21	0.27
Time-LLM	Uni	0.21	0.21	0.21	0.27
	Multi	0.19	0.20	0.20	0.27

Table 17: Traffic

NeurIPS Paper Checklist

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper's contributions and scope?

Answer: [\[Yes\]](#)

Justification: In the introduction and abstract, we clearly outline the main contributions, scope, and experimental support of the paper. The claims align with the research results, and we also include a discussion of the limitations.

Guidelines:

- The answer NA means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A No or NA answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [\[Yes\]](#)

Justification: Appendix B

Guidelines:

- The answer NA means that the paper has no limitation while the answer No means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate "Limitations" section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren't acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. Theory Assumptions and Proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [N/A]

Justification: NA

Guidelines:

- The answer NA means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental Result Reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [Yes]

Justification: Section 4 and Appendix K

Guidelines:

- The answer NA means that the paper does not include experiments.
- If the paper includes experiments, a No answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
 - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
 - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
 - (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [Yes]

Justification: Abstract

Guidelines:

- The answer NA means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://nips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so “No” is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://nips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. Experimental Setting/Details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer, etc.) necessary to understand the results?

Answer: [Yes]

Justification: Section 4 and Appendix K

Guidelines:

- The answer NA means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

7. Experiment Statistical Significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [Yes]

Justification: Section 4 and Appendix K

Guidelines:

- The answer NA means that the paper does not include experiments.
- The authors should answer "Yes" if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)

- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g. negative error rates).
- If error bars are reported in tables or plots, The authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments Compute Resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [Yes]

Justification: Section 4 and Appendix K

Guidelines:

- The answer NA means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

9. Code Of Ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

Answer: [Yes]

Justification: Section 6

Guidelines:

- The answer NA means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer No, they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. Broader Impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [Yes]

Justification: Section 5

Guidelines:

- The answer NA means that there is no societal impact of the work performed.
- If the authors answer NA or No, they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.

- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pretrained language models, image generators, or scraped datasets)?

Answer: [N/A]

Justification: NA

Guidelines:

- The answer NA means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [Yes]

Justification: Appendix V

Guidelines:

- The answer NA means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, paperswithcode.com/datasets has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.

- If this information is not available online, the authors are encouraged to reach out to the asset’s creators.

13. **New Assets**

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [\[Yes\]](#)

Justification: Section 2 and Appendix D

Guidelines:

- The answer NA means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. **Crowdsourcing and Research with Human Subjects**

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: [\[N/A\]](#)

Justification: NA

Guidelines:

- The answer NA means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

15. **Institutional Review Board (IRB) Approvals or Equivalent for Research with Human Subjects**

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

Answer: [\[N/A\]](#)

Justification: NA

Guidelines:

- The answer NA means that the paper does not involve crowdsourcing nor research with human subjects.
- Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper.
- We recognize that the procedures for this may vary significantly between institutions and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines for their institution.
- For initial submissions, do not include any information that would break anonymity (if applicable), such as the institution conducting the review.

Additional Supplementary Materials for Time-MMD Dataset

P Dataset Documentations

The dataset is provided in csv format.

Each numerical data point contains the following attribute:

- start time
- end time
- values of target variables
- values of other variables

Each textual data point contains the following attribute:

- start time
- end time
- extracted fact text (content & data source)
- extracted prediction text (content & data source)

Q Intended Uses

Time-MMD is intended for researchers in machine learning and related fields to develop novel methods for multimodal time-series analysis. We also aim to help developers to train or fine-tune multi-modal foundation time-series models using our dataset.

R Accessing

Time-MMD is publicly available in <https://github.com/AdityaLab/Time-MMD>. A detailed demo on how the dataset can be read and used can be found at <https://github.com/AdityaLab/MM-TSFlib>. A Croissant card of Time-MMD can be found in Google Drive.

S Statement

The authors will bear all responsibility in case of violation of rights.

T Hosting and Maintenance Plan

Time-MMD is hosted and version-tracked via GitHub. It will be permanently available under the link <https://github.com/AdityaLab/Time-MMD>. The download link of all the datasets can be found in the GitHub repository. We plan to regularly update the cutoff time of the data in Time-MMD at a frequency of every 6 months.

U Time-MMD Community

Time-MMD is a community-driven and open-source initiative. We are committed and have resources to maintain and actively develop Time-MMD in the future. We plan to grow Time-MMD to include more languages, domains, and multimodal time series analysis tasks. We welcome external contributors.

V Licensing

We license our dataset by ODC-By v1.0. While parsing the data, we follow the terms of both Google Search API³¹ and WebScraper³². We further make sure that all data sources obtained in Time-MMD are in public domain and authorized for non-commercial distribution. The authors will bear all responsibility in case of violation of rights.

³¹<https://developers.google.com/terms>

³²<https://webscraper.io/extension-privacy-policy>