

Chimera: Effectively Modeling Multivariate Time Series with 2-Dimensional State Space Models

Ali Behrouz^{*}, Michele Santacatterina[†], and Ramin Zabih^{*}

^{*}Department of Computer Science, Cornell University

[†]NYU Grossman School of Medicine

ab2947@cornell.edu, santam13@nyu.edu, rdz@cs.cornell.edu

Abstract

Modeling multivariate time series is a well-established problem with a wide range of applications from healthcare to financial markets. It, however, is extremely challenging as it requires methods to (1) have high expressive power of representing complicated dependencies along the time axis to capture both long-term progression and seasonal patterns, (2) capture the inter-variate dependencies when it is informative, (3) dynamically model the dependencies of variate and time dimensions, and (4) have efficient training and inference for very long sequences. Traditional State Space Models (SSMs) are classical approaches for univariate time series modeling due to their simplicity and expressive power to represent *linear* dependencies. They, however, have fundamentally limited expressive power to capture non-linear dependencies, are slow in practice, and fail to model the inter-variate information flow. Despite recent attempts to improve the expressive power of SSMs by using deep structured SSMs, the existing methods are either limited to univariate time series, fail to model complex patterns (e.g., seasonal patterns), fail to dynamically model the dependencies of variate and time dimensions, and/or are input-independent. We present Chimera, an expressive variation of the 2-dimensional SSMs with careful design of parameters to maintain high expressive power while keeping the training complexity linear. Using two SSM heads with different discretization processes and input-dependent parameters, Chimera is provably able to learn long-term progression, seasonal patterns, and desirable dynamic autoregressive processes. To improve the efficiency of complex 2D recurrence, we present a fast training using a new 2-dimensional parallel selective scan. We further present and discuss 2-dimensional Mamba and Mamba-2 as the spacial cases of our 2D SSM. Our experimental evaluation shows the superior performance of Chimera on extensive and diverse benchmarks, including ECG and speech time series classification, long-term and short-term time series forecasting, and time series anomaly detection.

1 Introduction

Modeling time series is a well-established problem with a wide range of applications from healthcare (Behrouz, Delavari, and Hashemi 2024; Ivanov et al. 1999) to financial markets (Gajamannage, Park, and Jayathilake 2023; Pincus and Kalman 2004) and energy management (H. Zhou et al. 2021). The complex nature of time series data, its diverse domains of applicability, and its broad range of tasks (e.g., classification (Behrouz, Delavari, and Hashemi 2024; Wagner et al. 2020), imputation (Luo and X. Wang 2024; H. Wu, Hu, et al. 2023), anomaly detection (Behrouz, Delavari, and Hashemi 2024; Su et al. 2019), and forecasting (H. Zhou et al. 2021)), however, raise fundamental challenges to design effective and generalizable models: (1) The higher-order, seasonal, and long-term patterns in time series require an effective model to be able to expressively capture complex and autoregressive dependencies; (2) In the presence of multiple variates of time series, an effective model need to capture the complex dynamics of the dependencies between time and variate axes. More specifically, most existing multivariate models seem to suffer from overfitting especially when the target time series is not correlated with other covariates (Zeng et al. 2023a). Accordingly, an effective model needs to adaptively learn to select (resp. filter) informative (resp. irrelevant) variates; (3) The diverse set of domains and tasks requires effective models to be free of manual pre-processing and domain knowledge and instead adaptively learn them; and (4) Due to the processing of very long sequences, effective methods need efficient training and inference.

Classical methods (e.g., State Space Models (Aoki 2013; Harvey 1990), ARIMA (Bartholomew 1971), SARIMA (Bender and Simonovic 1994), Exponential Smoothing (ETS) (Winters 1960)) require manual data preprocessing and model selection,

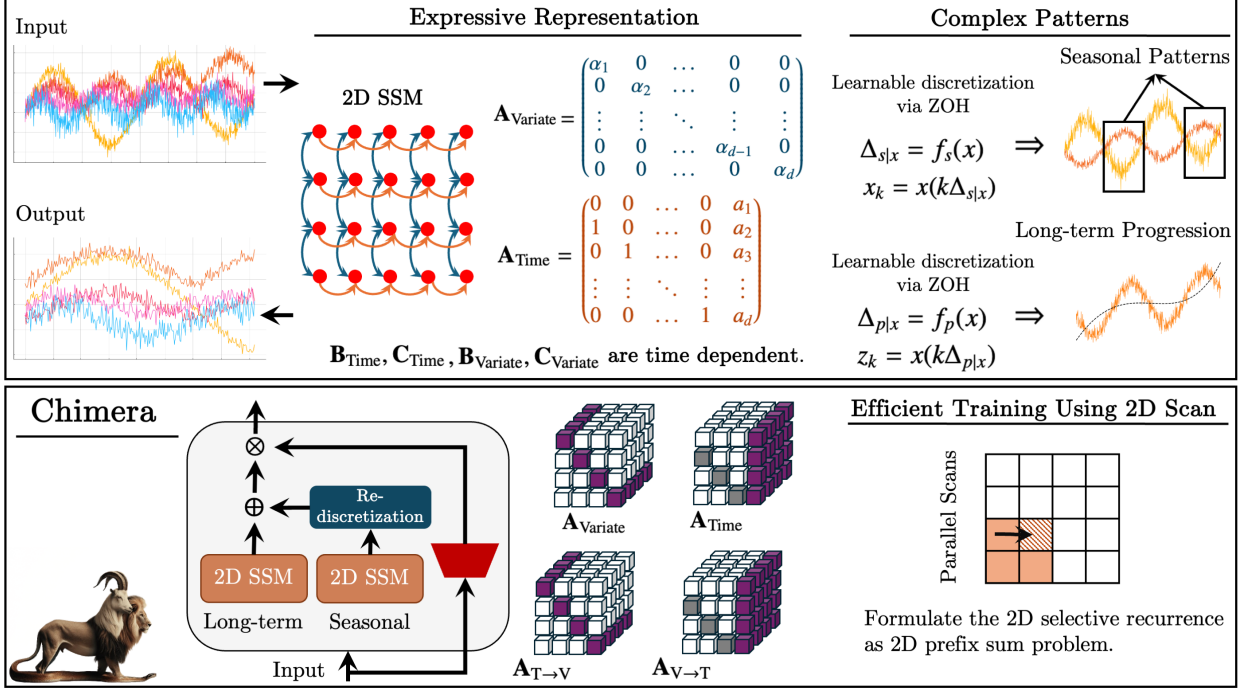


Figure 1: **The Overview of Contributions and Architecture of Chimera.** We present a 2-dimensional SSM with careful and expressive parameterization. It uses different learnable discretization processes to learn seasonal and long-term progression patterns, and leverages a parallelizable and fast training process by re-formulating the 2D input dependent recurrence as a 2D prefix sum problem.

and often are not able to capture complex *non-linear* dynamics. The raise of deep learning methods and more specifically Transformers (Vaswani et al. 2017) has led to significant research efforts to address the limitation of classical methods and develop effective deep models (Z. Chen et al. 2023; Kitaev, Kaiser, and Levskaya 2020; Lim and Zohren 2021; S. Liu et al. 2021; Yong Liu, Hu, et al. 2024; Woo et al. 2022; H. Wu, J. Wu, et al. 2022; H. Wu, Xu, et al. 2021; Y. Zhang and Yan 2023; T. Zhou, Z. Ma, Wen, X. Wang, et al. 2022). Unfortunately, most existing deep models struggle to achieve all the above four criteria. The main body of research in this direction has focused on designing attention modules that use the special traits of time series (Woo et al. 2022; H. Wu, Xu, et al. 2021). However, the inherent permutation equivariance of attentions contradicts the causal nature of time series and often results in suboptimal performance compared to simple linear methods (Zeng et al. 2023a). Moreover, they often either overlook difference of seasonal and long-term trend or use non-learnable methods to handle them (Woo et al. 2022).

A considerable subset of deep models overlook the importance of modeling the dependencies of variates (Nie et al. 2023; Zeng et al. 2023a; M. Zhang et al. 2023). These dependencies, however, are not always useful; specifically when the target time series is not correlated with other covariates (S.-A. Chen et al. 2023). Despite several studies exploring the importance of learning cross variate dependencies (S.-A. Chen et al. 2023; Yong Liu, Hu, et al. 2024; Y. Zhang and Yan 2023), there has been no universal standard and the conclusion has been different depending on the domain and benchmarks. Accordingly, we argue that an effective model need to *adaptively* learn to capture the dependencies of variates in a data-dependent manner. In this direction, recently, Yong Liu, Hu, et al. (2024) argue that attention mechanisms are more effective when they are used across variates, showing the importance of modeling complex non-linear dependencies across the variate axis in a data-dependent manner. However, the quadratic complexity of Transformers challenges the model on multivariate time series with a large number of variates (e.g., brain activity signals (Behrouz, Delavari, and Hashemi 2024) or traffic forecasting (H. Zhou et al. 2021)), limiting the efficient training and inference (see Table 3, and Table 5).

The objective of this study is to develop a provably expressive model for multivariate time series that not only can model the dynamics of the dependencies along both time and variates, but it also takes advantage of fast training and inference. To this end, we present a Chimera, a three-headed two-dimensional State Space Model (SSM) that is based on linear layers along (i) time, (ii) variates, (iii) time \rightarrow variate, and (iv) variate \rightarrow time. Chimera has a careful parameterization based on the pair of companion and diagonal matrices (see Figure 1), which is provably expressive to recover both

classical methods (Bartholomew 1971; Bender and Simonovic 1994; Winters 1960), linear attentions, and recent SSM-based models (Behrouz, Santacatterina, and Zabih 2024; Nguyen et al. 2022). It further uses an adaptive module based on a 2D SSM with an especially designed discretization process to capture seasonal patterns. While our theoretical results and design of Chimera guarantee the first three criteria of an effective model, due to its 2D recurrence, the naive implementation of Chimera results in slow training. To address this issue, we reformulate its 2D recurrence as the prefix sum problem with a 2-dimensional associative operators. This new formulation can be done in parallel and has hardware-friendly implementation, resulting in much faster training and inference.

We discuss new variants of our 2D SSM in Section 3.2 by limiting its transition matrices. The resulted models can be seen as the generalization of Mamba (Gu and Dao 2023) and Mamba-2 (Dao and Gu 2024) to 2-dimensional data. While the main focus of this paper is on time series data, these presented models due to their 2D inductive bias are potentially suitable for other high dimensional data modalities such as images, videos, and multi-channel audio.

In our experimental evaluation, we explore the performance of Chimera in a wide range of tasks: ECG and audio speech time series classification, long- and short-term time series forecasting, and anomaly detection tasks. We find that Chimera achieve superior or on par performance with state-of-the-art methods, while having faster training and less memory consumption. We perform a case study on the human brain activity signals (Behrouz, Delavari, and Hashemi 2024) to show (1) the effectiveness of Chimera and (2) evaluate the importance of modeling the dynamics of the variates dependencies.

2 Preliminaries

Notations. In this paper we mainly focus on classification and forecasting tasks. Note that anomaly detection can be seen as a binary classification task, where 0 means “normal” and 1 means “anomaly”. We let $\mathbf{X} = \{\mathbf{x}_1, \dots, \mathbf{x}_N\} \in \mathbb{R}^{N \times T}$ be the input sequences, where N is the number of variates and T is the time steps. We use $\mathbf{x}_{v,t}$ to refer to the value of the series v at time t . In classification (anomaly detection) tasks, we aim to classify input sequences and for forecasting tasks, given an input sequence \mathbf{x}_i , we aim to predict $\tilde{\mathbf{x}}_i \in \mathbb{R}^{1 \times H}$, i.e., the next H time steps for variate \mathbf{x}_i , where H is called horizon. In 2D SSMs formulation, for a 2-dimensional vector $\mathbf{x} \in \mathbb{C}^1$, we use $\mathbf{x}^{(1)}$ and $\mathbf{x}^{(2)}$ to refer to its real and imaginary components, respectively.

Multi-Dimensional State Space Models. We build our approach on the continuous State Space Model (SSM) but later we make each component of Chimera discrete by a designed discretization process. For additional discussion on 1D SSMs see Appendix A. Given parameters $\mathbf{A}_{\tau_1} \in \mathbb{R}^{N^{(\tau_1)} \times N^{(\tau_1)}}$, $\mathbf{B}_{\tau_2} \in \mathbb{C}^{N^{(\tau_2)} \times 1}$, and $\mathbf{C} \in \mathbb{C}^{N_1 \times N_2}$ for $\tau_1 \in \{1, \dots, 4\}$ and $\tau_2 \in \{1, 2\}$, the general form of the time-invariant 2D SSM is the map $\mathbf{x} \in \mathbb{C}^1 \mapsto \mathbf{y} \in \mathbb{C}^1$ defined by the linear Partial Differential Equation (PDE) with initial condition $h(0, 0) = 0$:

$$\frac{\partial}{\partial t^{(1)}} h(t^{(1)}, t^{(2)}) = \left(\mathbf{A}_1 h^{(1)}(t^{(1)}, t^{(2)}), \mathbf{A}_2 h^{(2)}(t^{(1)}, t^{(2)}) \right) + \mathbf{B}_1 \mathbf{x}(t^{(1)}, t^{(2)}), \quad (1)$$

$$\frac{\partial}{\partial t^{(2)}} h(t^{(1)}, t^{(2)}) = \left(\mathbf{A}_3 h^{(1)}(t^{(1)}, t^{(2)}), \mathbf{A}_4 h^{(2)}(t^{(1)}, t^{(2)}) \right) + \mathbf{B}_2 \mathbf{x}(t^{(1)}, t^{(2)}), \quad (2)$$

$$\mathbf{y}(t^{(1)}, t^{(2)}) = \langle \mathbf{C}, \mathbf{x}(t^{(1)}, t^{(2)}) \rangle. \quad (3)$$

Contrary to the multi-dimensional SSMs discussed by Gu and Dao (2023) and Gu, Goel, and Re (2022), in which multi-dimension refers to the dimension of the input but with one time variable, the above formulation uses two variables, meaning that the mapping is from a 2D grid to a 2D grid.

(Seasonal) Autoregressive Process. Autoregressive process is a basic yet essential premise for time series modeling, which models the causal nature of time series. Given $p \in \mathbb{N}$, $\mathbf{x}_k \in \mathbb{R}^d$, the simple linear autoregressive relationships between \mathbf{x}_k and its past samples $\mathbf{x}_{k-1}, \mathbf{x}_{k-2}, \dots, \mathbf{x}_{k-p}$ can be modeled as $\mathbf{x}_k = \phi_1 \mathbf{x}_{k-1} + \phi_2 \mathbf{x}_{k-2} + \dots + \phi_p \mathbf{x}_{k-p}$, where ϕ_1, \dots, ϕ_p are coefficients. This is called AR(p). Similarly, in the presence of seasonal patterns, the seasonal autoregressive process, SAR(p, q, s), is:

$$\mathbf{x}_k = \phi_1 \mathbf{x}_{k-1} + \phi_2 \mathbf{x}_{k-2} + \dots + \phi_p \mathbf{x}_{k-p} + \eta_1 \mathbf{x}_{k-s} + \eta_2 \mathbf{x}_{k-2s} + \dots + \eta_q \mathbf{x}_{k-qs}, \quad (4)$$

where s is the frequency of seasonality, and ϕ_1, \dots, ϕ_p and η_1, \dots, η_q are coefficients. Note that one can simply extend the above formulation to multivariate time series by letting coefficients to be vectors and replace the product with element-wise product.

3 Chimera: A Three-headed 2-Dimensional State Space Model

In this section, we first present a mathematical model for multivariate time series data and then based on this model, we present a neural architecture that can satisfy all the criteria discussed in §1.

3.1 Motivations & Chimera Model

SSMs have been long-standing methods for modeling time series (Aoki 2013; Harvey 1990), mainly due to their simplicity and expressive power to represent complicated and autoregressive dependencies. Their states, however, are the function of a single-variable (e.g., time). Multivariate time series, on the other hand, require capturing dependencies along both time and variate dimensions, requiring the current state of the model to be the function of both time and variate. Classical 2D SSMs (Eising 1978; Fornasini and Marchesini 1978; Hinamoto 1980; Kung et al. 1977), however, struggle to achieve good performance compared to recent advanced deep learning methods as they are : (1) only able to capture linear dependencies, (2) discrete by design, having a pre-determined resolution, and so cannot simply model seasonal patterns, (3) slow in practice for large datasets, (4) their update parameters are static and cannot capture the dynamics of dependencies. Deep learning-based methods (S.-A. Chen et al. 2023; Yong Liu, Hu, et al. 2024; H. Zhou et al. 2021), on the other hand, potentially are able to address a subset of the above limitations, while having their own drawbacks (discussed in §1). In this section, we start with *continuous* SSMs due to their connection to both classical methods (Aoki 2013; Harvey 1990) and recent breakthrough in deep learning (Gu and Dao 2023; Gu, Goel, and Re 2022). We then discuss our contributions on how to take the advantages of the best of both worlds, addressing all the abovementioned limitations.

Discrete 2D SSM. We use 2-dimensional SSMs, introduced in Equation 1-3, to model multivariate time series, where the first axis corresponds to the time dimension and the second axis is the variates. Accordingly, each state is a function of both time and variates. The first stage is to transform the continuous form of 2D SSMs to discrete form. Given the step size Δ_1 and Δ_2 , which represent the resolution of the input along the axes, discrete form of the input is defined as $\mathbf{x}_{k,\ell} = \mathbf{x}(k\Delta_1, \ell\Delta_2)$. Using Zero-Order Hold (ZOH) method, we can discretize the input as (see Appendix C for details):

$$\begin{pmatrix} h_{k,\ell+1}^{(1)} \\ h_{k+1,\ell}^{(2)} \end{pmatrix} = \begin{pmatrix} \bar{\mathbf{A}}_1 & \bar{\mathbf{A}}_2 \\ \bar{\mathbf{A}}_3 & \bar{\mathbf{A}}_4 \end{pmatrix} \begin{pmatrix} h_{k,\ell}^{(1)} \\ h_{k,\ell}^{(2)} \end{pmatrix} + \begin{pmatrix} \bar{\mathbf{B}}_1 \\ \bar{\mathbf{B}}_2 \end{pmatrix} \otimes \begin{pmatrix} \bar{\mathbf{x}}_{k,\ell+1} \\ \bar{\mathbf{x}}_{k+1,\ell} \end{pmatrix}, \quad (5)$$

where $\bar{\mathbf{A}}_i = \exp\left(\Delta \lfloor \frac{i+1}{2} \rfloor \mathbf{A}_i\right)$ for $i = 1, 2, 3, 4$, $\bar{\mathbf{B}}_1 = \begin{bmatrix} \mathbf{A}_1^{-1} (\bar{\mathbf{A}}_1 - \mathbf{I}) \mathbf{B}_1^{(1)} \\ \mathbf{A}_2^{-1} (\bar{\mathbf{A}}_2 - \mathbf{I}) \mathbf{B}_1^{(2)} \end{bmatrix}$, and $\bar{\mathbf{B}}_2 = \begin{bmatrix} \mathbf{A}_3^{-1} (\bar{\mathbf{A}}_3 - \mathbf{I}) \mathbf{B}_2^{(1)} \\ \mathbf{A}_4^{-1} (\bar{\mathbf{A}}_4 - \mathbf{I}) \mathbf{B}_2^{(2)} \end{bmatrix}$.

Note that this formulation can also be viewed as the modification of the discrete Roesser’s SSM model (Kung et al. 1977) when we add a lag of 1 in the inputs $\begin{pmatrix} \bar{\mathbf{x}}_{i,j} \\ \bar{\mathbf{x}}_{i,j} \end{pmatrix}$. This modification, however, misses the discretization step, which is an important step in our model. We later use the discretization step to (1) empower the model to select (resp. filter) relevant (resp. irrelevant) information, (2) adaptively adjust the resolution of the method, capturing seasonal patterns.

From now on, we use t (resp. v) to refer to the index along the time (resp. variate) dimension. Therefore, for the sake of simplicity, we reformulate Equation 5 as follows:

$$h_{v,t+1}^{(1)} = \bar{\mathbf{A}}_1 h_{v,t}^{(1)} + \bar{\mathbf{A}}_2 h_{v,t}^{(2)} + \bar{\mathbf{B}}_1 \mathbf{x}_{v,t+1}, \quad (6)$$

$$h_{v+1,t}^{(2)} = \bar{\mathbf{A}}_3 h_{v,t}^{(1)} + \bar{\mathbf{A}}_4 h_{v,t}^{(2)} + \bar{\mathbf{B}}_2 \mathbf{x}_{v+1,t}, \quad (7)$$

$$\mathbf{y}_{v,t} = \mathbf{C}_1 h_{v,t}^{(1)} + \mathbf{C}_2 h_{v,t}^{(2)}, \quad (8)$$

where $\bar{\mathbf{A}}_1, \bar{\mathbf{A}}_2, \bar{\mathbf{A}}_3, \bar{\mathbf{A}}_4 \in \mathbb{R}^{N \times N}$, $\bar{\mathbf{B}}_1, \bar{\mathbf{B}}_2 \in \mathbb{R}^{N \times 1}$, and $\mathbf{C}_1, \mathbf{C}_2 \in \mathbb{R}^{1 \times N}$ are parameters of the model, $h_{v,t}^{(1)}, h_{v,t}^{(2)} \in \mathbb{R}^{N \times d}$ are hidden states, and $\mathbf{x}_{v,t} \in \mathbb{R}^{1 \times d}$ is the input. In this formulation, intuitively, $h_{v,t}^{(1)}$ is the hidden state that carries cross-time information (each state depends on its previous time stamp but within the same variate), where $\bar{\mathbf{A}}_1$ and $\bar{\mathbf{A}}_2$ control the emphasis on past cross-time and cross-variate information, respectively. Similarly, $h_{v,t}^{(2)}$ is the hidden state that carries cross-variate information (each state depends on other variates but with the same time stamp). Later in this section, we discuss to modify the model to bi-directional setting along the variate dimension, to enhance information flow along this non-causal dimension.

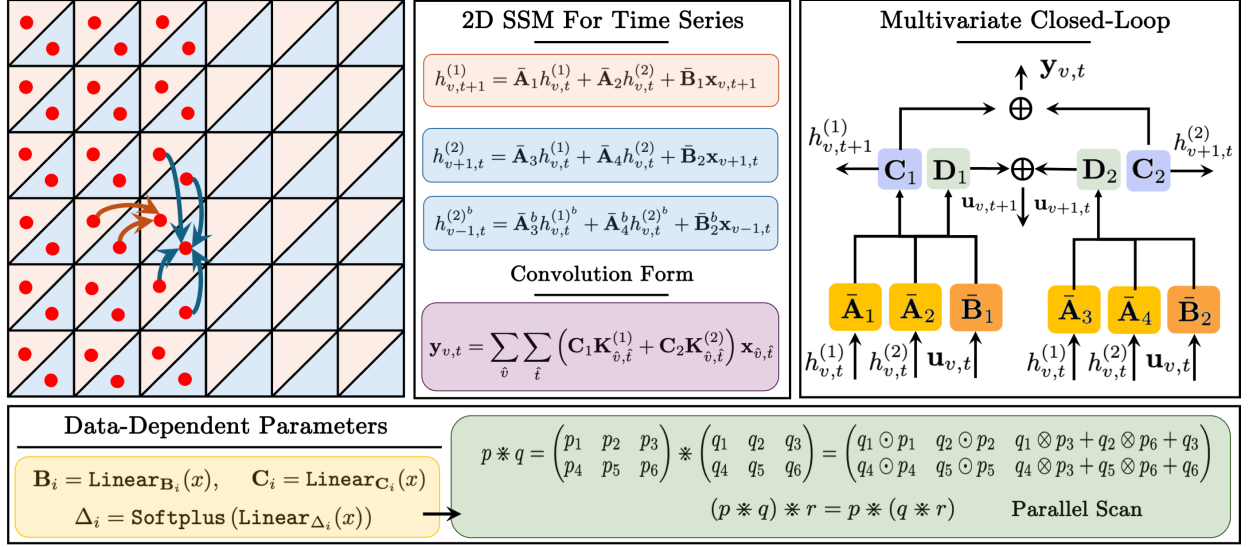


Figure 2: **Different forms of Chimera.** (**Top-Left**) Chimera has a recurrence form (bi-directional along the variates), which also can be computed as a global convolution in training. (**Top-Right**) In forecasting, we present the multivariate closed-loop to improve the performance for long horizons. (**Bottom**) Using data-dependent parameters, Chimera training can be done as a parallel 2D scan.

Interpretation of Discretization. Time series data are often sampled from an underlying continuous process (Hebart et al. 2023; Warden 2018). In these cases, variable Δ_1 in the discretization of the time axis can be interpreted as resolution or the sampling rate from the underlying continuous data. However, discretization along the variate axis, which is discrete by its nature, or when working directly with discrete data (A. E. Johnson et al. 2023) is an unintuitive process, and raise questions about its significance. The discretization step in 1D SSMs has deep connections to gating mechanisms of RNNs (Gu and Dao 2023; Tallec and Ollivier 2018), automatically ensures that the model is normalized (Gu, I. Johnson, Timalsina, et al. 2023), and results in desirable properties such as resolution invariance (Nguyen et al. 2022).

Proposition 3.1. *The 2D discrete SSM introduced in Equation 6-8 with parameters $(\{\bar{A}_i\}, \{\bar{B}_i\}, \{\bar{C}_i\}, k\Delta_1, \ell\Delta_2)$ evolves at a rate k (resp. ℓ) times as fast as the 2D discrete SSM with parameters $(\{\bar{A}_i\}, \{\bar{B}_i\}, \{\bar{C}_i\}, \Delta_1, \ell\Delta_2)$ (resp. $(\{\bar{A}_i\}, \{\bar{B}_i\}, \{\bar{C}_i\}, k\Delta_1, \Delta_2)$).*

Accordingly, parameters Δ_1 can be viewed as the controller of the length of dependencies that the model captures. That is, based on the above result, we see the discretization along the time axis as the setting of the resolution or sampling rate: while small Δ_1 can capture long-term progression, larger Δ_1 captures seasonal patterns. For now, we see the discretization along the variate axis as a mechanism similar to gating in RNNs (Gu and Dao 2023; Gu, Gulcehre, et al. 2020), where Δ_2 controls the length of the model context. Larger values of Δ_2 means less context window, ignoring other variates, while smaller values of Δ_2 means more emphasis on the dependencies of variates. Later, inspired by Gu and Dao (2023), we discuss making Δ_2 as the function of the input, resulting in a selection mechanism that filters irrelevant variates.

Structure of Transition Matrices. For Chimera to be expressive and able to recover autoregressive process, hidden states $h_{v,t}^{(1)}$ should carry information about *past* time stamps. While making all the parameters in A_i learnable allows the model to learn any arbitrary structure for A_i , previous studies show that this is not possible unless the structure of transition matrices are restricted (Gu, Goel, Gupta, et al. 2022; Gu, I. Johnson, Goel, et al. 2021). To this end, inspired by M. Zhang et al. (2023) that argue that companion matrices are effective to capture the dependencies along the time dimension, we restrict A_1 and A_2 matrices to have companion structure:

$$\mathbf{A}_i = \begin{pmatrix} 0 & 0 & \dots & 0 & a_1^{(i)} \\ 1 & 0 & \dots & 0 & a_2^{(i)} \\ 0 & 1 & \dots & 0 & a_3^{(i)} \\ \vdots & \vdots & \ddots & \vdots & \vdots \\ 0 & 0 & \dots & 1 & a_N^{(i)} \end{pmatrix} = \underbrace{\begin{pmatrix} 0 & 0 & \dots & 0 & 0 \\ 1 & 0 & \dots & 0 & 0 \\ 0 & 1 & \dots & 0 & 0 \\ \vdots & \vdots & \ddots & \vdots & \vdots \\ 0 & 0 & \dots & 1 & 0 \end{pmatrix}}_{\text{Shift Matrix}} + \underbrace{\begin{pmatrix} 0 & 0 & \dots & 0 & a_1^{(i)} \\ 0 & 0 & \dots & 0 & a_2^{(i)} \\ 0 & 0 & \dots & 0 & a_3^{(i)} \\ \vdots & \vdots & \ddots & \vdots & \vdots \\ 0 & 0 & \dots & 0 & a_N^{(i)} \end{pmatrix}}_{\text{Low-rank Matrix}}, \quad (9)$$

for $i = 1, 2$. Note that these two matrices are responsible to fuse the information along the time axis (see Figure 2). Not only this formulation is shown to be effective for capturing dependencies along the time dimension (M. Zhang et al. 2023) (also see Theorem 3.4), but it also can help us to compute the power of \mathbf{A}_1 and \mathbf{A}_2 faster in the convolutional form, as discussed by M. Zhang et al. (2023). Also, for \mathbf{A}_3 and \mathbf{A}_4 , we observe that even a simpler structure of diagonal matrices is effective to fuse information along the variate dimension. Not only these simple structured matrices make the training of the model faster, but they also are proven to be effective (Gu, Goel, Gupta, et al. 2022).

Bi-Directionality. The causal nature of the 2D SSM result in limited information flow along the variate dimension as variate are not ordered. To overcome this challenge, inspired by the bi-directional 1D SSMs (J. Wang et al. 2023), we use two different modules for forward and backward pass along the variate dimension:

$$\begin{aligned} h_{v,t+1}^{(1)f} &= \bar{\mathbf{A}}_1^f h_{v,t}^{(1)} + \bar{\mathbf{A}}_2^f h_{v,t}^{(2)f} + \bar{\mathbf{B}}_1^f \mathbf{x}_{v,t+1}, \\ h_{v,t+1}^{(1)b} &= \bar{\mathbf{A}}_1^b h_{v,t}^{(1)} + \bar{\mathbf{A}}_2^b h_{v,t}^{(2)} + \bar{\mathbf{B}}_1^b \mathbf{x}_{v,t+1}, \end{aligned} \quad (10)$$

$$\begin{aligned} h_{v+1,t}^{(2)f} &= \bar{\mathbf{A}}_3^f h_{v,t}^{(1)} + \bar{\mathbf{A}}_4^f h_{v,t}^{(2)f} + \bar{\mathbf{B}}_2^f \mathbf{x}_{v+1,t}, \\ h_{v-1,t}^{(2)b} &= \bar{\mathbf{A}}_3^b h_{v,t}^{(1)b} + \bar{\mathbf{A}}_4^b h_{v,t}^{(2)b} + \bar{\mathbf{B}}_2^b \mathbf{x}_{v-1,t}, \end{aligned} \quad (11)$$

$$\mathbf{y}_{v,t}^f = \mathbf{C}_1^f h_{v,t}^{(1)f} + \mathbf{C}_2^f h_{v,t}^{(2)f}, \quad (12)$$

$$\mathbf{y}_{v,t}^b = \mathbf{C}_1^b h_{v,t}^{(1)b} + \mathbf{C}_2^b h_{v,t}^{(2)b}, \quad (13)$$

$$\mathbf{y}_{v,t} = \mathbf{y}_{v,t}^f + \mathbf{y}_{v,t}^b, \quad (14)$$

where $\bar{\mathbf{A}}_1^\tau, \bar{\mathbf{A}}_2^\tau, \bar{\mathbf{A}}_3^\tau, \bar{\mathbf{A}}_4^\tau \in \mathbb{R}^{N \times N}$, $\bar{\mathbf{B}}_1^\tau, \bar{\mathbf{B}}_2^\tau \in \mathbb{R}^{N \times 1}$, and $\mathbf{C}_1^\tau, \mathbf{C}_2^\tau \in \mathbb{R}^{1 \times N}$ are parameters of the model, $h_{v,t}^{(1)\tau}, h_{v,t}^{(2)\tau} \in \mathbb{R}^{N \times d}$ are hidden states, $\mathbf{x}_{v,t} \in \mathbb{R}^{1 \times d}$ is the input, and $\tau \in \{f, b\}$. Figure 2 illustrates the bi-directional recurrence process in Chimera. For the sake of simplicity, we continue with unidirectional pass, but adapting them for bi-directional setting is simple as we use two separate blocks, each of which for a direction.

Convolution Form. Similar to 1D SSMs (Gu, Goel, and Re 2022), our *data-independent* formulation can be viewed as a convolution with a kernel \mathbf{K} . This formulation not only results in faster training by providing the ability of parallel processing, but it also connect Chimera with very recent studies of modern convolution-based architecture for time series (Luo and X. Wang 2024). Applying the recurrent rules in Equation 6-8, we can write the output as:

$$\mathbf{y}_{v,t} = \sum_{1 \leq \hat{v} \leq v} \sum_{1 \leq \hat{t} \leq t} \left(\mathbf{C}_1 \mathbf{K}_{\hat{v}, \hat{t}}^{(1)} + \mathbf{C}_2 \mathbf{K}_{\hat{v}, \hat{t}}^{(2)} \right) \mathbf{x}_{\hat{v}, \hat{t}}, \quad (15)$$

where kernels $\mathbf{K}_{\hat{v}, \hat{t}}^{(\tau)} = \sum_{(z_1, \dots, z_5) \in \mathbf{P}^{(\tau)}} q_i \bar{\mathbf{A}}_1^{p_1} \bar{\mathbf{A}}_2^{p_2} \bar{\mathbf{A}}_3^{p_3} \bar{\mathbf{A}}_4^{p_4} \bar{\mathbf{B}}_{p_5}$, and $\mathbf{P}^{(\tau)}$ is the partitioning of the paths from the starting point to (\hat{v}, \hat{t}) for $\tau \in \{1, 2\}$. As discussed by Baron, Zimmerman, and Wolf (2024), if the power of $\bar{\mathbf{A}}_i$ s are given and cached, calculating the partitioning of all paths can be done very efficiently (near-linearly) as it the generalization of pascal triangle. To calculate the power of $\bar{\mathbf{A}}_i$, note that we use diagonal matrices as the structure of $\bar{\mathbf{A}}_3$, and $\bar{\mathbf{A}}_4$, and so computing their powers is very fast. On the other hand, for $\bar{\mathbf{A}}_1$ and $\bar{\mathbf{A}}_2$ with companion structures, we can use sparse matrix multiplication, which results in linear complexity in terms of the sequence length.

Data-Dependent Parameters. As discussed earlier, parameters $\bar{\mathbf{A}}_1$ and $\bar{\mathbf{A}}_2$ controls the emphasis on past cross-time and cross-variate information. Similarly, parameters $\bar{\mathbf{B}}_1$ and $\bar{\mathbf{B}}_2$ controls the emphasis on the current input and historical data. Since these parameters are data-independent, one can interpret them as a global feature of the system. In complex systems (e.g., human neural activity), however, the emphasis depends on the current input, requiring these parameters to be the function of the input (see §4.1). The input-dependency of parameters allows the model to select relevant and filter irrelevant information for each input data, providing a similar mechanism as transformers (Gu and Dao 2023). Additionally, as we argue earlier, depending on the data, the model needs to adaptively learn if mixing information along the variates is useful. Making parameters input-dependent further overcomes this challenge and lets our model to mix relevant and filter irrelevant variates for the modeling of a variate of interest. One of our main technical contributions is to let $\bar{\mathbf{B}}_i$, $\bar{\mathbf{C}}_i$, and Δ_i for $i \in \{1, 2\}$ be the function of the input $\mathbf{x}_{v,t}$. This input-dependent 2D SSM, unfortunately, does not have the convolution form, limiting the scalability and efficiency of the training. We overcome this challenge by computing the model recurrently with a new 2D scan.

2D Selective Scan. Inspired by the scanning in 1D SSMs (Gu and Dao 2023; Smith, Warrington, and Linderman 2023), we present an algorithm to decrease the sequential steps that are required to calculate hidden states. Given p, q , each of which with 6 elements, we first define operation \ast as: (\odot is matrix-matrix and \otimes is matrix-vector multiplication)

$$p \ast q = \begin{pmatrix} p_1 & p_2 & p_3 \\ p_4 & p_5 & p_6 \end{pmatrix} \ast \begin{pmatrix} q_1 & q_2 & q_3 \\ q_4 & q_5 & q_6 \end{pmatrix} = \begin{pmatrix} q_1 \odot p_1 & q_2 \odot p_2 & q_1 \otimes p_3 + q_2 \otimes p_6 + q_3 \\ q_4 \odot p_4 & q_5 \odot p_5 & q_4 \otimes p_3 + q_5 \otimes p_6 + q_6 \end{pmatrix}$$

The proofs of the next two theorems are in [Appendix E](#).

Theorem 3.2. *Operator \ast is associative: Given p, q , and r , we have: $(p \ast q) \ast r = p \ast (q \ast r)$.*

Theorem 3.3. *2D SSM recurrence can be done in parallel using parallel prefix sum algorithms with associative operator \ast .*

3.2 New Variants of 2D SSM: 2D Mamba and 2D Mamba-2

[Figure 2 \(Top-Left\)](#) shows the recurrence form of our 2D SSM. Each small square is a state of the system, i.e., the state of a variate at a certain time stamp. 2D SSM considers two hidden states for each state (represented by two colors: light red and blue), encoding the information along the time (red) and variate (blue), respectively. Furthermore, each arrow represents a transition matrix A_i that decides how information need to be fused. In this section, we discuss different variants of our 2D SSM by limiting its parameters.

2D Mamba. We let $A_2 = A_3 = \mathbf{0}$ in the formulation of our 2D SSM. The resulting model is equivalent to:

$$h_{v,t+1}^{(1)} = \bar{A}_1 h_{v,t}^{(1)} + \bar{B}_1 x_{v,t+1}, \quad (16)$$

$$h_{v+1,t}^{(2)} = \bar{A}_4 h_{v,t}^{(2)} + \bar{B}_2 x_{v+1,t}, \quad (17)$$

$$y_{v,t} = C_1 h_{v,t}^{(1)} + C_2 h_{v,t}^{(2)}, \quad (18)$$

where $\bar{A}_1 = \exp(\Delta_1 A_1)$, $\bar{A}_2 = \exp(\Delta_2 A_2)$, $\bar{B}_1 = \begin{bmatrix} A_1^{-1} (\bar{A}_1 - I) B_1^{(1)} \\ \mathbf{0} \end{bmatrix}$, and $\bar{B}_2 = \begin{bmatrix} \mathbf{0} \\ A_4^{-1} (\bar{A}_4 - I) B_2^{(2)} \end{bmatrix}$. This formulation with data-dependent parameters, is equivalent to using two S6 blocks (Gu and Dao 2023) each of which along a dimension. Notably, these two S6 blocks are not separate as the output $y_{v,t}$ is based on both hidden states $h_{v,t}^{(1)}$ and $h_{v,t}^{(2)}$, capturing 2D inductive bias.

2D Mamba-2. Recently, Dao and Gu (2024) present Mamba-2 that re-formulates S6 block using structured semi-separable matrices, resulting in more efficient training and ability of having larger recurrent state sizes. Although we leave the exploration of how generic 2D SSMs can be re-formulated by tensors (see [Section 5](#) for further discussion), the special case of $A_2 = A_3 = \mathbf{0}$, similar to the above formulation, can be re-formulated as two SSD blocks (Dao and Gu 2024) each of which along a dimension. Furthermore, for bi-directionality along the variates, one can use quasi-separable structured matrices, which inherently captures bi-directionality as discussed by Behrouz, Santacatterina, and Zabih (2024):

$$y_{v,t} = \underbrace{\begin{pmatrix} C_{1,v,1} \bar{B}_{1,v,1} & \mathbf{0} & \dots & \mathbf{0} \\ C_{1,v,2} A_{1,v,2} \bar{B}_{1,v,1} & C_{1,v,2} \bar{B}_{1,v,2} & \dots & \mathbf{0} \\ \vdots & \vdots & \ddots & \vdots \\ C_{1,v,t} (\prod_{i=2}^t A_{1,v,i}) \bar{B}_{1,v,1} & C_{1,v,t} (\prod_{i=3}^t A_{1,v,i}) \bar{B}_{1,v,2} & \dots & C_{1,v,t} \bar{B}_{1,v,t} \end{pmatrix}}_{\text{SSD Block}} x_{v,:} \quad (19)$$

$$+ \underbrace{\begin{pmatrix} \gamma_1 & C'_{2,v-1,t} A'_{4,v-1,t} \bar{B}'_{2,v,t} & \dots & C'_{2,v,t} \left(\prod_{i=1}^{v-1} A'_{4,i,t} \right) \bar{B}'_{2,1,t} \\ C_{2,v,t} A_{4,v,t} \bar{B}_{2,1,t} & \gamma_2 & \dots & C'_{2,v-1,t} \left(\prod_{i=2}^{v-1} A'_{4,i,t} \right) \bar{B}'_{2,2,t} \\ \vdots & \vdots & \ddots & \vdots \\ C_{2,v,t} \left(\prod_{i=2}^v A_{4,i,t} \right) \bar{B}_{2,1,t} & C_{2,v,t} \left(\prod_{i=3}^v A_{4,i,t} \right) \bar{B}_{2,2,t} & \dots & \gamma_v \end{pmatrix}}_{\text{Quasi-Separable Block}} x_{:,t}, \quad (20)$$

where $x_{v,:}$ and $x_{:,t}$ are the vectors when we fix v and t in input x , respectively.

3.3 Chimera Neural Architecture

In this section, we use a stack of our 2D SSMs, with non-linearity in between, to enhance the expressive power and capabilities of the abovementioned 2D SSM. To this end, similar to deep SSM models (M. Zhang et al. 2023), we allow all parameters to be learnable and in each layer we use multiple 2D SSMs, each of which with its own responsibility. Also, in the data-dependent variant of Chimera, we let parameters \mathbf{B}_i , \mathbf{C}_i , and Δ_i for $i \in \{1, 2\}$ be the function of the input \mathbf{x} :

$$\mathbf{B}_i = \text{Linear}_{\mathbf{B}_i}(\mathbf{x}), \quad \mathbf{C}_i = \text{Linear}_{\mathbf{C}_i}(\mathbf{x}), \quad \Delta_i = \text{Softplus}(\text{Linear}_{\Delta_i}(\mathbf{x})). \quad (21)$$

Chimera follows the commonly used decomposition of time series, and decomposes them into trend components and seasonal patterns. it, however, uses special traits of 2D SSM to capture these terms.

Seasonal Patterns. To capture the multi-resolution seasonal patterns, we take advantage of the discretization process. **Proposition 3.1** states that if $\mathbf{x}(v, t) \mapsto \mathbf{y}(v, t)$ with parameters $(\{\tilde{\mathbf{A}}_i\}, \{\tilde{\mathbf{B}}_i\}, \{\tilde{\mathbf{C}}_i\}, \Delta_1, \Delta_2)$ then $\mathbf{x}(v, kt) \mapsto \mathbf{y}(v, kt)$ with $(\{\tilde{\mathbf{A}}_i\}, \{\tilde{\mathbf{B}}_i\}, \{\tilde{\mathbf{C}}_i\}, k\Delta_1, \Delta_2)$. Accordingly, we use 2D-SSM(.) module with a separate learnable Δ_s that is responsible to learn the best resolution to capture seasonal patterns. Another interpretation for this module is based on SAR(p, q, s) (Equation 4). In this case, Δ_s aims to learn a proper parameter s to capture seasonal patterns. Since we expect the resolution before and after this module matches, we add additional re-discretization module (a simple linear layer), after this module.

Trend Components. The second module of Chimera, 2D-SSM_t(.) simply uses a sequence of multiple 2D SSMs to learn trend components. Proper combination of the outputs of this and the previous modules can capture both seasonal and trend components.

Both Modules Together. We followed previous studies (Toner and Darlow 2024) and consider residual connection modeling for learning trend and seasonal patterns. Given input data $\tilde{\mathbf{X}}_0 = \mathbf{X}$, and $\ell = 0, \dots, \mathcal{L}$, we have:

$$\hat{\mathbf{X}}_{\ell+1} = \text{2D-SSM}_t(\tilde{\mathbf{X}}_\ell), \quad (22)$$

$$\tilde{\mathbf{X}}_{\ell+1} = \text{Re-Discretization}\left(\text{2D-SSM}_s(\tilde{\mathbf{X}}_\ell - \hat{\mathbf{X}}_{\ell+1})\right). \quad (23)$$

Figure 1 illustrate the architecture of Chimera. Due to the ability of our 2D SSM to recover smoothing techniques (see Theorem 3.4), this combination of modules for trend and seasonal patterns can be viewed as a generalization of traditional methods that use moving average with residual connection to model seasonality (Toner and Darlow 2024).

Gating with Linear Mapping. Inspired by the success of gated recurrent and SSM-based models (Gu and Dao 2023; Qin, Yang, and Zhong 2023), we use a head of a fully connected layer with Swish (Ramachandran, Zoph, and Le 2017), resulting in SwiGLU variant (Touvron et al. 2023). While we validate the significance of this head, this

Closed-Loop 2D SSM Decoder. To enhance the generalizability and the ability of our model for longer-horizon, we extend the closed-loop decoder module (M. Zhang et al. 2023), which is similar to autoregression, to multivariate time series. We use distinct processes for the inputs and outputs, using additional matrices \mathbf{D}_1 and \mathbf{D}_2 in each decoder 2D SSM, we model future input time-steps explicitly:

$$\mathbf{y}_{v,t} = \mathbf{C}_1 \mathbf{h}_{v,t}^{(1)} + \mathbf{C}_2 \mathbf{h}_{v,t}^{(2)}, \quad (24)$$

$$\mathbf{u}_{v,t} = \mathbf{D}_1 \mathbf{h}_{v,t}^{(1)} + \mathbf{D}_2 \mathbf{h}_{v,t}^{(2)}, \quad (25)$$

where $\mathbf{u}_{v,t}$ is the next input and $\mathbf{y}_{v,t}$ is the output. Note that the other parts (recurrence) are the same as Equation 6. Figure 2 illustrate the architecture of closed-loop 2D SSM.

3.4 Theoretical Justification

In this section, we provide some theoretical evidences for the performance of Chimera. These results are mostly revisiting the theorems by M. Zhang et al. (2023) and Baron, Zimmerman, and Wolf (2024), and extending them for Chimera. In the first theorem, we show that Chimera recovers several classic methods, and pre-processing steps as it can recover SpaceTime (M. Zhang et al. 2023) and additionally because of its design, it can recover SARIMA (Bender and Simonovic 1994):

Theorem 3.4. *Chimera can represent seasonal autoregressive process, SARIMA (Bender and Simonovic 1994), SpaceTime (M. Zhang et al. 2023), and so ARIMA (Bartholomew 1971), exponential smoothing (Winters 1960), and controllable linear time-invariant systems (C.-T. Chen 1984).*

Theorem 3.5. *Chimera can represent S4nd (Nguyen et al. 2022), TSM2 (Behrouz, Santacatterina, and Zabih 2024), and TSMixer (S.-A. Chen et al. 2023).*

Next theorem compares the expressiveness of Chimera with some existing 2D deep SSMs. Since Chimera can recover 2DSSM (Baron, Zimmerman, and Wolf 2024), it can express full-rank kernels with a constant number of parameters:

Theorem 3.6. *Similar to 2DSSM (Baron, Zimmerman, and Wolf 2024), Chimera can express full-rank kernels with $O(1)$ parameters, while existing deep SSMs (Behrouz, Santacatterina, and Zabih 2024; Nguyen et al. 2022) require $O(N)$ parameters to express N -rank kernels.*

4 Experiments

Goals and Baselines. We evaluate Chimera on a wide range of time series tasks. In § 4.1 we compare Chimera with the state-of-the-art general multivariate time series models (Behrouz, Santacatterina, and Zabih 2024; Das et al. 2023; Lim and Zohren 2021; M. Liu et al. 2022; Yong Liu, Hu, et al. 2024; Luo and X. Wang 2024; Badri N. Patro and Vijay S. Agneeswaran 2024; Woo et al. 2022; H. Wu, Hu, et al. 2023; H. Wu, Xu, et al. 2021; Y. Zhang and Yan 2023; T. Zhou, Z. Ma, Wen, X. Wang, et al. 2022) on long-term forecasting and classification tasks. In the next part, we test the performance of Chimera in short-term forecasting. In § 4.1 we perform a case study on human neural activity to classify seen images, which requires capturing complex dynamic dependencies of variates, to test the ability of Chimera in capturing cross-variate information and the significance of data-dependency. In § 4.2, we evaluate the significance of the Chimera’s components by performing ablation studies. In § 4.2, we evaluate whether the superior performance of Chimera coincide with its efficiency. Finally, we test the Chimera’s generalizability on unseen variates and further evaluate its ability to filter irrelevant context in § 4.3. The details and additional experiments are in Appendix G.

Table 1: Average Performance on long-term forecasting task. The first and second results are highlighted in **red** (bold) and orange (underline). Full results are reported in Appendix G.

	Chimera (ours)		TSM2 2024		Simba 2024		TCN 2024		iTransformer 2024		RLinear 2023		PatchTST 2023		Crossformer 2023		TiDE 2023		TimesNet 2023		DLinear 2023	
	MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE
ETTh1	0.345	0.377	0.361	-	0.383	0.396	<u>0.351</u>	<u>0.381</u>	0.407	0.410	0.414	0.407	0.387	0.400	0.513	0.496	0.419	0.419	0.400	0.406	0.403	0.407
ETTh2	0.250	<u>0.316</u>	0.267	-	0.271	0.327	<u>0.253</u>	0.314	0.288	0.332	0.286	0.327	0.281	0.326	0.757	0.610	0.358	0.404	0.291	0.333	0.350	0.401
ETTm1	0.405	<u>0.424</u>	0.403	-	0.441	0.432	<u>0.404</u>	0.420	0.454	0.447	0.446	0.434	0.469	0.454	0.529	0.522	0.541	0.507	0.458	0.450	0.456	0.452
ETTm2	0.318	0.375	0.333	-	0.361	0.391	<u>0.322</u>	<u>0.379</u>	0.383	0.407	0.374	0.398	0.387	0.407	0.942	0.684	0.611	0.550	0.414	0.427	0.559	0.515
ECL	0.154	0.249	0.169	-	0.185	0.274	<u>0.156</u>	<u>0.253</u>	0.178	0.270	0.219	0.298	0.205	0.290	0.244	0.334	0.251	0.344	0.192	0.295	0.212	0.300
Exchange	<u>0.311</u>	0.358	0.443	-	-	-	0.302	<u>0.366</u>	0.360	0.403	0.378	0.417	0.367	0.404	0.940	0.707	0.370	0.413	0.416	0.443	0.354	0.414
Traffic	<u>0.403</u>	<u>0.286</u>	0.420	-	0.493	0.291	0.398	0.270	0.428	0.282	0.626	0.378	0.481	0.304	0.550	0.304	0.760	0.473	0.620	0.336	0.625	0.383
Weather	0.219	0.258	0.239	-	0.255	0.280	<u>0.224</u>	<u>0.264</u>	0.258	0.278	0.272	0.291	0.259	0.281	0.259	0.315	0.271	0.320	0.259	0.287	0.265	0.317
1 st Count	5	5	1	-	0	0	<u>2</u>	<u>3</u>	0	0	0	0	0	0	0	0	0	0	0	0	0	0

4.1 Main Results: Classification and Forecasting

Long-Term Forecasting. We perform experiments in long-term forecasting task on benchmark datasets (H. Zhou et al. 2021). Table 1 reports the average of results over different horizons (for the results of each see Table 8). Chimera shows outstanding performance, achieving the best or the second best results in all the datasets and outperforms baselines in 5 out of 8 benchmarks. Notably, it surpasses extensively studied MLP-based and Transformer-based models while being more efficient (see Table 3, Figure 4, and Appendix G), providing a better balance of performance and efficiency. It further significantly outperforms recurrent models, including very recent Mamba-based architectures (Behrouz, Santacatterina,

and Zabih 2024; Badri N. Patro and Vijay S. Agneeswaran 2024), unleashing the potential of classical models, SSMs, when are carefully designed in deep learning settings.

Classification and Anomaly Detection. We evaluate the performance of Chimera in ECG classification on PTB-XL dataset (Wagner et al. 2020) (see Table 2), speech classification (Warden 2018)(Table 3), 10 multivariate datasets from UEA Time Series Classification Archive (Bagnall et al. 2018) (see Figure 3 and Table 10), and anomaly detection tasks on five widely-used benchmarks: SMD (Su et al. 2019), SWaT (Mathur and Tippenhauer 2016), PSM (Abdulaal, Z. Liu, and Lancewicki 2021) and SMAP (Hundman et al. 2018) (see Figure 3 and Table 11). For each benchmark, we use the state-of-the-art methods that are applicable to the task as the baselines. Table 2 reports the performance of Chimera and baselines on ECG classification tasks. Chimera outperforms all the baselines in 4/6 tasks, while achieving the second best results on the remaining tasks. Since these tasks are univariate time series, we attribute the outstanding performance of Chimera, specifically compared to SpaceTime (M. Zhang et al. 2023), to its ability of capturing seasonal patterns and its input-dependent parameters, resulting in dynamically learn dependencies.

Table 3 reports the results on speech audio classification task, which require long-range modeling of time series. Due to the length of the sequence (16K), LSSL (Gu, Goel, and Re 2022) and Transformer (Vaswani et al. 2017) has out of memory (OOM) issue, showing the efficiency of Chimera compared to alternative backbones.

Finally, we report the summary of the results in multivariate time series classification and anomaly detection tasks in Figure 3. The full list of results can be found in Table 10 and Table 11. Chimera shows outstanding performance, achieving highest average accuracy and F1 score in classification and anomaly detection tasks even compared to very recent and state-of-the-art methods (Luo and X. Wang 2024; H. Wu, Hu, et al. 2023).

Table 2: ECG statement classification on PTB-XL (100 Hz version).

Tasks	All	Diag	Sub-diag	Super-diag	Form	Rhythm
Chimera	0.941	0.947	0.935	<u>0.930</u>	0.901	<u>0.975</u>
SpaceTime (M. Zhang et al. 2023)	0.936	<u>0.941</u>	<u>0.933</u>	0.929	0.883	0.967
S4 (Gu, Goel, and Re 2022)	<u>0.938</u>	0.939	0.929	0.931	0.895	0.977
Inception	0.925	0.931	0.930	0.921	<u>0.899</u>	0.953
xRN-101	0.925	0.937	0.929	0.928	0.896	0.957
LSTM	0.907	0.927	0.928	0.927	0.851	0.953
Transformer	0.857	0.876	0.882	0.887	0.771	0.831

Table 3: Speech classification.

Method	Acc. (%)
Chimera	98.40
SpaceTime	97.29
S4	<u>98.32</u>
LSSL	OOM
WaveGan-D	96.25
Transformer	OOM

Short-Term Forecasting. Our evaluation on short-term forecasting tasks on M4 benchmark datasets (Godahewa et al. 2021) reports in Table 4 (Full list in Table 9), which also shows the superior performance of Chimera compared to baselines.

Table 4: Short-term forecasting task on the M4 dataset. Full results are reported in Appendix G.

Models		Chimera (ours)	ModernTCN 2024	PatchTST 2023	TimesNet 2023	N-HiTS 2022	N-BEATS* 2019	ETS* 2022	LightTS 2022	DLinear 2023	FED* 2022	Stationary 2022	Auto* 2021	Pyra* 2021	In* 2021	Re* 2020	LSTM 1997
Weighted Average	SMAPE	11.618	<u>11.698</u>	11.807	11.829	11.927	11.851	14.718	13.525	13.639	12.840	12.780	12.909	16.987	14.086	18.200	160.031
	MASE	1.528	<u>1.556</u>	1.590	1.585	1.613	1.599	2.408	2.111	2.095	1.701	1.756	1.771	3.265	2.718	4.223	25.788
	OWA	0.827	<u>0.838</u>	0.851	0.851	0.861	0.855	1.172	1.051	1.051	0.918	0.930	0.939	1.480	1.230	1.775	12.642

Case Study of Brain Activity. Input dependency is a must to capture the dynamic of dependencies. To support this claim, we use BVFC (Behrouz, Delavari, and Hashemi 2024) (multivariate time series only), which aim to classify seen images by its corresponding brain activity response. This task, requires focusing more on the dependencies of brain units and their responses rather than the actual time series. Also, since each window corresponds to a specific image, the model needs to capture the dependencies based on the current window, requiring to be input-dependent. Results are reported in Table 5. Chimera significantly outperforms all the baselines including our Chimera but without data-dependent parameters (convolution form). Due to the large number of brain units, i.e., 9K, in the first dataset, transformer-based methods face OOM issue. However, they are also data-dependent and so shows the second best results in second and third datasets. This results support the significance of data-dependency in Chimera.

4.2 Ablation Study and Efficiency

To evaluate the significance of the Chimera’s design, we perform ablation studies and remove one of the components at each time, keeping other parts unchanged. Table 6 reports the results. The first row reports the Chimera’s performance, while row 2 uses unidirectional recurrence along the variate dimension, row 3 removes the gating mechanism, row 4 uses convolution form (data-independent), and row 5 removes the module for seasonal patterns. The results show that all the components of Chimera contributes to its performance.

Table 5: Image classification by brain activity (Acc. %).

Method	Chimera (ours)	Chimera (ind.) (ours)	SpaceTime 2023	S4 2022	iTrans. 2024	Trans. 2017	DLinear 2023
BVFC (9K)	69.41	<u>62.36</u>	41.20	40.89	OOM	OOM	39.74
BVFC (1K)	58.99	50.25	34.31	35.19	<u>54.18</u>	43.60	33.09
BVFC (400)	51.08	45.17	33.58	33.76	<u>48.22</u>	38.05	32.73

Table 6: Ablation study on the Chimera’s design.

Method	ETTh1		ETTm1		ETTh2	
	MSE	MAE	MSE	MAE	MSE	MAE
Chimera	0.405	0.424	0.345	0.377	0.318	0.375
Uni-directional	0.409	0.429	0.354	0.385	0.326	0.381
w/o Gating	0.418	0.433	0.351	0.384	0.321	0.379
Input-independent	0.471	0.498	0.361	0.389	0.372	0.401
w/o seasonal	0.426	0.431	0.357	0.382	0.331	0.386




Figure 3: Classification and anomaly detection performance. Full list with additional baselines is in Appendix G.




Figure 4: Wall-clock scaling.

Length of Time Series. We perform experiments on the effect of the sequence length on the efficiency of Chimera and baselines. The results are reported in Figure 4. Chimera scales linearly with respect to the sequence length and has smoother scaling than S4 (Gu, Goel, and Re 2022) and Transformers (Vaswani et al. 2017). These results also highlight the significance of our algorithm that uses 2D parallel scans for training Chimera. This algorithm results in $\approx \times 4$ faster training, which is very closed to the convolutional format *without* data dependency. Chimera also has a close running time to SpaceTime (M. Zhang et al. 2023), which has 1D recurrent.

4.3 Selection Mechanism Along Time and Variate

Variate Generalization. We argue that the data-dependency with discretization allows the model to filter the irrelevant context based on the input, resulting in more generalizability. Inspired by Yong Liu, Hu, et al. (2024), we train our model (and baseline) on 20% of variates and evaluate its generalizability to unseen variates. The results are reported in Figure 5. Chimera has on par generalizability compared to Transformers (when applied along the variate dimension), which we attributes to its data-dependent parameters as Chimera with convolution form performs poorly on unseen variates.

Context Filtering. Increasing the lookback length does not necessarily result in better performance for Transformers (Yong Liu, Hu, et al. 2024). Due to the selection mechanism of Chimera, we expect it to filter irrelevant information and monotonically performs better. Figure 6 reports the Chimera’s performance (w/ and w/o data-dependency) and transformer-based baselines (H. Wu, J. Wu, et al. 2022; H. Zhou et al. 2021) while varying the lookback length. Chimera due to its selection mechanism monotonically performs better with increasing the lookback.




Figure 5: Selection results in generalization to unseen variates.




Figure 6: Effect of lookback length.

5 Conclusion and Future Work

This paper presents Chimera, a three-headed 2-dimensional SSM model with provably high expressive power. Chimera is based on 2D SSMs with careful design of parameters that allows it to dynamically and simultaneously capture the dependencies along both time and variate dimensions. We provide different views of our 2D SSM for efficient training, and present a data-dependent formulation with a fast implementation using 2D scans. Chimera uses two different modules to capture trend and seasonal patterns and its discretization process allows these modules to adjust the resolution at each time stamp and for each variate. Our experimental and theoretical results support the effectiveness and efficiency of Chimera in a wide range of tasks.

Other Data Modalities. While the parameterization of Chimera is designed to expressively model time series data, the overall architecture of Chimera and our data-dependent 2D SSM with its 2D scan form in training are potentially applicable for other higher dimensional data types, e.g., images, videos, multi-channel speech, etc. Despite recent attempts to design effective SSM-based vision models (Badri Narayana Patro and Vijay Srinivas Agneeswaran 2024), the existing models suffer from the lack of 2D spatial inductive bias. Our 2D SSM, however, is able to provide 2D inductive bias, potentially being more effective than existing 1D selective SSMs. Accordingly, a promising direction is to explore the potential of 2D selective SSMs for other high dimensional data modalities and different tasks.

Variants of Chimera. As discussed in Section 3.2, different variants of Chimera result in the extension of well-known architectures like Mamba (Gu and Dao 2023) to 2-dimensional data, or extension of methods like S4ND (Nguyen et al. 2022), and 2DSSM (Baron, Zimerman, and Wolf 2024) to have data-dependent weights. Despite the fact that our formulation of the 2D SSM with discretization and data dependent parameters provides a more general framework to extend SSMs to higher-dimensional data, it does not necessarily mean that for any data modalities and network size, its generic form can achieve the best result. While our experimental evaluation is limited to the generic form of our 2D SSM and Chimera, it is a promising future direction to see if limiting transition matrix A_i (i.e., 2D Mamba, 2D Mamba-2) can result in more powerful models. We leave the experimental evaluations of these spacial cases of our 2D SSM for future work.

Efficiency. While our 2D scan decreases the number of required recurrence to compute the hidden states, its still based on a naive implementation of parallel scan. There is a potential for further improvement of 2D parallel scan’s efficiency by using more hardware-aware implementations similar to selective scan by Gu and Dao (2023).

References

- [1] Ahmed Abdulaal, Zhuanghua Liu, and Tomer Lancewicki. “Practical approach to asynchronous multivariate time series anomaly detection and localization”. In: *Proceedings of the 27th ACM SIGKDD conference on knowledge discovery & data mining*. 2021, pp. 2485–2494.
- [2] Masanao Aoki. *State space modeling of time series*. Springer Science & Business Media, 2013.
- [3] Anthony Bagnall, Hoang Anh Dau, Jason Lines, Michael Flynn, James Large, Aaron Bostrom, Paul Southam, and Eamonn Keogh. “The UEA multivariate time series classification archive, 2018”. In: *arXiv preprint arXiv:1811.00075* (2018).
- [4] Ethan Baron, Itamar Zimerman, and Lior Wolf. “A 2-Dimensional State Space Layer for Spatial Inductive Bias”. In: *The Twelfth International Conference on Learning Representations*. 2024. URL: <https://openreview.net/forum?id=BGkqypmGvm>.
- [5] David J Bartholomew. *Time Series Analysis Forecasting and Control*. 1971.

- [6] Ali Behrouz, Parsa Delavari, and Farnoosh Hashemi. “Unsupervised Representation Learning of Brain Activity via Bridging Voxel Activity and Functional Connectivity”. In: *International conference on machine learning (ICML)*. 2024.
- [7] Ali Behrouz and Farnoosh Hashemi. “Graph Mamba: Towards Learning on Graphs with State Space Models”. In: *arXiv preprint arXiv:2402.08678* (2024).
- [8] Ali Behrouz, Michele Santacatterina, and Ramin Zabih. “Mambamixer: Efficient selective state space models with dual token and channel selection”. In: *arXiv preprint arXiv:2403.19888* (2024).
- [9] Michael Bender and Slobodan Simonovic. “Time-series modeling for long-range stream-flow forecasting”. In: *Journal of Water Resources Planning and Management* 120.6 (1994), pp. 857–870.
- [10] George EP Box and Gwilym M Jenkins. “Some recent advances in forecasting and control”. In: *Journal of the Royal Statistical Society. Series C (Applied Statistics)* 17.2 (1968), pp. 91–109.
- [11] Cristian Challu, Kin G Olivares, Boris N Oreshkin, Federico Garza, Max Mergenthaler, and Artur Dubrawski. “N-HiTS: Neural Hierarchical Interpolation for Time Series Forecasting”. In: *arXiv preprint arXiv:2201.12886* (2022).
- [12] Si-An Chen, Chun-Liang Li, Nate Yoder, Sercan O Arik, and Tomas Pfister. “Tsmixer: An all-mlp architecture for time series forecasting”. In: *arXiv preprint arXiv:2303.06053* (2023).
- [13] Chi-Tsong Chen. *Linear system theory and design*. Saunders college publishing, 1984.
- [14] Zonglei Chen, Minbo Ma, Tianrui Li, Hongjun Wang, and Chongshou Li. “Long sequence time-series forecasting with deep learning: A survey”. In: *Information Fusion* 97 (2023), p. 101819.
- [15] Junyoung Chung, Caglar Gulcehre, KyungHyun Cho, and Yoshua Bengio. “Empirical evaluation of gated recurrent neural networks on sequence modeling”. In: *arXiv preprint arXiv:1412.3555* (2014).
- [16] Tri Dao and Albert Gu. “Transformers are SSMS: Generalized Models and Efficient Algorithms Through Structured State Space Duality”. In: *International Conference on Machine Learning (ICML)*. 2024.
- [17] Abhimanyu Das, Weihao Kong, Andrew Leach, Shaan K Mathur, Rajat Sen, and Rose Yu. “Long-term Forecasting with TiDE: Time-series Dense Encoder”. In: *Transactions on Machine Learning Research* (2023). ISSN: 2835-8856. URL: <https://openreview.net/forum?id=pCbC3aQB5W>.
- [18] Rikus Eising. “Realization and stabilization of 2-D systems”. In: *IEEE Transactions on Automatic Control* 23.5 (1978), pp. 793–799.
- [19] Ettore Fornasini and Giovanni Marchesini. “Doubly-indexed dynamical systems: State-space models and structural properties”. In: *Mathematical systems theory* 12.1 (1978), pp. 59–72.
- [20] Jean-Yves Franceschi, Aymeric Dieuleveut, and Martin Jaggi. “Unsupervised Scalable Representation Learning for Multivariate Time Series”. In: *NeurIPS*. 2019.
- [21] Kelum Gajamannage, Yonggi Park, and Dilhani I Jayathilake. “Real-time forecasting of time series in financial markets using sequentially trained dual-LSTMs”. In: *Expert Systems with Applications* 223 (2023), p. 119879.
- [22] Rakshitha Godahewa, Christoph Bergmeir, Geoffrey I Webb, Rob J Hyndman, and Pablo Montero-Manso. “Monash time series forecasting archive”. In: *arXiv preprint arXiv:2105.06643* (2021).
- [23] Albert Gu and Tri Dao. “Mamba: Linear-time sequence modeling with selective state spaces”. In: *arXiv preprint arXiv:2312.00752* (2023).
- [24] Albert Gu, Tri Dao, Stefano Ermon, Atri Rudra, and Christopher Ré. “Hippo: Recurrent memory with optimal polynomial projections”. In: *Advances in neural information processing systems* 33 (2020), pp. 1474–1487.
- [25] Albert Gu, Karan Goel, Ankit Gupta, and Christopher Ré. “On the Parameterization and Initialization of Diagonal State Space Models”. In: *Advances in Neural Information Processing Systems*. Ed. by Alice H. Oh, Alekh Agarwal, Danielle Belgrave, and Kyunghyun Cho. 2022. URL: <https://openreview.net/forum?id=yJE7iQSAep>.
- [26] Albert Gu, Karan Goel, and Christopher Re. “Efficiently Modeling Long Sequences with Structured State Spaces”. In: *International Conference on Learning Representations*. 2022. URL: <https://openreview.net/forum?id=uYLFoz1v1AC>.
- [27] Albert Gu, Caglar Gulcehre, Thomas Paine, Matt Hoffman, and Razvan Pascanu. “Improving the gating mechanism of recurrent neural networks”. In: *International Conference on Machine Learning*. PMLR. 2020, pp. 3800–3809.
- [28] Albert Gu, Isys Johnson, Karan Goel, Khaled Saab, Tri Dao, Atri Rudra, and Christopher Ré. “Combining recurrent, convolutional, and continuous-time models with linear state space layers”. In: *Advances in neural information processing systems* 34 (2021), pp. 572–585.
- [29] Albert Gu, Isys Johnson, Aman Timalsina, Atri Rudra, and Christopher Re. “How to Train your HIPPO: State Space Models with Generalized Orthogonal Basis Projections”. In: *International Conference on Learning Representations*. 2023. URL: <https://openreview.net/forum?id=klk170Q3KB>.
- [30] Andrew C Harvey. “Forecasting, structural time series models and the Kalman filter”. In: *Cambridge university press* (1990).

- [31] Martin N Hebart, Oliver Contier, Lina Teichmann, Adam H Rockter, Charles Y Zheng, Alexis Kidder, Anna Corriveau, Maryam Vaziri-Pashkam, and Chris I Baker. “THINGS-data, a multimodal collection of large-scale datasets for investigating object representations in human brain and behavior”. In: *Elife* 12 (2023), e82580.
- [32] Ts Hinamoto. “Realizations of a state-space model from two-dimensional input-output map”. In: *IEEE Transactions on Circuits and Systems* 27.1 (1980), pp. 36–44.
- [33] S. Hochreiter and J. Schmidhuber. “Long Short-Term Memory”. In: *Neural Comput.* (1997).
- [34] Kyle Hundman, Valentino Constantinou, Christopher Laporte, Ian Colwell, and Tom Soderstrom. “Detecting spacecraft anomalies using lstms and nonparametric dynamic thresholding”. In: *Proceedings of the 24th ACM SIGKDD international conference on knowledge discovery & data mining*. 2018, pp. 387–395.
- [35] Romain Ilbert, Ambroise Odonnat, Vasilii Feofanov, Aladin Virmaux, Giuseppe Paolo, Themis Palpanas, and Ievgen Redko. “Unlocking the Potential of Transformers in Time Series Forecasting with Sharpness-Aware Minimization and Channel-Wise Attention”. In: *arXiv preprint arXiv:2402.10198* (2024).
- [36] Plamen Ch Ivanov, Luis A Nunes Amaral, Ary L Goldberger, Shlomo Havlin, Michael G Rosenblum, Zbigniew R Struzik, and H Eugene Stanley. “Multifractality in human heartbeat dynamics”. In: *Nature* 399.6735 (1999), pp. 461–465.
- [37] Alistair EW Johnson, Lucas Bulgarelli, Lu Shen, Alvin Gayles, Ayad Shammout, Steven Horng, Tom J Pollard, Sicheng Hao, Benjamin Moody, Brian Gow, et al. “MIMIC-IV, a freely accessible electronic health record dataset”. In: *Scientific data* 10.1 (2023), p. 1.
- [38] Angelos Katharopoulos, Apoorv Vyas, Nikolaos Pappas, and François Fleuret. “Transformers are rnns: Fast autoregressive transformers with linear attention”. In: *International conference on machine learning*. PMLR. 2020, pp. 5156–5165.
- [39] Nikita Kitaev, Lukasz Kaiser, and Anselm Levskaya. “Reformer: The Efficient Transformer”. In: *ICLR*. 2020.
- [40] Sun-Yuan Kung, Bernard C Levy, Martin Morf, and Thomas Kailath. “New results in 2-D systems theory, Part II: 2-D state-space models—realization and the notions of controllability, observability, and minimality”. In: *Proceedings of the IEEE* 65.6 (1977), pp. 945–961.
- [41] Guokun Lai, Wei-Cheng Chang, Yiming Yang, and Hanxiao Liu. “Modeling long-and short-term temporal patterns with deep neural networks”. In: *SIGIR*. 2018.
- [42] Shiyang Li, Xiaoyong Jin, Yao Xuan, Xiyu Zhou, Wenhui Chen, Yu-Xiang Wang, and Xifeng Yan. “Enhancing the Locality and Breaking the Memory Bottleneck of Transformer on Time Series Forecasting”. In: *NeurIPS*. 2019.
- [43] Zhe Li, Shiyi Qi, Yiduo Li, and Zenglin Xu. “Revisiting long-term time series forecasting: An investigation on linear mapping”. In: *arXiv preprint arXiv:2305.10721* (2023).
- [44] Bryan Lim and Stefan Zohren. “Time-series forecasting with deep learning: a survey”. In: *Philosophical Transactions of the Royal Society A* 379.2194 (2021), p. 20200209.
- [45] Minhao Liu, Ailing Zeng, Muxi Chen, Zhijian Xu, Qiuxia Lai, Lingna Ma, and Qiang Xu. “Scinet: Time series modeling and forecasting with sample convolution and interaction”. In: *Advances in Neural Information Processing Systems* 35 (2022), pp. 5816–5828.
- [46] Shizhan Liu, Hang Yu, Cong Liao, Jianguo Li, Weiyao Lin, Alex X Liu, and Schahram Dustdar. “Pyraformer: Low-complexity pyramidal attention for long-range time series modeling and forecasting”. In: *International conference on learning representations*. 2021.
- [47] Yong Liu, Tengge Hu, Haoran Zhang, Haixu Wu, Shiyu Wang, Lintao Ma, and Mingsheng Long. “iTransformer: Inverted Transformers Are Effective for Time Series Forecasting”. In: *The Twelfth International Conference on Learning Representations*. 2024. URL: <https://openreview.net/forum?id=JePFAI8fah>.
- [48] Yong Liu, Haixu Wu, Jianmin Wang, and Mingsheng Long. “Non-stationary transformers: Exploring the stationarity in time series forecasting”. In: *Advances in Neural Information Processing Systems* 35 (2022), pp. 9881–9893.
- [49] Yong Liu, Haixu Wu, Jianmin Wang, and Mingsheng Long. “Non-stationary Transformers: Rethinking the Stationarity in Time Series Forecasting”. In: *NeurIPS*. 2022.
- [50] Yue Liu, Yunjie Tian, Yuzhong Zhao, Hongtian Yu, Lingxi Xie, Yaowei Wang, Qixiang Ye, and Yunfan Liu. “Vmamba: Visual state space model”. In: *arXiv preprint arXiv:2401.10166* (2024).
- [51] Donghao Luo and Xue Wang. “ModernTCN: A Modern Pure Convolution Structure for General Time Series Analysis”. In: *The Twelfth International Conference on Learning Representations*. 2024. URL: <https://openreview.net/forum?id=vpJMJerXHU>.
- [52] Jun Ma, Feifei Li, and Bo Wang. “U-Mamba: Enhancing Long-range Dependency for Biomedical Image Segmentation”. In: *arXiv preprint arXiv:2401.04722* (2024).

- [53] Eric Martin and Chris Cundy. “Parallelizing Linear Recurrent Neural Nets Over Sequence Length”. In: *International Conference on Learning Representations*. 2018. URL: <https://openreview.net/forum?id=HyUNwulC->.
- [54] Aditya P Mathur and Nils Ole Tippenhauer. “SWaT: A water treatment testbed for research and training on ICS security”. In: *2016 international workshop on cyber-physical systems for smart water networks (CySWater)*. IEEE. 2016, pp. 31–36.
- [55] Eric Nguyen, Karan Goel, Albert Gu, Gordon Downs, Preey Shah, Tri Dao, Stephen Baccus, and Christopher Ré. “S4nd: Modeling images and videos as multidimensional signals with state spaces”. In: *Advances in neural information processing systems* 35 (2022), pp. 2846–2861.
- [56] Yuqi Nie, Nam H Nguyen, Phanwadee Sinthong, and Jayant Kalagnanam. “A Time Series is Worth 64 Words: Long-term Forecasting with Transformers”. In: *The Eleventh International Conference on Learning Representations*. 2023. URL: <https://openreview.net/forum?id=Jbdc0vT0col>.
- [57] Boris N Oreshkin, Dmitri Carпов, Nicolas Chapados, and Yoshua Bengio. “N-BEATS: Neural basis expansion analysis for interpretable time series forecasting”. In: *ICLR* (2019).
- [58] Badri N. Patro and Vijay S. Agneeswaran. *SiMBA: Simplified Mamba-Based Architecture for Vision and Multivariate Time series*. 2024. arXiv: [2403.15360](https://arxiv.org/abs/2403.15360) [cs.CV].
- [59] Badri Narayana Patro and Vijay Srinivas Agneeswaran. “Mamba-360: Survey of State Space Models as Transformer Alternative for Long Sequence Modelling: Methods, Applications, and Challenges”. In: *arXiv preprint arXiv:2404.16112* (2024).
- [60] Steve Pincus and Rudolf E Kalman. “Irregularity, volatility, risk, and financial market time series”. In: *Proceedings of the National Academy of Sciences* 101.38 (2004), pp. 13709–13714.
- [61] Zhen Qin, Songlin Yang, and Yiran Zhong. “Hierarchically gated recurrent neural network for sequence modeling”. In: *Advances in Neural Information Processing Systems* 36 (2023).
- [62] Prajit Ramachandran, Barret Zoph, and Quoc V Le. “Searching for activation functions”. In: *arXiv preprint arXiv:1710.05941* (2017).
- [63] David Salinas, Valentin Flunkert, Jan Gasthaus, and Tim Januschowski. “DeepAR: Probabilistic forecasting with autoregressive recurrent networks”. In: *International journal of forecasting* 36.3 (2020), pp. 1181–1191.
- [64] Yair Schiff, Chia-Hsiang Kao, Aaron Gokaslan, Tri Dao, Albert Gu, and Volodymyr Kuleshov. “Caduceus: Bi-directional equivariant long-range dna sequence modeling”. In: *arXiv preprint arXiv:2403.03234* (2024).
- [65] Imanol Schlag, Kazuki Irie, and Jürgen Schmidhuber. “Linear transformers are secretly fast weight programmers”. In: *International Conference on Machine Learning*. PMLR. 2021, pp. 9355–9366.
- [66] Jimmy T.H. Smith, Andrew Warrington, and Scott Linderman. “Simplified State Space Layers for Sequence Modeling”. In: *The Eleventh International Conference on Learning Representations*. 2023. URL: <https://openreview.net/forum?id=Ai8Hw3AXqks>.
- [67] Ya Su, Youjian Zhao, Chenhao Niu, Rong Liu, Wei Sun, and Dan Pei. “Robust anomaly detection for multivariate time series through stochastic recurrent neural network”. In: *Proceedings of the 25th ACM SIGKDD international conference on knowledge discovery & data mining*. 2019, pp. 2828–2837.
- [68] Corentin Tallec and Yann Ollivier. “Can recurrent neural networks warp time?” In: *International Conference on Learning Representations*. 2018. URL: <https://openreview.net/forum?id=SJcKhk-Ab>.
- [69] William Toner and Luke Darlow. “An Analysis of Linear Time Series Forecasting Models”. In: *International conference on machine learning (ICML)* (2024).
- [70] Hugo Touvron, Thibaut Lavril, Gautier Izacard, Xavier Martinet, Marie-Anne Lachaux, Timothée Lacroix, Baptiste Rozière, Naman Goyal, Eric Hambro, Faisal Azhar, et al. “Llama: Open and efficient foundation language models”. In: *arXiv preprint arXiv:2302.13971* (2023).
- [71] Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N Gomez, Łukasz Kaiser, and Illia Polosukhin. “Attention is all you need”. In: *Advances in neural information processing systems* 30 (2017).
- [72] Patrick Wagner, Nils Strodthoff, Ralf-Dieter Boussejot, Dieter Kreiseler, Fatima I Lunze, Wojciech Samek, and Tobias Schaeffter. “PTB-XL, a large publicly available electrocardiography dataset”. In: *Scientific data* 7.1 (2020), pp. 1–15.
- [73] Junxiong Wang, Jing Nathan Yan, Albert Gu, and Alexander Rush. “Pretraining Without Attention”. In: *Findings of the Association for Computational Linguistics: EMNLP 2023*. Ed. by Houda Bouamor, Juan Pino, and Kalika Bali. Singapore: Association for Computational Linguistics, Dec. 2023, pp. 58–69. DOI: [10.18653/v1/2023.findings-emnlp.5](https://doi.org/10.18653/v1/2023.findings-emnlp.5). URL: <https://aclanthology.org/2023.findings-emnlp.5>.
- [74] Pete Warden. “Speech commands: A dataset for limited-vocabulary speech recognition”. In: *arXiv preprint arXiv:1804.03209* (2018).

- [75] Peter R Winters. “Forecasting sales by exponentially weighted moving averages”. In: *Management science* 6.3 (1960), pp. 324–342.
- [76] Gerald Woo, Chenghao Liu, Doyen Sahoo, Akshat Kumar, and Steven C. H. Hoi. “ETSformer: Exponential Smoothing Transformers for Time-series Forecasting”. In: *arXiv preprint arXiv:2202.01381* (2022).
- [77] Haixu Wu, Tengge Hu, Yong Liu, Hang Zhou, Jianmin Wang, and Mingsheng Long. “TimesNet: Temporal 2D-Variation Modeling for General Time Series Analysis”. In: *The Eleventh International Conference on Learning Representations*. 2023. URL: https://openreview.net/forum?id=ju_Uqw384Oq.
- [78] Haixu Wu, Jialong Wu, Jiehui Xu, Jianmin Wang, and Mingsheng Long. “Flowformer: Linearizing Transformers with Conservation Flows”. In: *ICML*. 2022.
- [79] Haixu Wu, Jiehui Xu, Jianmin Wang, and Mingsheng Long. “Autoformer: Decomposition transformers with auto-correlation for long-term series forecasting”. In: *Advances in neural information processing systems* 34 (2021), pp. 22419–22430.
- [80] Jiehui Xu, Haixu Wu, Jianmin Wang, and Mingsheng Long. “Anomaly Transformer: Time Series Anomaly Detection with Association Discrepancy”. In: *ICLR*. 2021.
- [81] Songlin Yang, Bailin Wang, Yikang Shen, Rameswar Panda, and Yoon Kim. “Gated linear attention transformers with hardware-efficient training”. In: *International conference on machine learning (ICML)*. 2024.
- [82] Ailing Zeng, Muxi Chen, Lei Zhang, and Qiang Xu. “Are Transformers Effective for Time Series Forecasting?” In: *AAAI*. 2023.
- [83] Ailing Zeng, Muxi Chen, Lei Zhang, and Qiang Xu. “Are transformers effective for time series forecasting?” In: *Proceedings of the AAAI conference on artificial intelligence*. Vol. 37. 2023, pp. 11121–11128.
- [84] Michael Zhang, Khaled Kamal Saab, Michael Poli, Tri Dao, Karan Goel, and Christopher Re. “Effectively Modeling Time Series with Simple Discrete State Spaces”. In: *The Eleventh International Conference on Learning Representations*. 2023. URL: <https://openreview.net/forum?id=2EpjkjzdCAa>.
- [85] T. Zhang, Yizhuo Zhang, Wei Cao, J. Bian, Xiaohan Yi, Shun Zheng, and Jian Li. “Less Is More: Fast Multivariate Time Series Forecasting with Light Sampling-oriented MLP Structures”. In: *arXiv preprint arXiv:2207.01186* (2022).
- [86] Yunhao Zhang and Junchi Yan. “Crossformer: Transformer utilizing cross-dimension dependency for multivariate time series forecasting”. In: *The eleventh international conference on learning representations*. 2023.
- [87] Haoyi Zhou, Shanghang Zhang, Jieqi Peng, Shuai Zhang, Jianxin Li, Hui Xiong, and Wancai Zhang. “Informer: Beyond efficient transformer for long sequence time-series forecasting”. In: *Proceedings of the AAAI conference on artificial intelligence*. Vol. 35. 2021, pp. 11106–11115.
- [88] Tian Zhou, Ziqing Ma, Qingsong Wen, Liang Sun, Tao Yao, Wotao Yin, Rong Jin, et al. “Film: Frequency improved legendre memory model for long-term time series forecasting”. In: *Advances in Neural Information Processing Systems* 35 (2022), pp. 12677–12690.
- [89] Tian Zhou, Ziqing Ma, Qingsong Wen, Xue Wang, Liang Sun, and Rong Jin. “Fedformer: Frequency enhanced decomposed transformer for long-term series forecasting”. In: *International conference on machine learning*. PMLR. 2022, pp. 27268–27286.

A Background

A.1 1D Space State Models

1D Space State Models (SSMs) are linear time-invariant systems that map input sequence $x(t) \in \mathbb{R}^L \mapsto y(t) \in \mathbb{R}^L$ (Aoki 2013). SSMs use a latent state $h(t) \in \mathbb{R}^{N \times L}$, transition parameter $\mathbf{A} \in \mathbb{R}^{N \times N}$, and projection parameters $\mathbf{B} \in \mathbb{R}^{N \times 1}$, $\mathbf{C} \in \mathbb{R}^{1 \times N}$ to model the input and output as:

$$h'(t) = \mathbf{A} h(t) + \mathbf{B} x(t), \quad y(t) = \mathbf{C} h(t). \quad (26)$$

Most existing SSMs (Behrouz, Santacatterina, and Zabih 2024; Gu and Dao 2023; Gu, Goel, and Re 2022), first discretize the signals \mathbf{A} , \mathbf{B} , and \mathbf{C} . That is, using a parameter Δ and zero-order hold, the discretized formulation is defined as:

$$h_t = \bar{\mathbf{A}} h_{t-1} + \bar{\mathbf{B}} x_t, \quad y_t = \mathbf{C} h_t, \quad (27)$$

where $\bar{\mathbf{A}} = \exp(\Delta \mathbf{A})$ and $\bar{\mathbf{B}} = (\Delta \mathbf{A})^{-1} (\exp(\Delta \mathbf{A}) - \mathbf{I}) \cdot \Delta \mathbf{B}$. (Gu, Dao, et al. 2020) show that discrete SSMs can be interpreted as both convolutions and recurrent networks: i.e.,

$$\begin{aligned} \bar{\mathbf{K}} &= (\mathbf{C} \bar{\mathbf{B}}, \mathbf{C} \bar{\mathbf{A}} \bar{\mathbf{B}}, \dots, \mathbf{C} \bar{\mathbf{A}}^{L-1} \bar{\mathbf{B}}), \\ y &= x * \bar{\mathbf{K}}, \end{aligned} \quad (28)$$

which makes their training and inference very efficient as a convolution and recurrent model, respectively.

A.2 Data Dependency

Above discrete SSMs are based on data-independent parameters. That is, parameters Δ , $\bar{\mathbf{A}}$, $\bar{\mathbf{B}}$, and \mathbf{C} are time invariant and are the same for any input. Gu and Dao (2023) argue that this time invariance has the cost of limiting SSMs effectiveness in compressing context into a smaller state (Gu and Dao 2023). To overcome this challenge, they present a selective SSMs (S6) block that effectively selects relevant context by enabling dependence of the parameters $\bar{\mathbf{B}}$, $\bar{\mathbf{C}}$, and Δ on the input x_t , i.e.:

$$\bar{\mathbf{B}}_t = \text{Linear}_{\mathbf{B}}(x_t) \quad (29)$$

$$\bar{\mathbf{C}}_t = \text{Linear}_{\mathbf{C}}(x_t) \quad (30)$$

$$\Delta_t = \text{Softplus}(\text{Linear}_{\Delta}(x_t)), \quad (31)$$

where $\text{Linear}(\cdot)$ is a linear projection and $\text{Softplus}(\cdot) = \log(1 + \exp(\cdot))$. This data dependency comes at the cost of efficiency as the model cannot be trained as a convolution. To overcome this challenge, Gu and Dao (2023) show that the linear recurrence in Equation 1 can be formulated as an associative scan (Martin and Cundy 2018), which accepts efficient parallel algorithms.

B Additional Related Work

Classical Approach. Modeling time series data is a long-standing problem and has attracted much attention during the past 60 years. There have been several mathematical models to capture the time series traits like exponential smoothing (Winters 1960), autoregressive integrated moving average (ARIMA) (Bartholomew 1971), SARIMA (Bender and Simonovic 1994), Box-Jenkins method (Box and Jenkins 1968), and more recently state-space models (Aoki 2013; Harvey 1990). Despite their more interpretability, these methods usually fail to capture non-linear dependencies and also often require manually analyzing time series features (e.g., trend or seasonality), resulting in lack of generalizability.

Recurrent and Deep State Space Models. Another group of relevant studies to ours is deep sequence models. A common class of architectures for sequence modeling are recurrent neural networks such as like GRUs (Chung et al. 2014), DeepAR (Salinas et al. 2020), LSTMs (Hochreiter and J. Schmidhuber 1997). The main drawback of RNNs is their potential for vanishing/exploding gradients and also their slow training. Recently, linear attention methods with fast training attracted attention (Katharopoulos et al. 2020; Schlag, Irie, and Jürgen Schmidhuber 2021; Yang et al. 2024). Katharopoulos et al. (2020) show that these methods have recurrent formulation and can be fast in inference.

Recently, deep state space models have attracted much attention as the alternative of Transformers (Vaswani et al. 2017), due to their fast training and inference (Gu, Dao, et al. 2020). These methods are the combination of traditional SSMs with deep neural networks by directly parameterizing the layers of a neural network with multiple linear SSMs, and overcome common recurrent training drawbacks by leveraging the convolutional view of SSMs (Gu, Dao, et al. 2020; Gu, Goel, Gupta, et al. 2022; Gu, Goel, and Re 2022; Gu, I. Johnson, Goel, et al. 2021; Smith, Warrington, and Linderman 2023). Recently, Gu and Dao (2023) present a new formulation of deep SSMs by allowing the parameters to be the function of inputs. This architecture shows promising potential in various domains like NLP (Gu and Dao 2023), vision (Behrouz, Santacatterina, and Zabih 2024; Yue Liu et al. 2024; J. Ma, F. Li, and B. Wang 2024), graphs (Behrouz and Hashemi 2024), DNA modeling (Gu and Dao 2023; Schiff et al. 2024).

All the above methods are design for 1D data, meaning that the states depends on one variable. There are, however, a few studies that uses 2D SSMs in deep learning settings. S4ND (Nguyen et al. 2022) uses continuous signals to model images. These methods not only consider two separate SSM for the axes, but it also directly treat the system as a continuous system without discretization step. Furthermore, S4ND has data-independent parameters. Another similar approach is 2DSSM (Baron, Zimerman, and Wolf 2024), that models images as discrete signals. That is, the initial SSM model is discrete and again there is a lack of discretization step, which is important for time series as we discussed earlier. Also, their method again is based on data-independent parameters. Both S4ND and 2DSSM can be computed as a convolution. We, however, present a new scanning technique for fast training of 2D SSMs, even with input-dependent parameters.

Other methods. Transformer-based models have attracted much attention over recent years for multivariate time series forecasting, when modeling the complex relationships of co-variables or along the time dimension is required (Ilbert et al. 2024; Kitaev, Kaiser, and Levskaya 2020; S. Liu et al. 2021; Nie et al. 2023; H. Wu, Xu, et al. 2021; Zeng et al. 2023a; Y. Zhang and Yan 2023; H. Zhou et al. 2021; T. Zhou, Z. Ma, Wen, X. Wang, et al. 2022). Several studies have focused on designing more efficient and effective attentions with using special traits of time series (Woo et al. 2022). Some other studies have focused on extracting long-term information for better forecasting (Nie et al. 2023; T. Zhou, Z. Ma, Wen, Sun, et al. 2022). In addition to transformers, linear models also have shown promising results (S.-A. Chen et al. 2023; H. Wu, Hu, et al. 2023). For example, S.-A. Chen et al. (2023) present TSMixer, an all-MLP architecture for time series forecasting, with promising performance. Due to the expressive power of our 2D SSM, these linear methods sometimes can be viewed as a special case of 2D SSMs. Recently, convolution-based models for time series have shown promising results (Luo and X. Wang 2024). These methods by using global kernels enhance the global receptive field. Our data-independent formulation of Chimera is connected to this line of work as it can be written as a global convolution.

C Details of the Discretization

Given PDE with initial condition $h(0, 0) = 0$:

$$\frac{\partial}{\partial t^{(1)}} h^{(1)}(t^{(1)}, t^{(2)}) = \left(\mathbf{A}_1 h^{(1)}(t^{(1)}, t^{(2)}), \mathbf{A}_2 h^{(2)}(t^{(1)}, t^{(2)}) \right) + \mathbf{B}_1 \mathbf{x}(t^{(1)}, t^{(2)}), \quad (32)$$

$$\frac{\partial}{\partial t^{(1)}} h^{(2)}(t^{(1)}, t^{(2)}) = \left(\mathbf{A}_1 h^{(1)}(t^{(1)}, t^{(2)}), \mathbf{A}_2 h^{(2)}(t^{(1)}, t^{(2)}) \right) + \mathbf{B}_1 \mathbf{x}(t^{(1)}, t^{(2)}), \quad (33)$$

$$\frac{\partial}{\partial t^{(2)}} h^{(1)}(t^{(1)}, t^{(2)}) = \left(\mathbf{A}_3 h^{(1)}(t^{(1)}, t^{(2)}), \mathbf{A}_4 h^{(2)}(t^{(1)}, t^{(2)}) \right) + \mathbf{B}_2 \mathbf{x}(t^{(1)}, t^{(2)}), \quad (34)$$

$$\frac{\partial}{\partial t^{(2)}} h^{(2)}(t^{(1)}, t^{(2)}) = \left(\mathbf{A}_3 h^{(1)}(t^{(1)}, t^{(2)}), \mathbf{A}_4 h^{(2)}(t^{(1)}, t^{(2)}) \right) + \mathbf{B}_2 \mathbf{x}(t^{(1)}, t^{(2)}), \quad (35)$$

over the sampling intervals $[k\Delta t^{(1)}, (k+1)\Delta t^{(1)}]$ and $[\ell\Delta t^{(2)}, (\ell+1)\Delta t^{(2)}]$ we have:

$$\begin{aligned} & \int_{k\Delta t^{(1)}}^{(k+1)\Delta t^{(1)}} \frac{\partial}{\partial t^{(1)}} h^{(1)}(t^{(1)}, t^{(2)}) dt^{(1)} \\ &= \int_{k\Delta t^{(1)}}^{(k+1)\Delta t^{(1)}} \left(\mathbf{A}_1 h^{(1)}(t^{(1)}, t^{(2)}) + \mathbf{B}_1^{(1)} \mathbf{x}^{(1)}(t^{(1)}, t^{(2)}) \right) dt^{(1)} \end{aligned} \quad (36)$$

and so:

$$\begin{aligned}
& \int_{k\Delta t^{(1)}}^{(k+1)\Delta t^{(1)}} \frac{\partial}{\partial t^{(1)}} h^{(2)} \left(t^{(1)}, t^{(2)} \right) dt^{(1)} \\
&= \int_{k\Delta t^{(1)}}^{(k+1)\Delta t^{(1)}} \left(\mathbf{A}_2 h^{(2)} \left(t^{(1)}, t^{(2)} \right) + \mathbf{B}_1^{(2)} \mathbf{x}^{(2)} \left(t^{(1)}, t^{(2)} \right) \right) dt^{(1)}
\end{aligned} \tag{37}$$

Similarly, for the second equation we have:

$$\begin{aligned}
& \int_{\ell\Delta t^{(2)}}^{(\ell+1)\Delta t^{(2)}} \frac{\partial}{\partial t^{(2)}} h^{(1)} \left(t^{(1)}, t^{(2)} \right) dt^{(2)} \\
&= \int_{\ell\Delta t^{(2)}}^{(\ell+1)\Delta t^{(2)}} \left(\mathbf{A}_3 h^{(1)} \left(t^{(1)}, t^{(2)} \right) + \mathbf{B}_2^{(1)} \mathbf{x}^{(1)} \left(t^{(1)}, t^{(2)} \right) \right) dt^{(2)}
\end{aligned} \tag{38}$$

and so:

$$\begin{aligned}
& \int_{\ell\Delta t^{(2)}}^{(\ell+1)\Delta t^{(2)}} \frac{\partial}{\partial t^{(2)}} h^{(2)} \left(t^{(1)}, t^{(2)} \right) dt^{(2)} \\
&= \int_{\ell\Delta t^{(2)}}^{(\ell+1)\Delta t^{(2)}} \left(\mathbf{A}_4 h^{(2)} \left(t^{(1)}, t^{(2)} \right) + \mathbf{B}_2^{(2)} \mathbf{x}^{(2)} \left(t^{(1)}, t^{(2)} \right) \right) dt^{(2)}
\end{aligned} \tag{39}$$

Next, the integrals can be simplified as:

$$\begin{aligned}
& h^{(1)} \left((k+1)\Delta t^{(1)}, t^{(2)} \right) \\
&= e^{\mathbf{A}_1 \Delta t^{(1)}} h^{(1)} \left(k\Delta t^{(1)}, t^{(2)} \right) + \int_{k\Delta t^{(1)}}^{(k+1)\Delta t^{(1)}} e^{\mathbf{A}_1(t^{(1)}-k\Delta t^{(1)})} \mathbf{B}_1^{(1)} \mathbf{x}^{(1)} \left(t^{(1)}, t^{(2)} \right) dt^{(1)},
\end{aligned} \tag{40}$$

and

$$\begin{aligned}
& h^{(2)} \left((k+1)\Delta t^{(1)}, t^{(2)} \right) \\
&= e^{\mathbf{A}_2 \Delta t^{(1)}} h^{(2)} \left(k\Delta t^{(1)}, t^{(2)} \right) + \int_{k\Delta t^{(1)}}^{(k+1)\Delta t^{(1)}} e^{\mathbf{A}_2(t^{(1)}-k\Delta t^{(1)})} \mathbf{B}_1^{(2)} \mathbf{x}^{(2)} \left(t^{(1)}, t^{(2)} \right) dt^{(1)},
\end{aligned} \tag{41}$$

and similarly for the third and fourth equations we have:

$$\begin{aligned}
& h^{(1)} \left(t^{(1)}, (\ell+1)\Delta t^{(2)} \right) \\
&= e^{\mathbf{A}_3 \Delta t^{(2)}} h^{(1)} \left(t^{(1)}, \ell\Delta t^{(2)} \right) + \int_{\ell\Delta t^{(2)}}^{(\ell+1)\Delta t^{(2)}} e^{\mathbf{A}_3(t^{(2)}-\ell\Delta t^{(2)})} \mathbf{B}_2^{(1)} \mathbf{x}^{(1)} \left(t^{(2)}, t^{(1)} \right) dt^{(2)}
\end{aligned} \tag{42}$$

and

$$\begin{aligned}
& h^{(2)} \left(t^{(1)}, (\ell+1)\Delta t^{(2)} \right) \\
&= e^{\mathbf{A}_4 \Delta t^{(2)}} h^{(2)} \left(t^{(1)}, \ell\Delta t^{(2)} \right) + \int_{\ell\Delta t^{(2)}}^{(\ell+1)\Delta t^{(2)}} e^{\mathbf{A}_4(t^{(2)}-\ell\Delta t^{(2)})} \mathbf{B}_2^{(2)} \mathbf{x}^{(2)} \left(t^{(2)}, t^{(1)} \right) dt^{(2)}
\end{aligned} \tag{43}$$

Using ZOH assumption, we have:

$$\int_0^{\Delta t^{(1)}} e^{\mathbf{A}_1 s} ds = \mathbf{A}^{(1)-1} \left(e^{\mathbf{A}_1 \Delta t^{(1)}} - \mathbf{I} \right)$$

$$\int_0^{\Delta t^{(1)}} e^{\mathbf{A}_2 s} ds = \mathbf{A}^{(2)-1} \left(e^{\mathbf{A}_2 \Delta t^{(1)}} - \mathbf{I} \right) \quad (44)$$

$$\int_0^{\Delta t^{(2)}} e^{\mathbf{A}_3 s} ds = \mathbf{A}^{(3)-1} \left(e^{\mathbf{A}_3 \Delta t^{(2)}} - \mathbf{I} \right) \quad (45)$$

$$\int_0^{\Delta t^{(2)}} e^{\mathbf{A}_4 s} ds = \mathbf{A}^{(4)-1} \left(e^{\mathbf{A}_4 \Delta t^{(2)}} - \mathbf{I} \right) \quad (46)$$

Accordingly, the discretized form is as follows:

$$h_{k+1,\ell}^{(1)} = e^{\mathbf{A}_1 \Delta t^{(1)}} h_{k,\ell}^{(1)} + \mathbf{A}^{(1)-1} \left(e^{\mathbf{A}_1 \Delta t^{(1)}} - \mathbf{I} \right) \mathbf{B}_1^{(1)} \mathbf{x}_{k+1,\ell}^{(1)} \quad (47)$$

$$h_{k+1,\ell}^{(2)} = e^{\mathbf{A}_2 \Delta t^{(1)}} h_{k,\ell}^{(2)} + \mathbf{A}^{(2)-1} \left(e^{\mathbf{A}_2 \Delta t^{(1)}} - \mathbf{I} \right) \mathbf{B}_1^{(2)} \mathbf{x}_{k+1,\ell}^{(2)} \quad (48)$$

$$h_{k,\ell+1}^{(1)} = e^{\mathbf{A}_3 \Delta t^{(2)}} h_{k,\ell}^{(1)} + \mathbf{A}^{(3)-1} \left(e^{\mathbf{A}_3 \Delta t^{(2)}} - \mathbf{I} \right) \mathbf{B}_2^{(1)} \mathbf{x}_{k,\ell+1}^{(1)} \quad (49)$$

$$h_{k,\ell+1}^{(2)} = e^{\mathbf{A}_4 \Delta t^{(2)}} h_{k,\ell}^{(2)} + \mathbf{A}^{(4)-1} \left(e^{\mathbf{A}_4 \Delta t^{(2)}} - \mathbf{I} \right) \mathbf{B}_2^{(2)} \mathbf{x}_{k,\ell+1}^{(2)}, \quad (50)$$

which means that:

$$\bar{\mathbf{A}}_1 = \exp(\mathbf{A}_1 \Delta_1), \quad (51)$$

$$\bar{\mathbf{A}}_2 = \exp(\mathbf{A}_2 \Delta_1), \quad (52)$$

$$\bar{\mathbf{A}}_3 = \exp(\mathbf{A}_3 \Delta_2), \quad (53)$$

$$\bar{\mathbf{A}}_4 = \exp(\mathbf{A}_4 \Delta_2), \quad (54)$$

$$(55)$$

and

$$\bar{\mathbf{B}}_1 = \begin{bmatrix} \mathbf{A}^{(1)-1} (e^{\mathbf{A}_1 \Delta_1} - \mathbf{I}) \mathbf{B}_1^{(1)} \\ \mathbf{A}^{(2)-1} (e^{\mathbf{A}_2 \Delta_1} - \mathbf{I}) \mathbf{B}_1^{(2)} \end{bmatrix}, \quad (56)$$

$$\bar{\mathbf{B}}_2 = \begin{bmatrix} \mathbf{A}^{(3)-1} (e^{\mathbf{A}_3 \Delta_2} - \mathbf{I}) \mathbf{B}_2^{(1)} \\ \mathbf{A}^{(4)-1} (e^{\mathbf{A}_4 \Delta_2} - \mathbf{I}) \mathbf{B}_2^{(2)} \end{bmatrix}. \quad (57)$$

D Details of the Structure of Transition Matrices

Definition D.1 (Companion Matrix). *A matrix $A \in \mathbb{R}^{N \times N}$ has companion form if it can be written as:*

$$A = \begin{pmatrix} 0 & 0 & \dots & 0 & a_1 \\ 1 & 0 & \dots & 0 & a_2 \\ 0 & 1 & \dots & 0 & a_3 \\ \vdots & \vdots & \ddots & \vdots & \vdots \\ 0 & 0 & \dots & 0 & a_{N_1} \\ 0 & 0 & \dots & 1 & a_N \end{pmatrix}. \quad (58)$$

These matrices can be decompose into a shift and a low-rank matrix. That is:

$$A = \begin{pmatrix} 0 & 0 & \dots & 0 & a_1 \\ 1 & 0 & \dots & 0 & a_2 \\ 0 & 1 & \dots & 0 & a_3 \\ \vdots & \vdots & \ddots & \vdots & \vdots \\ 0 & 0 & \dots & 0 & a_{N_1} \\ 0 & 0 & \dots & 1 & a_N \end{pmatrix} = \underbrace{\begin{pmatrix} 0 & 0 & \dots & 0 & 0 \\ 1 & 0 & \dots & 0 & 0 \\ 0 & 1 & \dots & 0 & 0 \\ \vdots & \vdots & \ddots & \vdots & \vdots \\ 0 & 0 & \dots & 0 & 0 \\ 0 & 0 & \dots & 1 & 0 \end{pmatrix}}_{\text{Shift Matrix}} + \underbrace{\begin{pmatrix} 0 & 0 & \dots & 0 & a_1 \\ 0 & 0 & \dots & 0 & a_2 \\ 0 & 0 & \dots & 0 & a_3 \\ \vdots & \vdots & \ddots & \vdots & \vdots \\ 0 & 0 & \dots & 0 & a_{N_1} \\ 0 & 0 & \dots & 0 & a_N \end{pmatrix}}_{\text{Low-rank Matrix}}. \quad (59)$$

This formulation can help us to compute the power of A faster in the convolutional form, as discussed by M. Zhang et al. (2023).

E Theoretical Results

E.1 Proof of Theorem 3.2

In this part, we want to prove that \ast is associative. This operator is defined as:

$$p \ast q = \begin{pmatrix} p_1 & p_2 & p_3 \\ p_4 & p_5 & p_6 \end{pmatrix} \ast \begin{pmatrix} q_1 & q_2 & q_3 \\ q_4 & q_5 & q_6 \end{pmatrix} = \begin{pmatrix} q_1 \odot p_1 & q_2 \odot p_2 & q_1 \otimes p_3 + q_2 \otimes p_6 + q_3 \\ q_4 \odot p_4 & q_5 \odot p_5 & q_4 \otimes p_3 + q_5 \otimes p_6 + q_6 \end{pmatrix}$$

Accordingly, we have:

$$(p \ast q) \ast r = \begin{pmatrix} q_1 \odot p_1 & q_2 \odot p_2 & q_1 \otimes p_3 + q_2 \otimes p_6 + q_3 \\ q_4 \odot p_4 & q_5 \odot p_5 & q_4 \otimes p_3 + q_5 \otimes p_6 + q_6 \end{pmatrix} \ast \begin{pmatrix} r_1 & r_2 & r_3 \\ r_4 & r_5 & r_6 \end{pmatrix}, \quad (60)$$

re-using the definition of \ast , we have:

$$(p \ast q) \ast r = \begin{pmatrix} q_1 \odot p_1 & q_2 \odot p_2 & q_1 \otimes p_3 + q_2 \otimes p_6 + q_3 \\ q_4 \odot p_4 & q_5 \odot p_5 & q_4 \otimes p_3 + q_5 \otimes p_6 + q_6 \end{pmatrix} \ast \begin{pmatrix} r_1 & r_2 & r_3 \\ r_4 & r_5 & r_6 \end{pmatrix} \quad (61)$$

$$= \begin{pmatrix} r_1 \odot (q_1 \odot p_1) & r_2 \odot (q_2 \odot p_2) & r_1 \otimes (q_1 \otimes p_3 + q_2 \otimes p_6 + q_3) + r_2 \odot (q_4 \otimes p_3 + q_5 \otimes p_6 + q_6) + r_3 \\ r_4 \odot (q_4 \odot p_4) & r_5 \odot (q_5 \odot p_5) & r_4 \otimes (q_1 \otimes p_3 + q_2 \otimes p_6 + q_3) + r_4 \otimes (q_4 \otimes p_3 + q_5 \otimes p_6 + q_6) + r_6 \end{pmatrix} \quad (62)$$

Using the fact that \odot and \otimes are associative, we have:

$$(p \ast q) \ast r = \begin{pmatrix} q_1 \odot p_1 & q_2 \odot p_2 & q_1 \otimes p_3 + q_2 \otimes p_6 + q_3 \\ q_4 \odot p_4 & q_5 \odot p_5 & q_4 \otimes p_3 + q_5 \otimes p_6 + q_6 \end{pmatrix} \ast \begin{pmatrix} r_1 & r_2 & r_3 \\ r_4 & r_5 & r_6 \end{pmatrix} \quad (63)$$

$$= \begin{pmatrix} r_1 \odot (q_1 \odot p_1) & r_2 \odot (q_2 \odot p_2) & r_1 \otimes (q_1 \otimes p_3 + q_2 \otimes p_6 + q_3) + r_2 \odot (q_4 \otimes p_3 + q_5 \otimes p_6 + q_6) + r_3 \\ r_4 \odot (q_4 \odot p_4) & r_5 \odot (q_5 \odot p_5) & r_4 \otimes (q_1 \otimes p_3 + q_2 \otimes p_6 + q_3) + r_4 \otimes (q_4 \otimes p_3 + q_5 \otimes p_6 + q_6) + r_6 \end{pmatrix} \quad (64)$$

$$= \begin{pmatrix} p_1 & p_2 & p_3 \\ p_4 & p_5 & p_6 \end{pmatrix} \ast \begin{pmatrix} r_1 \odot q_1 & r_2 \odot q_2 & r_1 \otimes q_3 + r_2 \otimes q_6 + r_3 \\ r_4 \odot q_4 & r_5 \odot q_5 & r_4 \otimes q_3 + r_5 \otimes q_6 + r_6 \end{pmatrix} \quad (65)$$

$$= p \ast (q \ast r), \quad (66)$$

which proves the theorem.

E.2 Proof of Theorem 3.3

For each v, t , we can pre-compute $B_1 \mathbf{x}_{v,t}$ and $B_2 \mathbf{x}_{v,t+1}$. Accordingly, all the following parameters are pre-computed:

$$c_{v,t}^{(i,j,k,\ell)} = \begin{pmatrix} A_1 & A_2 & B_1 \mathbf{x}_{v+i,t+j} \\ A_3 & A_4 & B_2 \mathbf{x}_{v+k,t+\ell} \end{pmatrix}, \quad (67)$$

for all inputs $\mathbf{x}_{v,t}$ and $i, j, k, \ell \in \{0, 1\}$. Now, starting from $\begin{pmatrix} \mathbf{S}_{0,0}^{(1)} \\ \mathbf{S}_{0,0}^{(2)} \end{pmatrix} = \begin{pmatrix} I & I & 0 \\ I & I & 0 \end{pmatrix}$, we have:

$$\begin{pmatrix} \mathbf{S}_{0,1} \\ \mathbf{S}_{1,0} \end{pmatrix} = \begin{pmatrix} I & I & 0 \\ I & I & 0 \end{pmatrix} \ast \begin{pmatrix} \mathbf{A}_1 & \mathbf{A}_2 & \mathbf{B}_1 \mathbf{x}_{0,1} \\ \mathbf{A}_3 & \mathbf{A}_4 & \mathbf{B}_2 \mathbf{x}_{1,0} \end{pmatrix} \quad (68)$$

$$= \begin{pmatrix} \mathbf{A}_1 & \mathbf{A}_2 & \mathbf{B}_1 \mathbf{x}_{0,1} \\ \mathbf{A}_3 & \mathbf{A}_4 & \mathbf{B}_2 \mathbf{x}_{1,0} \end{pmatrix}. \quad (69)$$

Re-using operator \ast , we have:

$$\begin{pmatrix} \mathbf{S}_{1,1} \\ \mathbf{S}_{1,1} \end{pmatrix} = \begin{pmatrix} \mathbf{S}_{0,1} \\ \mathbf{S}_{1,0} \end{pmatrix} \ast \underbrace{\begin{pmatrix} \mathbf{A}_1 & \mathbf{A}_2 & \mathbf{B}_1 \mathbf{x}_{1,1} \\ \mathbf{A}_3 & \mathbf{A}_4 & \mathbf{B}_2 \mathbf{x}_{1,1} \end{pmatrix}}_{\text{Pre-computed}} \quad (70)$$

$$= \begin{pmatrix} \mathbf{A}_1^2 & \mathbf{A}_2^2 & \mathbf{A}_1 \mathbf{B}_1 \mathbf{x}_{0,1} + \mathbf{A}_2 \mathbf{B}_2 \mathbf{x}_{1,0} + \mathbf{B}_1 \mathbf{x}_{1,1} \\ \mathbf{A}_3^2 & \mathbf{A}_4^2 & \mathbf{A}_3 \mathbf{B}_1 \mathbf{x}_{0,1} + \mathbf{A}_4 \mathbf{B}_2 \mathbf{x}_{1,0} + \mathbf{B}_2 \mathbf{x}_{1,1} \end{pmatrix} \quad (71)$$

Looking at the third element of each row, these elements are calculating the hidden states of the recurrent (it can be shown by a straightforward induction). Accordingly, using this operation, we can recursively calculate the the outputs of 2D SSM.

However, using [Theorem 3.2](#), we know that this is an associative operation, so instead of calculating in the recurrent form, we can use parallel pre-fix sum make this computation parallel, decreasing the sequential operations required to calculate the hidden states. Note that since our above operation can model the problem as an parallel prefix, all the algorithms for this problem can be used to enhance the efficiency.

E.3 Proof of [Theorem 3.4](#)

To prove this theorem, we need to (1) show that Chimera can recover SpaceTime. Given this, since SpaceTime is capable of recovering ARIMA ([Bartholomew 1971](#)), exponential smoothing ([Winters 1960](#)), and controllable linear time-invariant systems ([C.-T. Chen 1984](#)), we can conclude that Chimera can also recover these methods. Then, (2) we need to prove that Chimera can recover SARIMA. This is the model that SpaceTime is not capable of recovering due to the additional seasonal terms.

Note that using $\mathbf{A}_2 = \mathbf{A}_3 = \mathbf{A}_4 = 0$, results in a 1D SSM, with companion matrix as the structure of \mathbf{A}_1 , which is SpaceTime. Accordingly, SpaceTime is a special case of Chimera when the recurrence only happen along the time direction.

Note that as discussed in [Proposition 3.1](#), multiplying the discretization parameter Δ results in multiplying the steps. Accordingly, using s as the Δ in our seasonal module and also letting $\mathbf{A}_2 = \mathbf{A}_3 = \mathbf{A}_4 = 0$ for the seasonal module, we can model the seasonal terms in the formulation of $\text{SAR}(p, q, s)$, meaning that Chimera can also recover SARIMA which is ARIMA with seasonal terms. Note that the reason that Chimera is capable of such modeling is that it uses two heads separately for trend and seasonal terms. Therefore, using different discretization parameters, each can model their own corresponding terms in $\text{SAR}(p, q, s)$.

E.4 Proof of [Theorem 3.5](#)

Similar to the above, using $\mathbf{A}_2 = \mathbf{A}_3 = 0$, our formulation is equivalent to S4D, while we use diagonal matrices as the structure of \mathbf{A}_1 . Similarly, as discussed by Behrouz, Santacatterina, and Zabih ([2024](#)), MambaMixer is equivalent to S4ND but on patched data. Using our [Theorem 5](#), we can recover linear layers, resulting in recovering TSMixer by setting $\mathbf{A}_2 = \mathbf{A}_3 = 0$.

E.5 Proof of [Theorem 3.6](#)

We in fact will show that restricting Chimera results in recovering 2DSSM ([Baron, Zimmerman, and Wolf 2024](#)). As discussed earlier, this method do not use discretization and initially starts from a discrete system. Also, it uses input-independent parameters. Therefore, we use $\text{Linear}_{\Delta_1}(\cdot) = \text{Linear}_{\Delta_2}(\cdot)$ as broadcast function, and restrict Chimera to have input-independent parameters, then Chimera can recover 2DSSM ([Baron, Zimmerman, and Wolf 2024](#)).

F Experimental Settings

We provide the description of datasets in [Table 7](#).

Table 7: Dataset descriptions. The dataset size is organized in (Train, Validation, Test).

Tasks	Dataset	Dim	Series Length	Dataset Size	Information (Frequency)
Forecasting (Long-term)	ETTm1, ETTm2	7	{96, 192, 336, 720}	(34465, 11521, 11521)	Electricity (15 mins)
	ETTTh1, ETTTh2	7	{96, 192, 336, 720}	(8545, 2881, 2881)	Electricity (15 mins)
	Electricity	321	{96, 192, 336, 720}	(18317, 2633, 5261)	Electricity (Hourly)
	Traffic	862	{96, 192, 336, 720}	(12185, 1757, 3509)	Transportation (Hourly)
	Weather	21	{96, 192, 336, 720}	(36792, 5271, 10540)	Weather (10 mins)
	Exchange	8	{96, 192, 336, 720}	(5120, 665, 1422)	Exchange rate (Daily)
	ILI	7	{24, 36, 48, 60}	(617, 74, 170)	Illness (Weekly)
Forecasting (short-term)	M4-Yearly	1	6	(23000, 0, 23000)	Demographic
	M4-Quarterly	1	8	(24000, 0, 24000)	Finance
	M4-Monthly	1	18	(48000, 0, 48000)	Industry
	M4-Weakly	1	13	(359, 0, 359)	Macro
	M4-Daily	1	14	(4227, 0, 4227)	Micro
	M4-Hourly	1	48	(414, 0, 414)	Other
Classification (UEA)	EthanolConcentration	3	1751	(261, 0, 263)	Alcohol Industry
	FaceDetection	144	62	(5890, 0, 3524)	Face (250Hz)
	Handwriting	3	152	(150, 0, 850)	Handwriting
	Heartbeat	61	405	(204, 0, 205)	Heart Beat
	JapaneseVowels	12	29	(270, 0, 370)	Voice
	PEMS-SF	963	144	(267, 0, 173)	Transportation (Daily)
	SelfRegulationSCP1	6	896	(268, 0, 293)	Health (256Hz)
	SelfRegulationSCP2	7	1152	(200, 0, 180)	Health (256Hz)
	SpokenArabicDigits	13	93	(6599, 0, 2199)	Voice (11025Hz)
	UWaveGestureLibrary	3	315	(120, 0, 320)	Gesture
Anomaly Detection	SMD	38	100	(566724, 141681, 708420)	Server Machine
	MSL	55	100	(44653, 11664, 73729)	Spacecraft
	SMAP	25	100	(108146, 27037, 427617)	Spacecraft
	SWaT	51	100	(396000, 99000, 449919)	Infrastructure
	PSM	25	100	(105984, 26497, 87841)	Server Machine

F.1 Baselines

In our experiments, we use the following baselines:

- [Table 8](#): TSM2 (Behrouz, Santacatterina, and Zabih [2024](#)), Simba (Badri N. Patro and Vijay S. Agneeswaran [2024](#)), TCN (Luo and X. Wang [2024](#)), iTransformer (Yong Liu, Hu, et al. [2024](#)), RLinear (Z. Li et al. [2023](#)), PatchTST (Nie et al. [2023](#)), Crossformer (Y. Zhang and Yan [2023](#)), TiDE (Das et al. [2023](#)), TimesNet (H. Wu, Hu, et al. [2023](#)), DLinear (Zeng et al. [2023b](#)), SCINet (M. Liu et al. [2022](#)), FEDformer (T. Zhou, Z. Ma, Wen, X. Wang, et al. [2022](#)), Stationary (Yong Liu, H. Wu, et al. [2022a](#)), Autoformer (H. Wu, Xu, et al. [2021](#))

- **Table 9:** ModernTCN (Luo and X. Wang 2024), PatchTST (Nie et al. 2023), TimesNet (H. Wu, Hu, et al. 2023), N-HiTS (Challu et al. 2022), N-BEATS* (Oreshkin et al. 2019), ETSformer (Woo et al. 2022), LightTS (T. Zhang et al. 2022), DLinear (Zeng et al. 2023a), FEDformer (T. Zhou, Z. Ma, Wen, X. Wang, et al. 2022), Stationary (Yong Liu, H. Wu, et al. 2022b), Autoformer (H. Wu, Xu, et al. 2021), Pyraformer (S. Liu et al. 2021), Informer (H. Zhou et al. 2021), Reformer (Kitaev, Kaiser, and Levskaya 2020), LSTM (Hochreiter and J. Schmidhuber 1997)
- **Table 10:** LSTM (Hochreiter and J. Schmidhuber 1997), LSTNet (Lai et al. 2018), LSSL (Gu, Goel, and Re 2022), Trans.former (Vaswani et al. 2017), Reformer (Kitaev, Kaiser, and Levskaya 2020), Informer (H. Zhou et al. 2021), Pyraformer (S. Liu et al. 2021), Autoformer (H. Wu, Xu, et al. 2021), Station. (Yong Liu, H. Wu, et al. 2022b), FEDformer (T. Zhou, Z. Ma, Wen, X. Wang, et al. 2022), ETSformer (Woo et al. 2022), Flowformer (H. Wu, J. Wu, et al. 2022), DLinear (Zeng et al. 2023a), LightTS. (T. Zhang et al. 2022), TimesNet (H. Wu, Hu, et al. 2023), PatchTST (Nie et al. 2023), MTCN (Luo and X. Wang 2024)

For the results of the baselines, we re-use the results reported by H. Wu, Hu, et al. (2023), or from the original cited papers.

G Additional Experimental Results

G.1 Long Term Forecasting Full Results

The complete results of long term forecasting are reported in Table 8.

G.2 Short-Term Forecasting

The complete results of short term forecasting are reported in Table 9.

G.3 Classification

The complete results of time series classification are reported in Table 10.

G.4 Anomaly Detection

The complete results of anomaly detection tasks are reported in Table 11.

Table 8: Long-term forecasting task with different horizons H . The first, second, and third best results are highlighted in **red** (bold), **orange** (underline), and **purple**.

		Chimera (ours)	TSM2 2024	Simba 2024	TCN 2024	iTransformer 2024	RLinear 2023	PatchTST 2023	Crossformer 2023	TiDE 2023	TimesNet 2023	DLinear 2023	SCINet 2022	FEDformer 2022	Stationary 2022	Autoformer 2021
		MSE MAE	MSE MAE	MSE MAE	MSE MAE	MSE MAE	MSE MAE	MSE MAE	MSE MAE	MSE MAE	MSE MAE	MSE MAE	MSE MAE	MSE MAE	MSE MAE	MSE MAE
ETm1	96	<u>0.293</u> 0.351	0.322 -	0.324 0.360	0.292 0.346	0.334 0.368	0.355 0.376	0.329 0.367	0.404 0.426	0.364 0.387	0.338 0.375	0.345 0.372	0.418 0.438	0.379 0.419	0.386 0.398	0.505 0.475
	192	0.329 0.362	0.349 -	0.363 0.382	<u>0.332</u> 0.368	0.377 0.391	0.391 0.392	0.367 0.385	0.450 0.451	0.398 0.404	0.374 0.387	0.380 0.389	0.439 0.450	0.426 0.441	0.459 0.444	0.553 0.496
	336	0.352 0.383	0.366 -	0.395 0.405	0.365 0.391	0.426 0.420	0.424 0.415	0.399 0.410	0.532 0.515	0.428 0.425	0.410 0.411	0.413 0.413	0.490 0.485	0.445 0.459	0.495 0.464	0.621 0.537
	720	0.408 0.412	0.407 -	0.451 0.437	0.416 0.417	0.491 0.459	0.487 0.450	0.454 0.439	0.666 0.589	0.487 0.461	0.478 0.450	0.474 0.453	0.595 0.550	0.543 0.490	0.585 0.516	0.671 0.561
	Avg	0.345 0.377	0.361 -	0.383 0.396	0.351 0.381	0.407 0.410	0.414 0.407	0.387 0.400	0.513 0.496	0.419 0.419	0.400 0.406	0.403 0.407	0.485 0.481	0.448 0.452	0.481 0.456	0.588 0.517
ETm2	96	0.168 0.261	0.173 -	0.177 0.263	0.166 0.256	0.180 0.264	0.182 0.265	0.175 0.259	0.287 0.366	0.207 0.305	0.187 0.267	0.193 0.292	0.286 0.377	0.203 0.287	0.192 0.274	0.255 0.339
	192	0.215 0.289	0.230 -	0.245 0.306	0.222 0.293	0.250 0.309	0.246 0.304	0.241 0.302	0.414 0.492	0.290 0.364	0.249 0.309	0.284 0.362	0.399 0.445	0.269 0.328	0.280 0.339	0.281 0.340
	336	0.278 0.337	0.279 -	0.304 0.343	0.272 0.324	0.311 0.348	0.307 0.342	0.305 0.343	0.597 0.542	0.377 0.422	0.321 0.351	0.369 0.427	0.637 0.591	0.325 0.366	0.334 0.361	0.339 0.372
	720	0.341 0.378	0.388 -	0.400 0.399	0.351 0.381	0.412 0.407	0.407 0.398	0.402 0.400	1.730 1.042	0.558 0.524	0.408 0.403	0.554 0.522	0.960 0.735	0.421 0.415	0.417 0.413	0.433 0.432
	Avg	0.250 0.316	0.267 -	0.271 0.327	0.253 0.314	0.288 0.332	0.286 0.327	0.281 0.326	0.757 0.610	0.358 0.404	0.291 0.333	0.350 0.401	0.571 0.537	0.305 0.349	0.306 0.347	0.327 0.371
ETTh1	96	0.362 0.391	0.375 -	0.379 0.395	0.368 0.394	0.386 0.405	0.386 0.395	0.414 0.419	0.423 0.448	0.479 0.464	0.384 0.402	0.386 0.400	0.654 0.599	0.376 0.419	0.513 0.491	0.449 0.459
	192	0.398 0.415	0.398 -	0.432 0.424	0.405 0.413	0.441 0.436	0.437 0.424	0.460 0.445	0.471 0.474	0.525 0.492	0.436 0.429	0.437 0.432	0.719 0.631	0.420 0.448	0.534 0.504	0.500 0.482
	336	0.402 0.416	0.419 -	0.473 0.443	0.391 0.412	0.487 0.458	0.479 0.446	0.501 0.466	0.570 0.546	0.565 0.515	0.491 0.469	0.481 0.459	0.778 0.659	0.459 0.465	0.588 0.535	0.521 0.496
	720	0.458 0.477	0.422 -	0.483 0.469	0.450 0.461	0.503 0.491	0.481 0.470	0.500 0.488	0.653 0.621	0.594 0.558	0.521 0.500	0.519 0.516	0.836 0.699	0.506 0.507	0.643 0.616	0.514 0.512
	Avg	0.405 0.424	0.403 -	0.441 0.432	0.404 0.420	0.454 0.447	0.446 0.434	0.469 0.454	0.529 0.522	0.541 0.507	0.458 0.450	0.456 0.452	0.747 0.647	0.440 0.460	0.570 0.537	0.496 0.487
ETTh2	96	0.257 0.325	0.253 -	0.290 0.339	0.263 0.332	0.297 0.349	0.288 0.338	0.302 0.348	0.745 0.584	0.400 0.440	0.340 0.374	0.333 0.387	0.707 0.621	0.358 0.397	0.476 0.458	0.346 0.388
	192	0.314 0.369	0.334 -	0.373 0.390	0.320 0.374	0.380 0.400	0.374 0.390	0.388 0.400	0.877 0.656	0.528 0.509	0.402 0.414	0.477 0.476	0.860 0.689	0.429 0.439	0.512 0.493	0.456 0.452
	336	0.316 0.381	0.347 -	0.376 0.406	0.313 0.376	0.428 0.432	0.415 0.426	0.426 0.433	1.043 0.731	0.643 0.571	0.452 0.452	0.594 0.541	1.000 0.744	0.496 0.487	0.552 0.551	0.482 0.486
	720	0.388 0.427	0.401 -	0.407 0.431	0.392 0.433	0.427 0.445	0.420 0.440	0.431 0.446	1.104 0.763	0.874 0.679	0.462 0.468	0.831 0.657	1.249 0.838	0.463 0.474	0.562 0.560	0.515 0.511
	Avg	0.318 0.375	0.333 -	0.361 0.391	0.322 0.379	0.383 0.407	0.374 0.398	0.387 0.407	0.942 0.684	0.611 0.550	0.414 0.427	0.559 0.515	0.954 0.723	0.437 0.449	0.526 0.516	0.450 0.459
ECL	96	0.132 0.234	0.142 -	0.165 0.253	0.129 0.226	0.148 0.240	0.201 0.281	0.181 0.270	0.219 0.314	0.237 0.329	0.168 0.272	0.197 0.282	0.247 0.345	0.193 0.308	0.169 0.273	0.201 0.317
	192	0.144 0.223	0.153 -	0.173 0.262	0.143 0.239	0.162 0.253	0.201 0.283	0.188 0.274	0.231 0.322	0.236 0.330	0.184 0.289	0.196 0.285	0.257 0.355	0.201 0.315	0.182 0.286	0.222 0.334
	336	0.156 0.259	0.175 -	0.188 0.277	0.161 0.259	0.178 0.269	0.215 0.298	0.204 0.293	0.246 0.337	0.249 0.344	0.198 0.300	0.209 0.301	0.269 0.369	0.214 0.329	0.200 0.304	0.231 0.338
	720	0.184 0.280	0.209 -	0.214 0.305	0.191 0.286	0.225 0.317	0.257 0.331	0.246 0.324	0.280 0.363	0.284 0.373	0.220 0.320	0.245 0.333	0.299 0.390	0.246 0.355	0.222 0.321	0.254 0.361
	Avg	0.154 0.249	0.169 -	0.185 0.274	0.156 0.253	0.178 0.270	0.219 0.298	0.205 0.290	0.244 0.334	0.251 0.344	0.192 0.295	0.212 0.300	0.268 0.365	0.214 0.327	0.193 0.296	0.227 0.338
Exchange	96	0.077 0.198	0.163 -	- -	0.080 0.196	0.086 0.206	0.093 0.217	0.088 0.205	0.256 0.367	0.094 0.218	0.107 0.234	0.088 0.218	0.267 0.396	0.148 0.278	0.111 0.237	0.197 0.323
	192	0.159 0.270	0.229 -	- -	0.166 0.288	0.177 0.299	0.184 0.307	0.176 0.299	0.470 0.509	0.184 0.307	0.226 0.344	0.176 0.315	0.351 0.459	0.271 0.315	0.219 0.335	0.300 0.369
	336	0.311 0.344	0.383 -	- -	0.307 0.398	0.331 0.417	0.351 0.432	0.301 0.397	1.268 0.883	0.349 0.431	0.367 0.448	0.313 0.427	1.324 0.853	0.460 0.427	0.421 0.476	0.509 0.524
	720	0.697 0.623	0.999 -	- -	0.656 0.582	0.847 0.691	0.886 0.714	0.901 0.714	1.767 1.068	0.852 0.698	0.964 0.746	0.839 0.695	1.058 0.797	1.195 0.695	1.092 0.769	1.447 0.941
	Avg	0.311 0.358	0.443 -	- -	0.302 0.366	0.360 0.403	0.378 0.417	0.367 0.404	0.940 0.707	0.370 0.413	0.416 0.443	0.354 0.414	0.750 0.626	0.519 0.429	0.461 0.454	0.613 0.539
Traffic	96	0.366 0.248	0.396 -	0.468 0.268	0.368 0.253	0.395 0.268	0.649 0.389	0.462 0.295	0.522 0.290	0.805 0.493	0.593 0.321	0.650 0.396	0.788 0.499	0.587 0.366	0.612 0.338	0.613 0.388
	192	0.394 0.292	0.408 -	0.413 0.317	0.379 0.261	0.417 0.276	0.601 0.366	0.466 0.296	0.530 0.293	0.756 0.474	0.617 0.336	0.598 0.370	0.789 0.505	0.604 0.373	0.613 0.340	0.616 0.382
	336	0.409 0.311	0.427 -	0.529 0.284	0.397 0.270	0.433 0.283	0.609 0.369	0.482 0.304	0.558 0.305	0.762 0.477	0.629 0.336	0.605 0.373	0.797 0.508	0.621 0.383	0.618 0.328	0.622 0.337
	720	0.443 0.294	0.449 -	0.564 0.297	0.440 0.296	0.467 0.302	0.647 0.387	0.514 0.322	0.589 0.328	0.719 0.449	0.640 0.350	0.645 0.394	0.841 0.523	0.626 0.382	0.653 0.355	0.660 0.408
	Avg	0.403 0.286	0.420 -	0.493 0.291	0.398 0.270	0.428 0.282	0.626 0.378	0.481 0.304	0.550 0.304	0.760 0.473	0.620 0.336	0.625 0.383	0.804 0.509	0.610 0.376	0.624 0.340	0.628 0.379
Weather	96	0.146 0.206	0.161 -	0.176 0.219	0.149 0.200	0.174 0.214	0.192 0.232	0.177 0.218	0.158 0.230	0.202 0.261	0.172 0.220	0.196 0.255	0.221 0.306	0.217 0.296	0.173 0.223	0.266 0.336
	192	0.189 0.239	0.208 -	0.222 0.260	0.196 0.245	0.221 0.254	0.240 0.271	0.225 0.259	0.206 0.277	0.242 0.298	0.219 0.261	0.237 0.296	0.261 0.340	0.276 0.336	0.245 0.285	0.307 0.367
	336	0.244 0.281	0.252 -	0.275 0.297	0.238 0.277	0.278 0.296	0.292 0.307	0.278 0.297	0.272 0.335	0.287 0.335	0.280 0.306	0.283 0.335	0.309 0.378	0.339 0.380	0.321 0.338	0.359 0.395
	720	0.297 0.309	0.337 -	0.350 0.349	0.314 0.334	0.358 0.347	0.364 0.353	0.354 0.348	0.398 0.418	0.351 0.386	0.365 0.359	0.345 0.381	0.377 0.427	0.403 0.428	0.414 0.410	0.419 0.428
	Avg	0.219 0.258	0.239 -	0.255 0.280	0.224 0.264	0.258 0.278	0.272 0.291	0.259 0.281	0.259 0.315	0.271 0.320	0.259 0.287	0.265 0.317	0.292 0.363	0.309 0.360	0.288 0.314	0.338 0.382

Table 9: Full results for the short-term forecasting task in the M4 dataset. *. in the Transformers indicates the name of *former. *Stationary* means the Non-stationary Transformer.

Models		Chimera (ours)	ModernTCN 2024	PatchTST 2023	TimesNet 2023	N-HiTS 2022	N-BEATS* 2019	ETS* 2022	LightTS 2022	DLinear 2023	FED* 2022	Stationary 2022	Auto* 2021	Pyra* 2021	In* 2021	Re* 2020	LSTM 1997
Yearly	SMAPE	13.107	<u>13.226</u>	13.258	13.387	13.418	13.436	18.009	14.247	16.965	13.728	13.717	13.974	15.530	14.727	16.169	176.040
	MASE	2.902	<u>2.957</u>	2.985	2.996	3.045	3.043	4.487	3.109	4.283	3.048	3.078	3.134	3.711	3.418	3.800	31.033
	OWA	0.767	<u>0.777</u>	0.781	0.786	0.793	0.794	1.115	0.827	1.058	0.803	0.807	0.822	0.942	0.881	0.973	9.290
Quarterly	SMAPE	9.892	<u>9.971</u>	10.179	10.100	10.202	10.124	13.376	11.364	12.145	10.792	10.958	11.338	15.449	11.360	13.313	172.808
	MASE	<u>1.105</u>	1.167	0.803	1.182	1.194	1.169	1.906	1.328	1.520	1.283	1.325	1.365	2.350	1.401	1.775	19.753
	OWA	<u>0.853</u>	0.878	0.803	0.890	0.899	0.886	1.302	1.000	1.106	0.958	0.981	1.012	1.558	1.027	1.252	15.049
Monthly	SMAPE	12.549	<u>12.556</u>	12.641	12.670	12.791	12.677	14.588	14.014	13.514	14.260	13.917	13.958	17.642	14.062	20.128	143.237
	MASE	0.914	<u>0.917</u>	0.930	0.933	0.969	0.937	1.368	1.053	1.037	1.102	1.097	1.103	1.913	1.141	2.614	16.551
	OWA	0.864	<u>0.866</u>	0.876	0.878	0.899	0.880	1.149	0.981	0.956	1.012	0.998	1.002	1.511	1.024	1.927	12.747
Others	SMAPE	4.685	<u>4.715</u>	4.946	4.891	5.061	4.925	7.267	15.880	6.709	4.954	6.302	5.485	24.786	24.460	32.491	186.282
	MASE	<u>3.007</u>	3.107	2.985	3.302	3.216	3.391	5.240	11.434	4.953	3.264	4.064	3.865	18.581	20.960	33.355	119.294
	OWA	0.983	<u>0.986</u>	1.044	1.035	1.040	1.053	1.591	3.474	1.487	1.036	1.304	1.187	5.538	5.013	8.679	38.411
Weighted Average	SMAPE	11.618	<u>11.698</u>	11.807	11.829	11.927	11.851	14.718	13.525	13.639	12.840	12.780	12.909	16.987	14.086	18.200	160.031
	MASE	1.528	<u>1.556</u>	1.590	1.585	1.613	1.599	2.408	2.111	2.095	1.701	1.756	1.771	3.265	2.718	4.223	25.788
	OWA	0.827	<u>0.838</u>	0.851	0.851	0.861	0.855	1.172	1.051	1.051	0.918	0.930	0.939	1.480	1.230	1.775	12.642

Table 10: Full results for the classification task (accuracy %). We omit “former” from the names of Transformer-based methods. For all methods, the standard deviation is less than 0.1%.

Datasets / Models	LSTM	LSTNet	LSSL	Trans.	Re.	In.	Pyra.	Auto.	Station.	FED.	/ETS.	/Flow.	DLinear	LightTS.	/TimesNet	PatchTST	MTCN	Chimera (ours)
	1997	2018	2022	2017	2020	2021	2021	2021	2022	2022	2022	2022	2023	2022	2023	2023	2024	
EthanolConcentration	32.3	39.9	31.1	32.7	31.9	31.6	30.8	31.6	32.7	31.2	28.1	33.8	32.6	29.7	35.7	32.8	36.3	39.8
FaceDetection	57.7	65.7	66.7	67.3	68.6	67.0	65.7	68.4	68.0	66.0	66.3	67.6	68.0	67.5	68.6	68.3	70.8	70.4
Handwriting	15.2	25.8	24.6	32.0	27.4	32.8	29.4	36.7	31.6	28.0	32.5	33.8	27.0	26.1	32.1	29.6	30.6	32.9
Heartbeat	72.2	77.1	72.7	76.1	77.1	80.5	75.6	74.6	73.7	73.7	71.2	77.6	75.1	75.1	78.0	74.9	77.2	81.3
JapaneseVowels	79.7	98.1	98.4	98.7	97.8	98.9	98.4	96.2	99.2	98.4	95.9	98.9	96.2	96.2	98.4	97.5	98.8	99.1
PEMS-SF	39.9	86.7	86.1	82.1	82.7	81.5	83.2	82.7	87.3	80.9	86.0	83.8	75.1	88.4	89.6	89.3	89.1	89.5
SelfRegulationSCP1	68.9	84.0	90.8	92.2	90.4	90.1	88.1	84.0	89.4	88.7	89.6	92.5	87.3	89.8	91.8	90.7	93.4	93.7
SelfRegulationSCP2	46.6	52.8	52.2	53.9	56.7	53.3	53.3	50.6	57.2	54.4	55.0	56.1	50.5	51.1	57.2	57.8	60.3	59.9
SpokenArabicDigits	31.9	100.0	100.0	98.4	97.0	100.0	99.6	100.0	100.0	100.0	100.0	98.8	81.4	100.0	99.0	98.3	98.7	100.0
UWaveGestureLibrary	41.2	87.8	85.9	85.6	85.6	85.6	83.4	85.9	87.5	85.3	85.0	86.6	82.1	80.3	85.3	85.8	86.7	86.7
Average Accuracy	48.6	71.8	70.9	71.9	71.5	72.1	70.8	71.1	72.7	70.7	71.0	73.0	67.5	70.4	73.6	72.5	<u>74.2</u>	75.3

Table 11: Full results for the anomaly detection task. The P, R and F1 represent the precision, recall and F1-score (%) respectively. A higher value of P, R and F1 indicates a better performance.

Datasets		SMD			MSL			SMAP			SWaT			PSM			Avg F1
Metrics		P	R	F1	P	R	F1	P	R	F1	P	R	F1	P	R	F1	(%)
LSTM	1997	78.52	65.47	71.41	78.04	86.22	81.93	91.06	57.49	70.48	78.06	91.72	84.34	69.24	99.53	81.67	77.97
Transformer	2017	83.58	76.13	79.56	71.57	87.37	78.68	89.37	57.12	69.70	68.84	96.53	80.37	62.75	96.56	76.07	76.88
LogTrans	2019	83.46	70.13	76.21	73.05	87.37	79.57	89.15	57.59	69.97	68.67	97.32	80.52	63.06	98.00	76.74	76.60
TCN	2019	84.06	79.07	81.49	75.11	82.44	78.60	86.90	59.23	70.45	76.59	95.71	85.09	54.59	99.77	70.57	77.24
Reformer	2020	82.58	69.24	75.32	85.51	83.31	84.40	90.91	57.44	70.40	72.50	96.53	82.80	59.93	95.38	73.61	77.31
Informer	2021	86.60	77.23	81.65	81.77	86.48	84.06	90.11	57.13	69.92	70.29	96.75	81.43	64.27	96.33	77.10	78.83
Anomaly*	2021	88.91	82.23	85.49	79.61	87.37	83.31	91.85	58.11	71.18	72.51	97.32	83.10	68.35	94.72	79.40	80.50
Pyraformer	2021	85.61	80.61	83.04	83.81	85.93	84.86	92.54	57.71	71.09	87.92	96.00	91.78	71.67	96.02	82.08	82.57
Autoformer	2021	88.06	82.35	85.11	77.27	80.92	79.05	90.40	58.62	71.12	89.85	95.81	92.74	99.08	88.15	93.29	84.26
LSSL	2022	78.51	65.32	71.31	77.55	88.18	82.53	89.43	53.43	66.90	79.05	93.72	85.76	66.02	92.93	77.20	76.74
Stationary	2022	88.33	81.21	84.62	68.55	89.14	77.50	89.37	59.02	71.09	68.03	96.75	79.88	97.82	96.76	97.29	82.08
DLinear	2023	83.62	71.52	77.10	84.34	85.42	84.88	92.32	55.41	69.26	80.91	95.30	87.52	98.28	89.26	93.55	82.46
ETSformer	2022	87.44	79.23	83.13	85.13	84.93	85.03	92.25	55.75	69.50	90.02	80.36	84.91	99.31	85.28	91.76	82.87
LightTS	2022	87.10	78.42	82.53	82.40	75.78	78.95	92.58	55.27	69.21	91.98	94.72	93.33	98.37	95.97	97.15	84.23
FEDformer	2022	87.95	82.39	85.08	77.14	80.07	78.57	90.47	58.10	70.76	90.17	96.42	93.19	97.31	97.16	97.23	84.97
TimesNet (I)	2023	87.76	82.63	85.12	82.97	85.42	84.18	91.50	57.80	70.85	88.31	96.24	92.10	98.22	92.21	95.21	85.49
TimesNet (R)	2023	88.66	83.14	85.81	83.92	86.42	85.15	92.52	58.29	71.52	86.76	97.32	91.74	98.19	96.76	97.47	86.34
CrossFormer	2023	83.6	76.61	79.70	84.68	83.71	84.19	92.04	55.37	69.14	88.49	93.48	90.92	97.16	89.73	93.30	83.45
PatchTST	2023	87.42	81.65	84.44	84.07	86.23	85.14	92.43	57.51	70.91	80.70	94.93	87.24	98.87	93.99	96.37	84.82
ModernTCN	2024	87.86	83.85	85.81	83.94	85.93	84.92	93.17	57.69	71.26	91.83	95.98	93.86	98.09	96.38	97.23	86.62
Chimera	(ours)	87.74	83.29	85.46	84.01	86.83	85.39	93.05	58.12	71.55	92.18	95.93	94.01	97.30	96.19	96.74	86.69